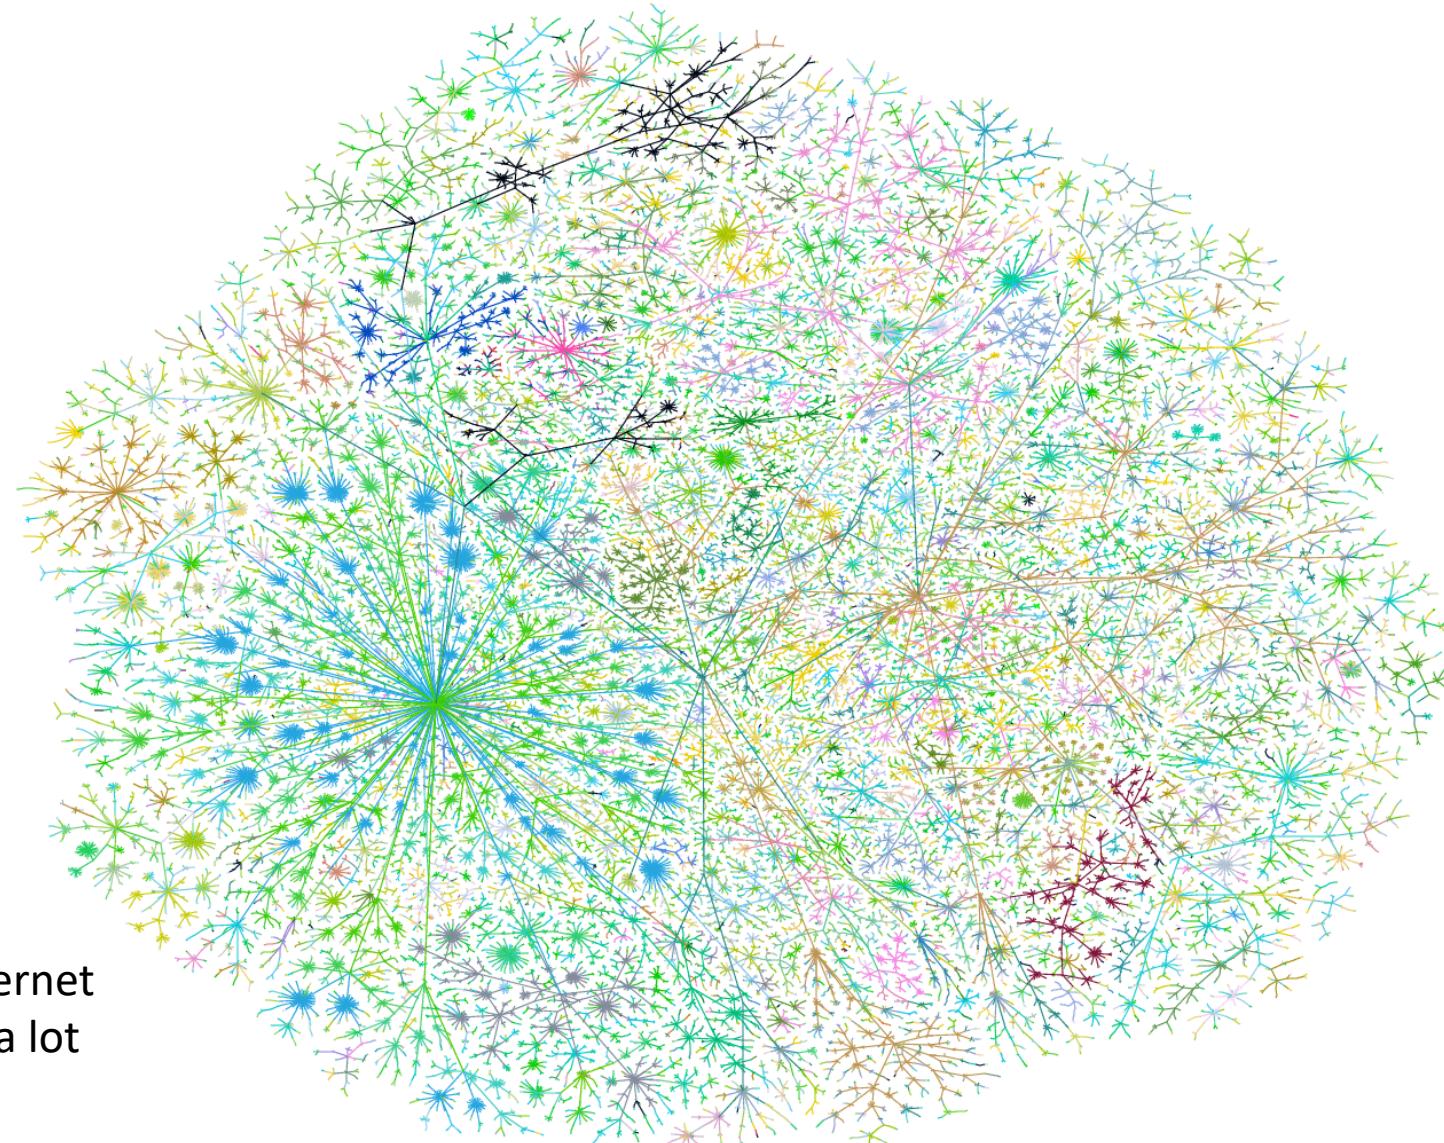


Graphs

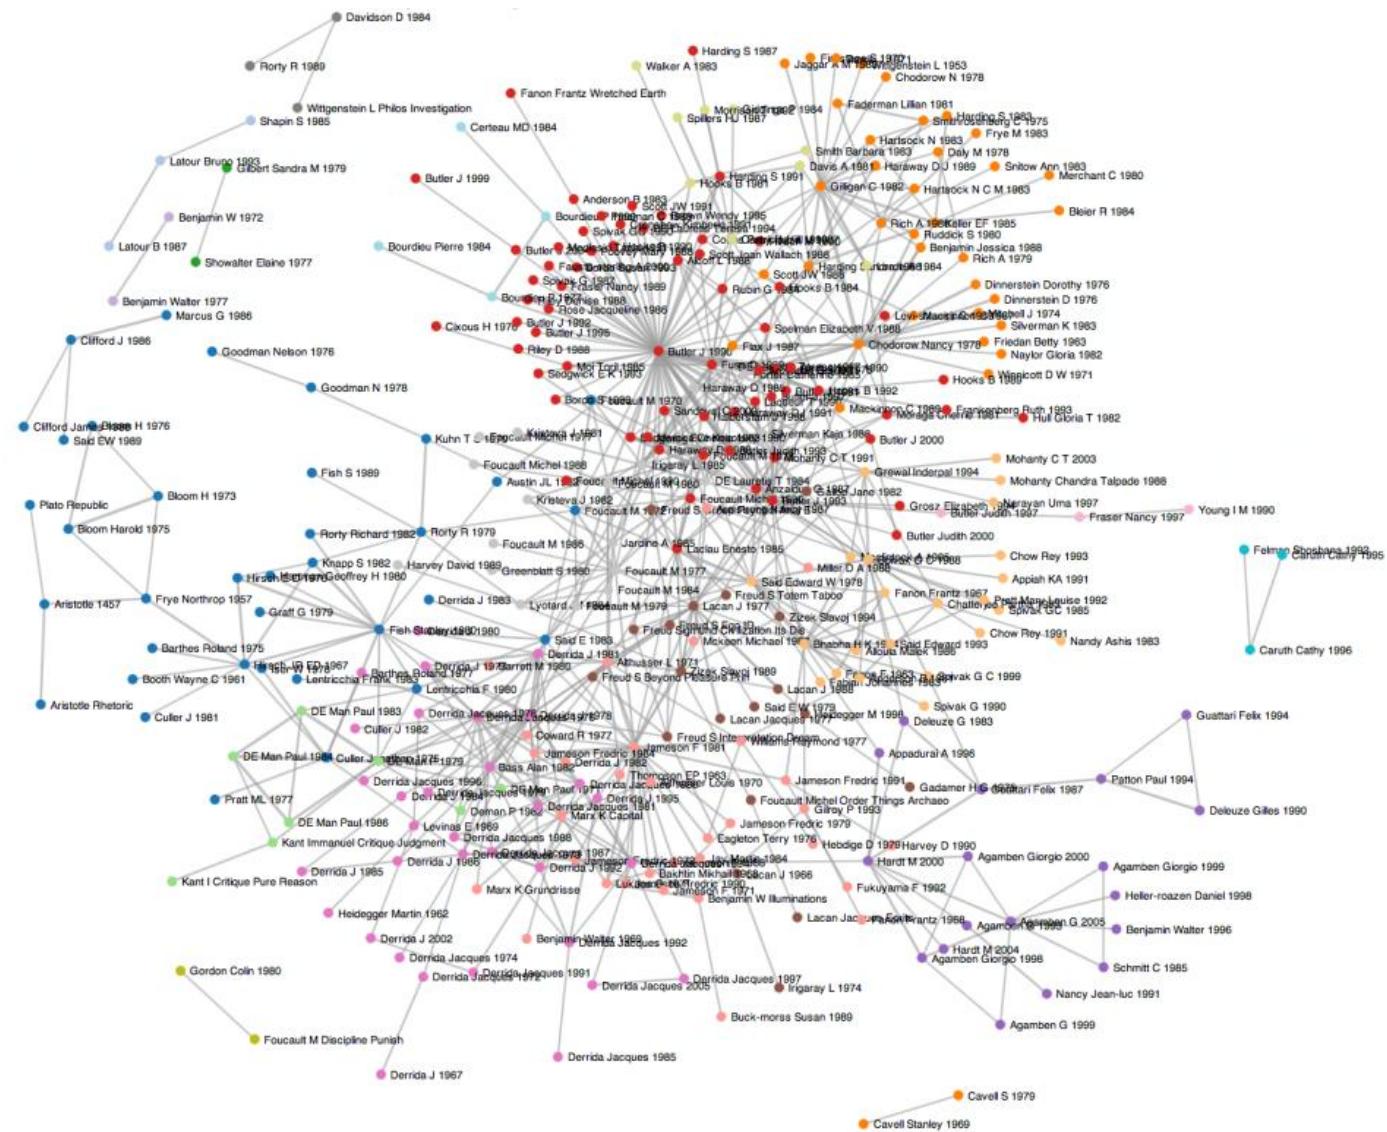
Graphs



Graph of the internet
(circa 1999...it's a lot
bigger now...)

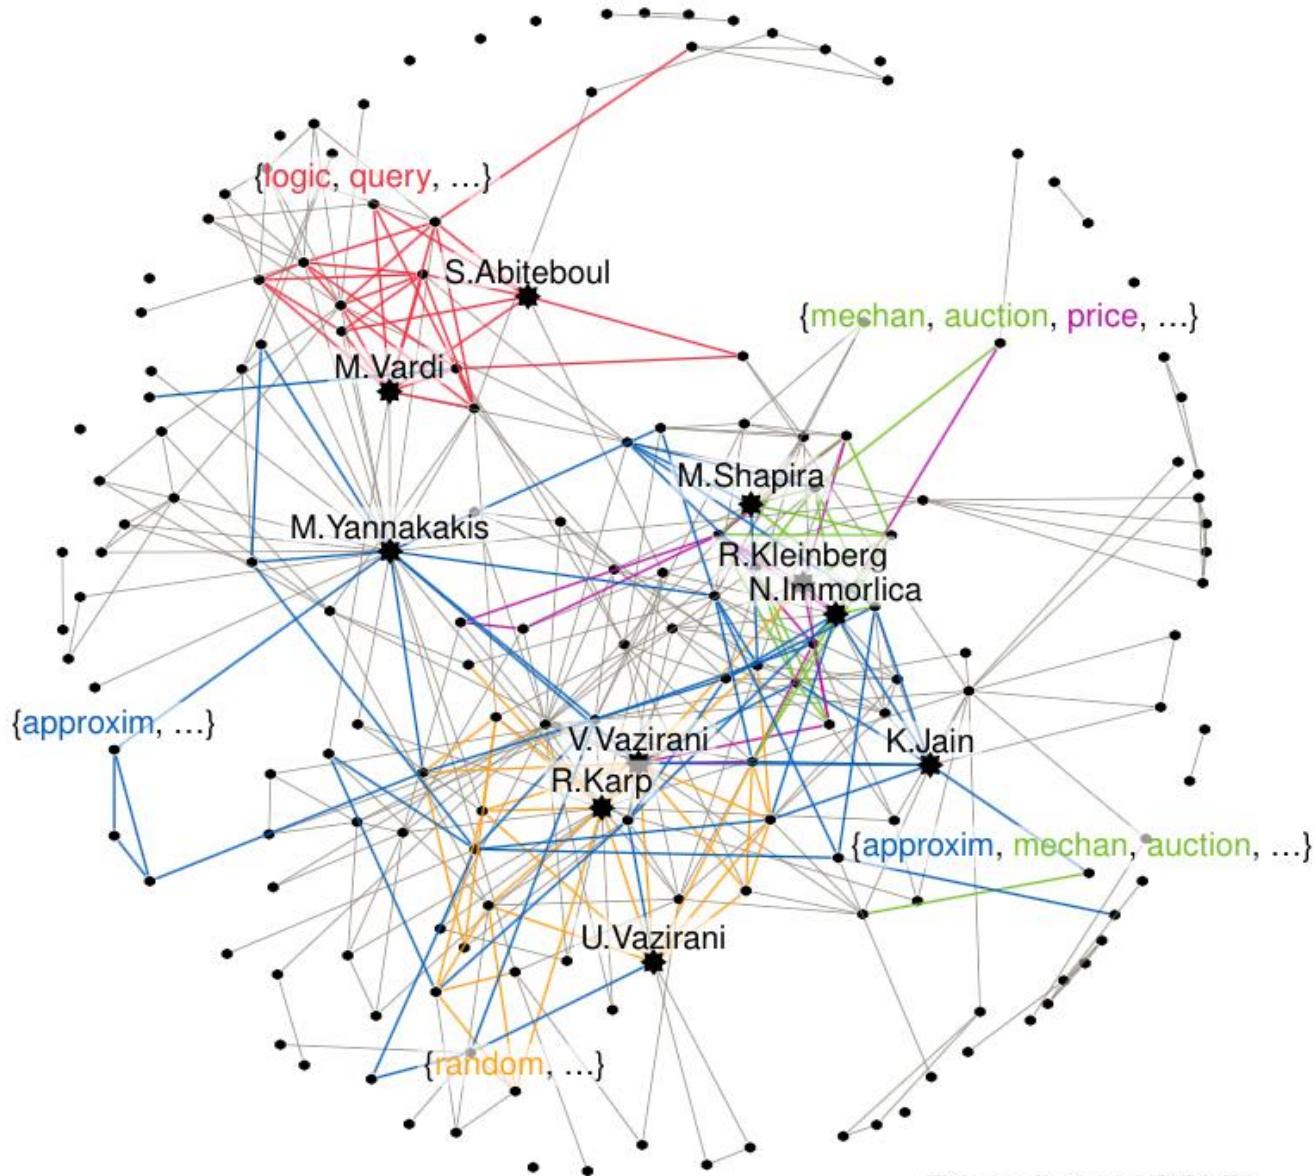
Graphs

Citation graph of literary theory academic papers



Graphs

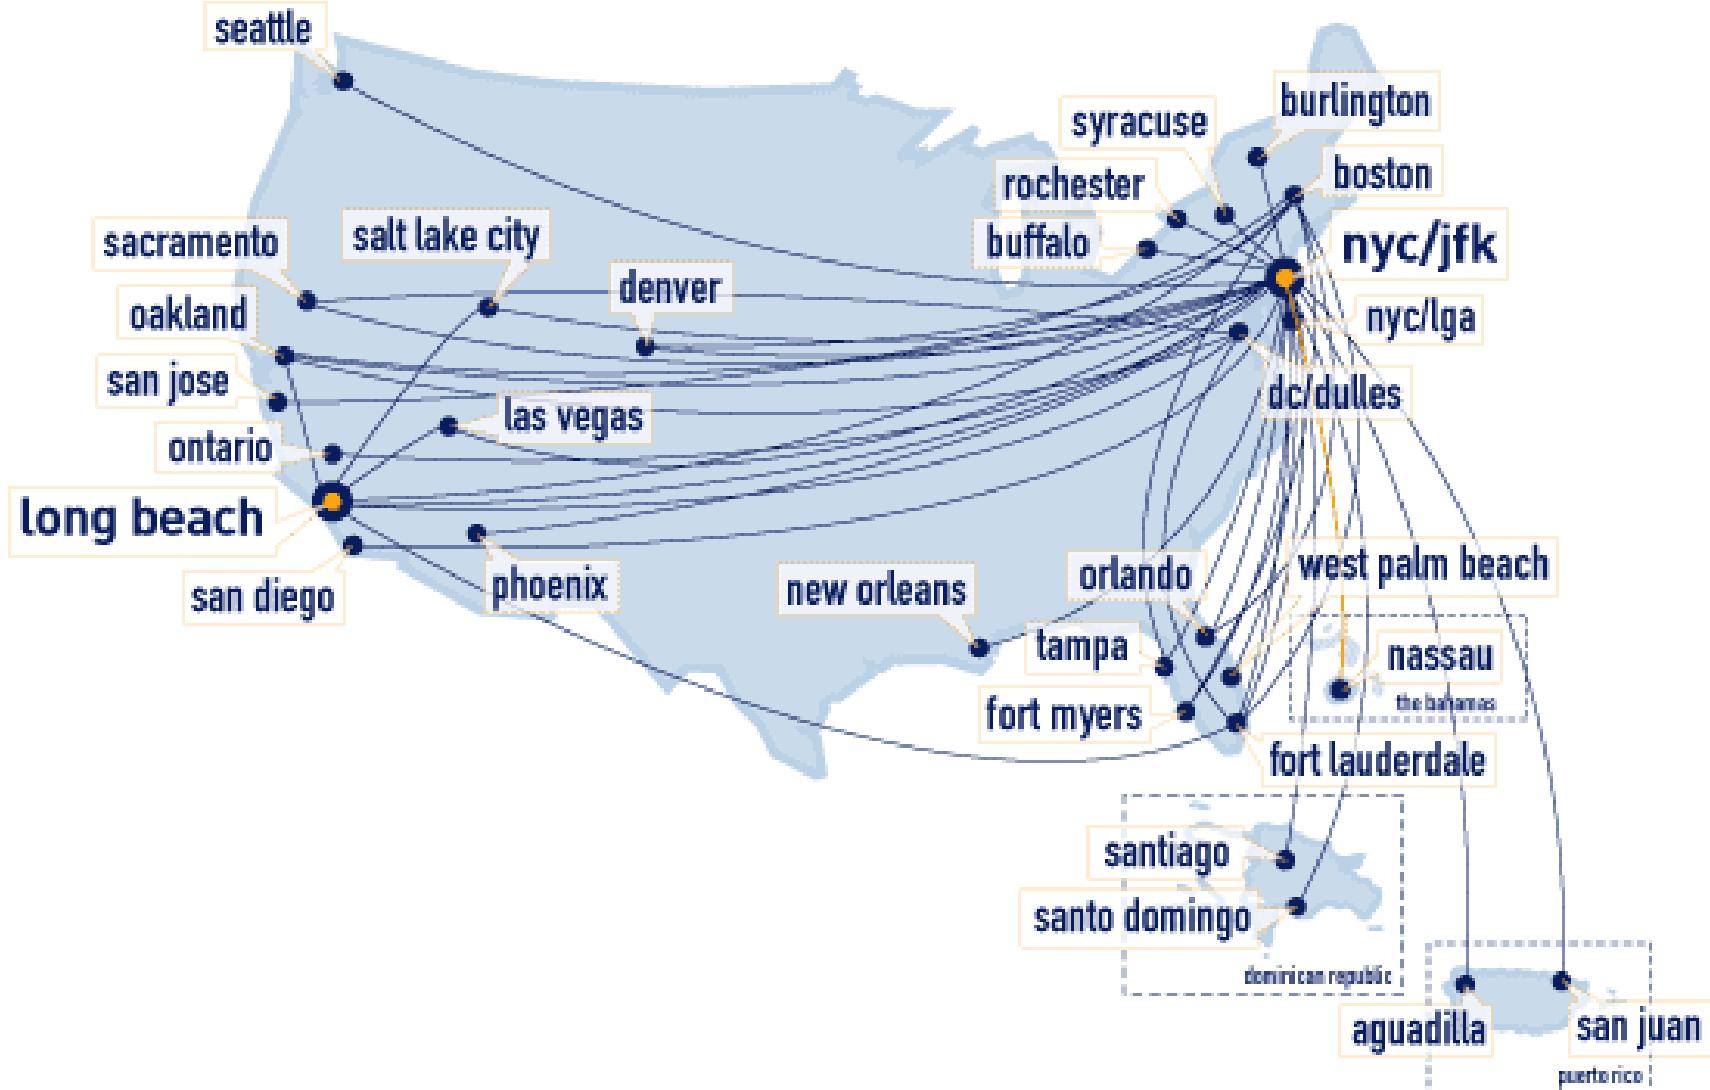
Theoretical Computer
Science academic
communities



Example from DBLP:
Communities within the co-authors of Christos H. Papadimitriou

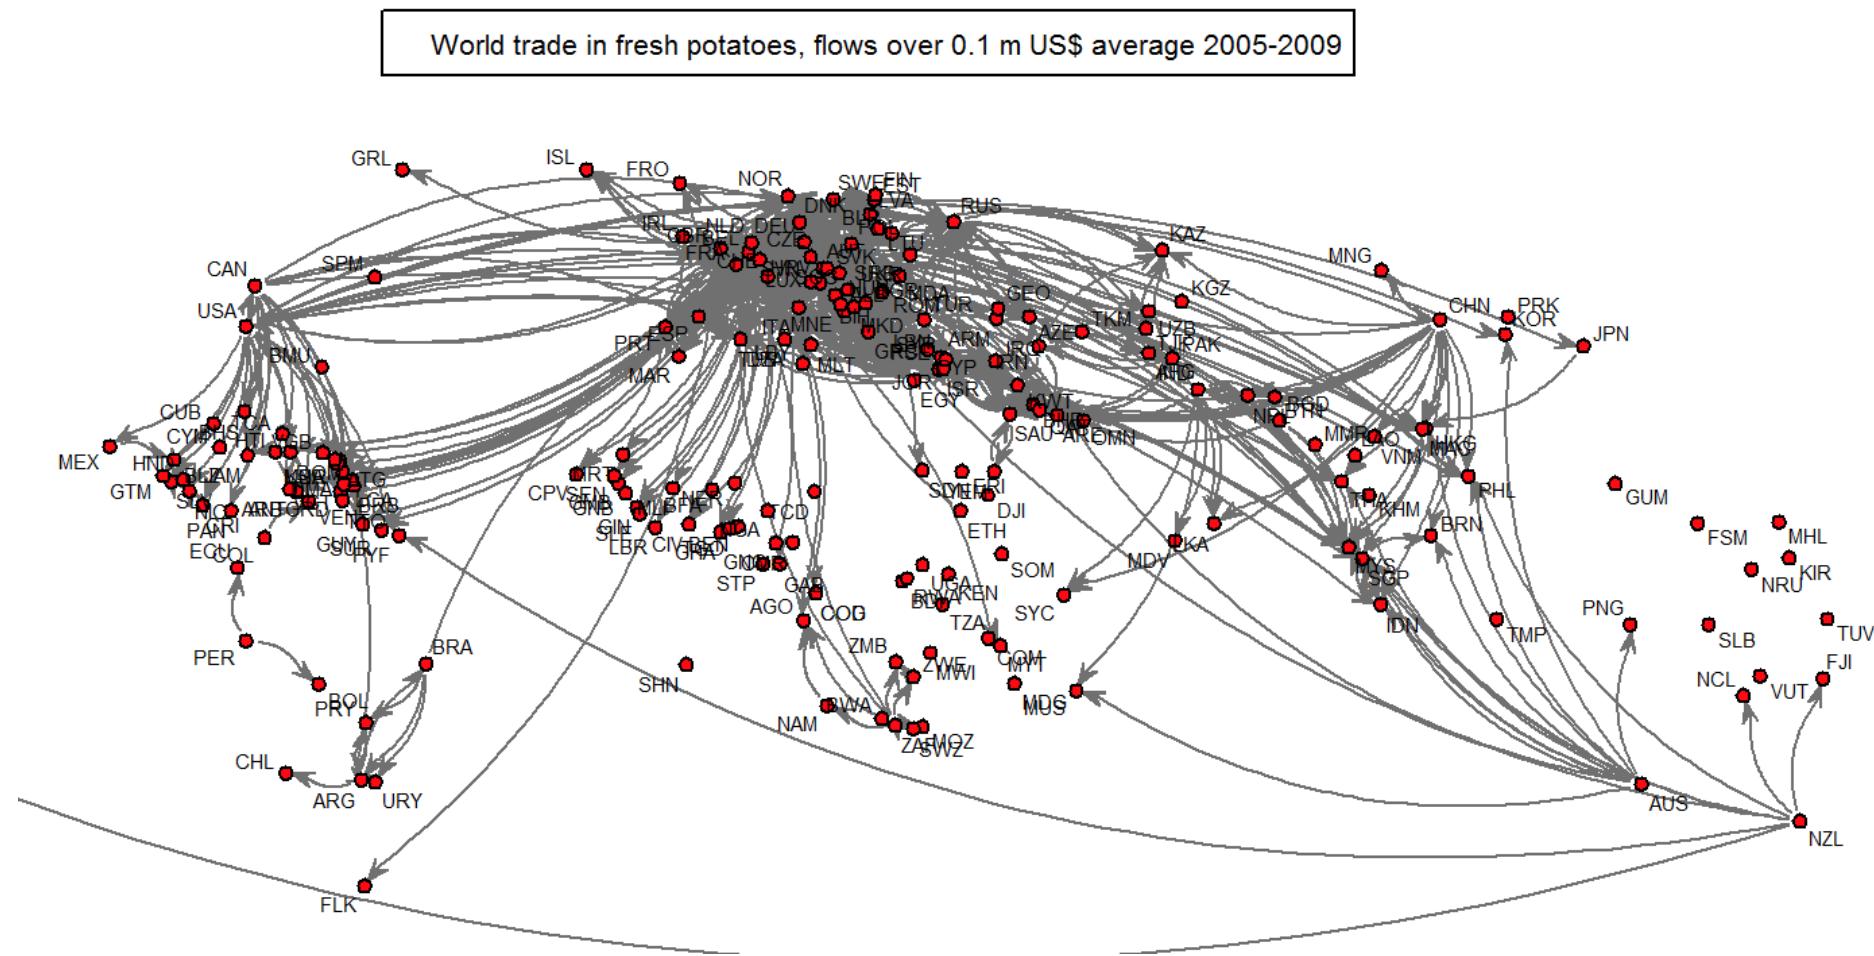
Graphs

jetblue flights

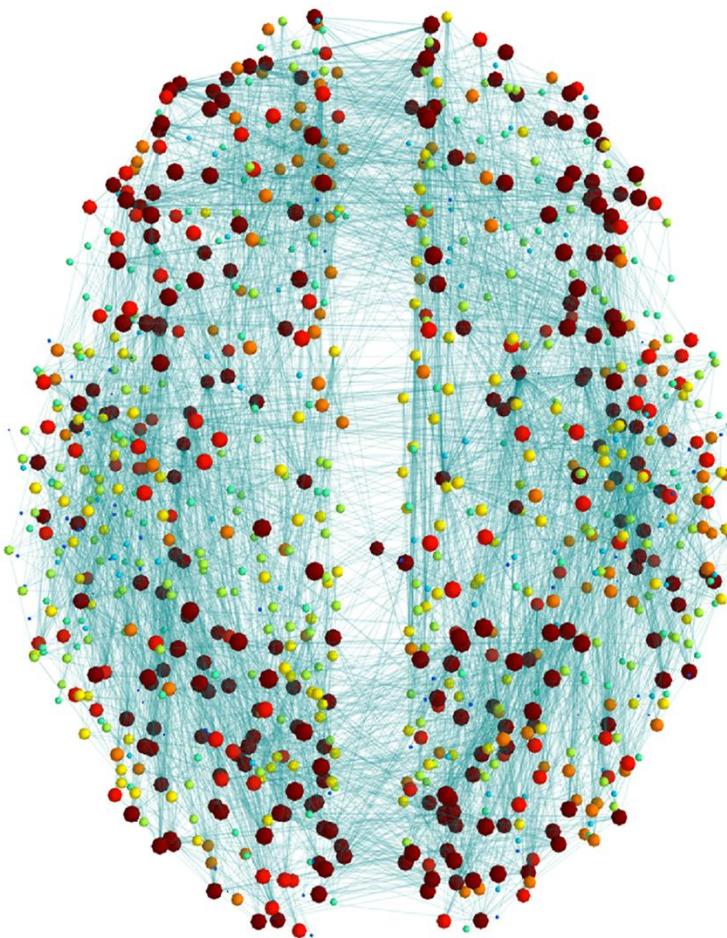
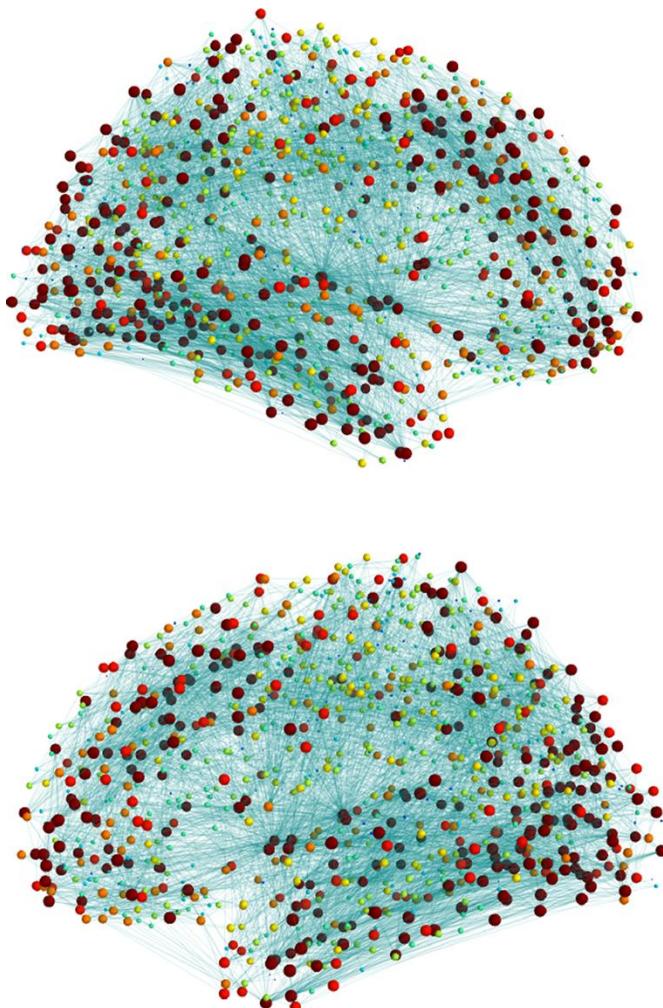


Graphs

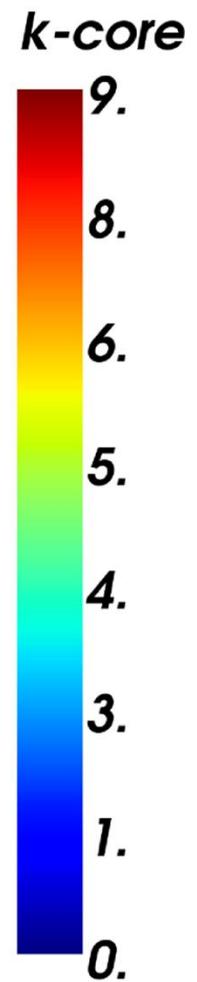
Potato trade



Graphs



Neural connections
in the brain

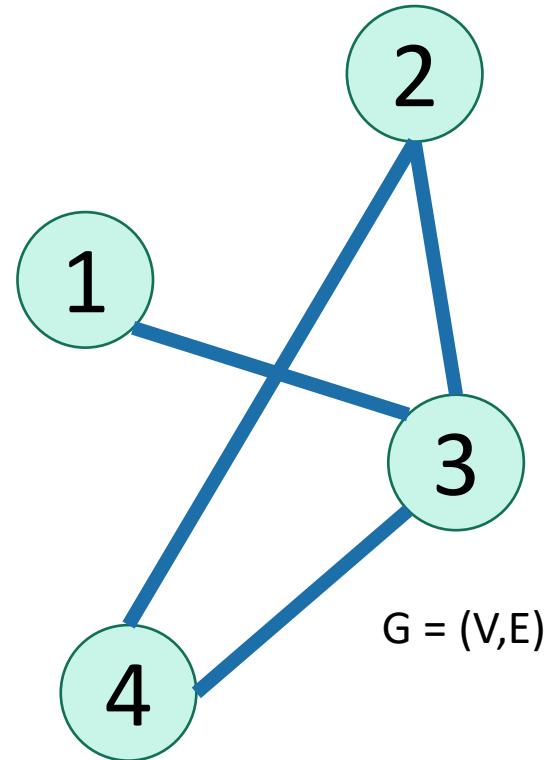


Graphs

- **There are a lot of graphs.**
- We want to answer questions about them.
 - Efficient routing?
 - Community detection/clustering?
 - An ordering that respects dependencies?
- This is what we'll do for the next several lectures.

Undirected Graphs

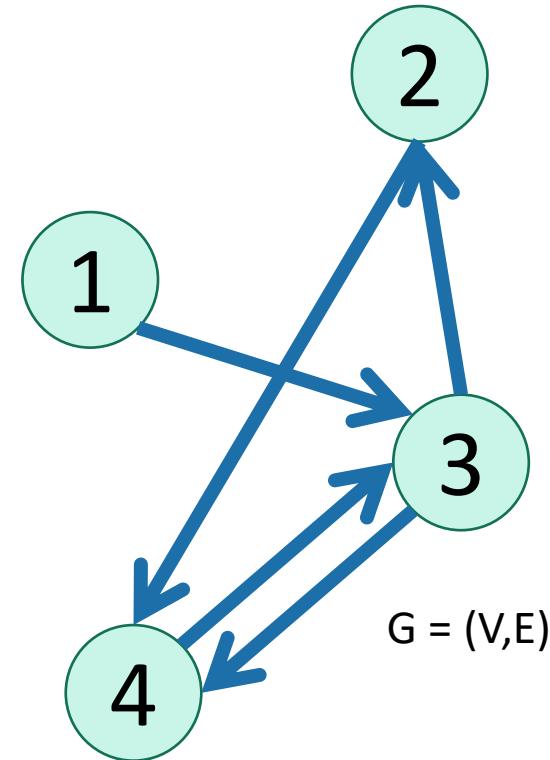
- Has **vertices** and **edges**
 - V is the set of vertices
 - E is the set of edges
 - Formally, a graph is $G = (V, E)$
- Example
 - $V = \{1, 2, 3, 4\}$
 - $E = \{ \{1, 3\}, \{2, 4\}, \{3, 4\}, \{2, 3\} \}$



- The **degree** of vertex 4 is 2.
 - There are 2 edges coming out.
- Vertex 4's **neighbors** are 2 and 3

Directed Graphs

- Has **vertices** and **edges**
 - V is the set of vertices
 - E is the set of **DIRECTED** edges
 - Formally, a graph is $G = (V,E)$
- Example
 - $V = \{1,2,3,4\}$
 - $E = \{ (1,3), (2,4), (3,4), (4,3), (3,2) \}$



- The **in-degree** of vertex 4 is 2.
- The **out-degree** of vertex 4 is 1.
- Vertex 4's **incoming neighbors** are 2
- Vertex 4's **outgoing neighbor** is 3.

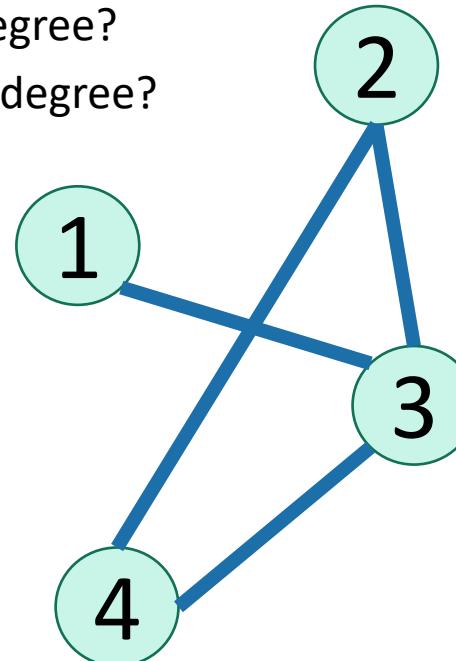
Graph Representation

How do we represent graphs?

- Option 1: adjacency matrix

$$\begin{matrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \end{matrix} & \left[\begin{matrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{matrix} \right] \end{matrix}$$

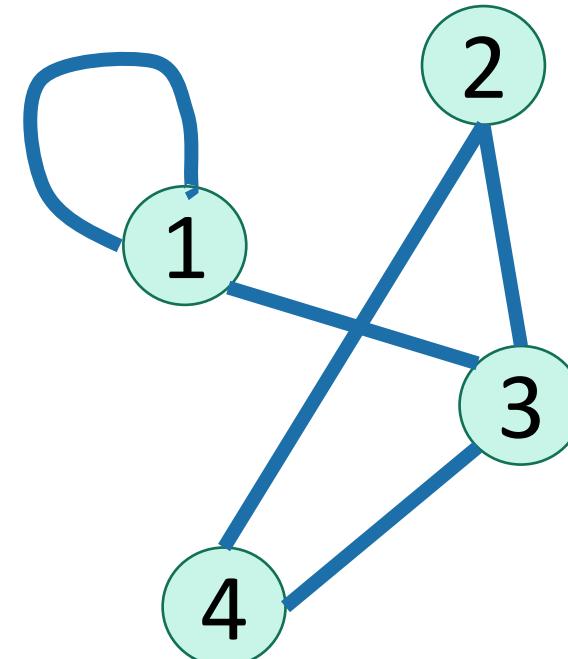
Indegree?
Outdegree?



How do we represent graphs?

- Option 1: adjacency matrix

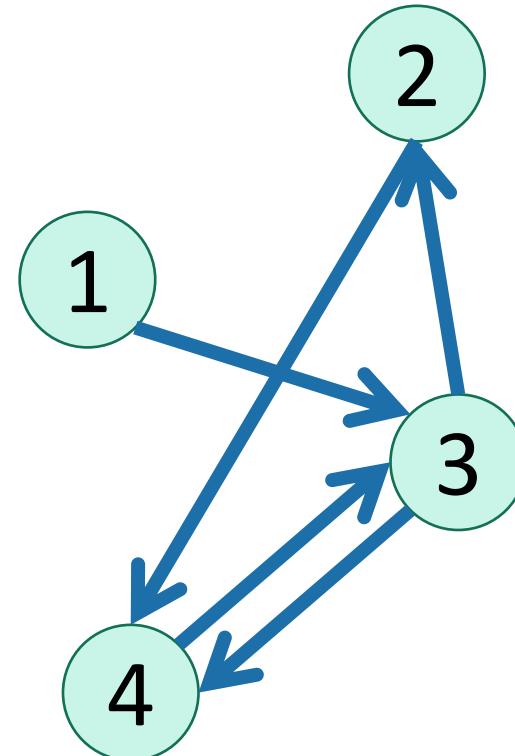
$$\begin{matrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \end{matrix} & \left[\begin{matrix} 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{matrix} \right] \end{matrix}$$



How do we represent graphs?

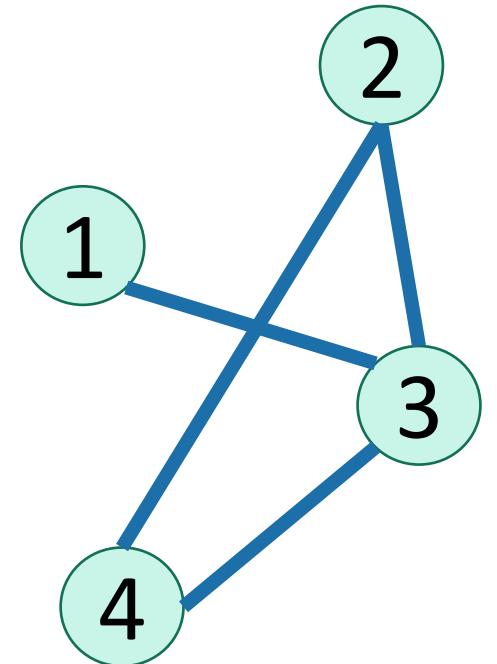
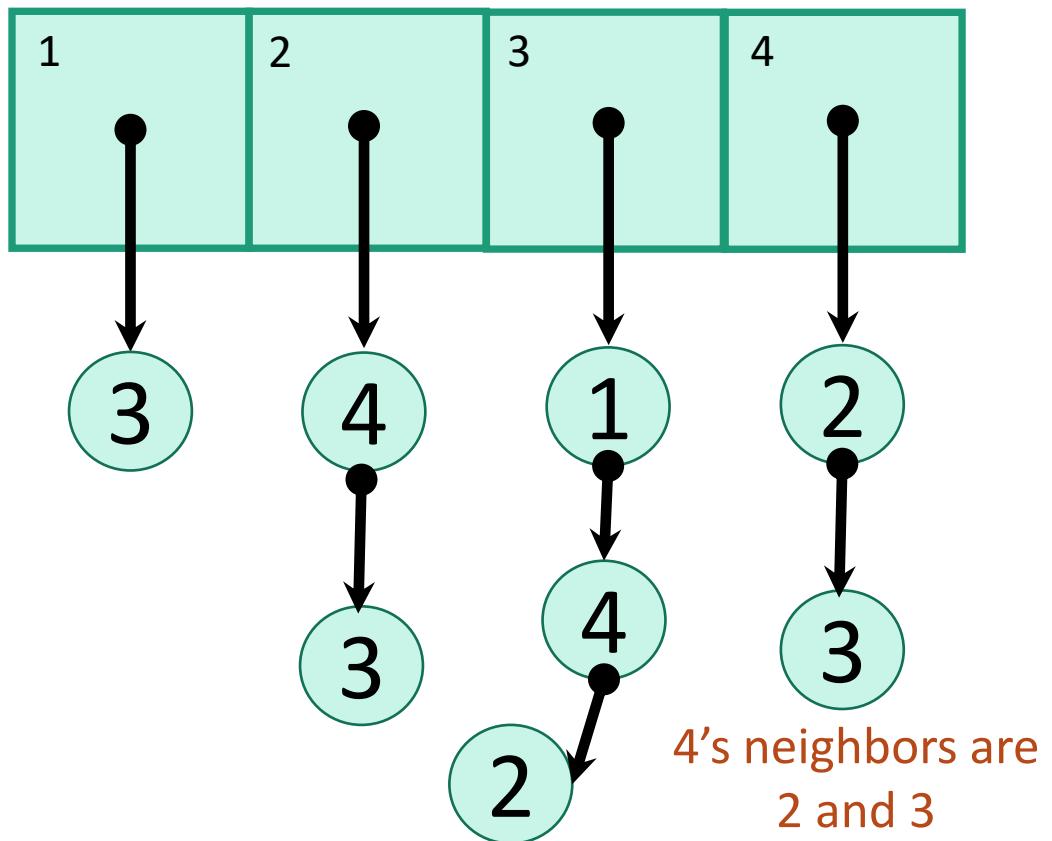
- Option 1: adjacency matrix

		Destination			
		1	2	3	4
Source	1	0	0	1	0
	2	0	0	0	1
3	0	1	0	1	
4	0	0	1	0	



How do we represent graphs?

- Option 2: linked lists.



How would you
modify this for
directed graphs?



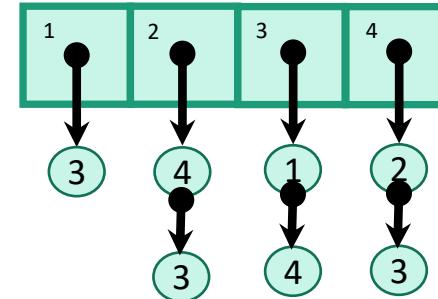
In either case

- May think of vertices storing other information
 - Attributes (name, IP address, ...)
 - helper info for algorithms that we will perform on the graph
- We will want to be able to do the following ops:
 - Edge Membership: Is edge e in E?
 - Neighbor Query: What are the neighbors of vertex v?

Trade-offs

Say there are n vertices
and m edges.

$$\begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$



Edge membership
Is $e = \{v,w\}$ in E ?

$O(1)$

$O(\deg(v))$ or
 $O(\deg(w))$

Neighbor query
Give me v 's neighbors.

$O(n)$

$O(\deg(v))$

Space requirements

$O(n^2)$

$O(n + m)$

We'll assume this
representation for
the rest of the class

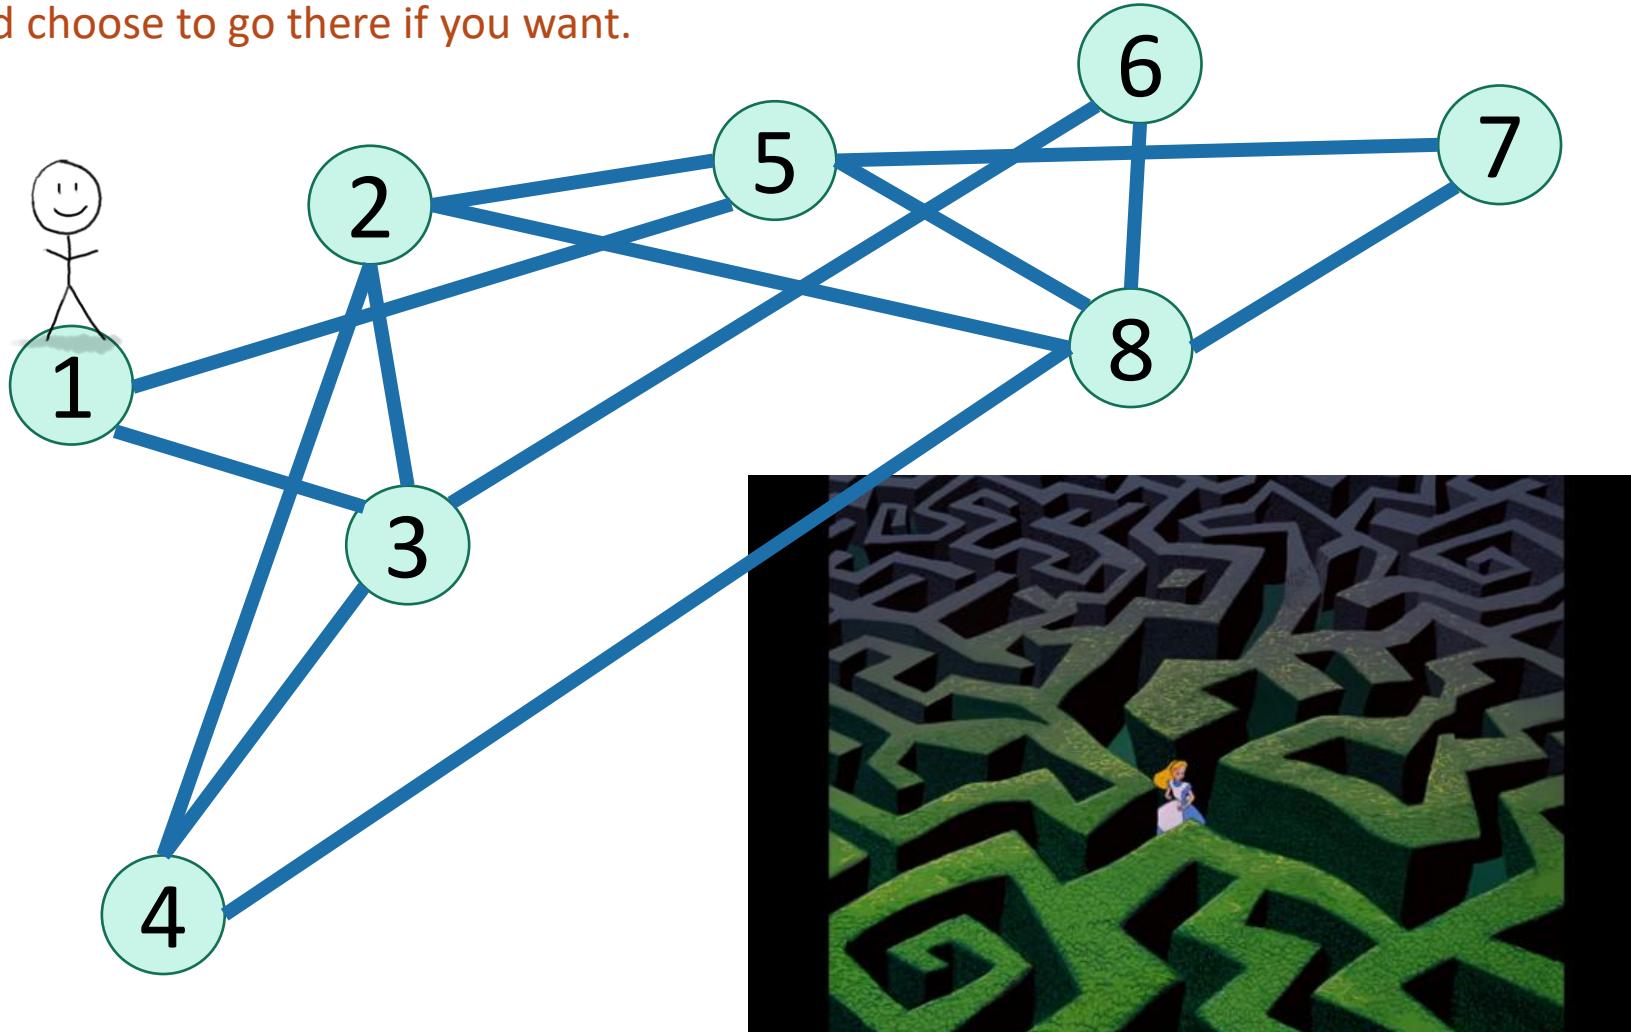
Assignment Question 1

- Design a Graph Class with following functions
 - EdgeMemberShip
 - NeighbourQuery
 - Add Edge
 - Constructor(List of Vertices)

Depth-first search

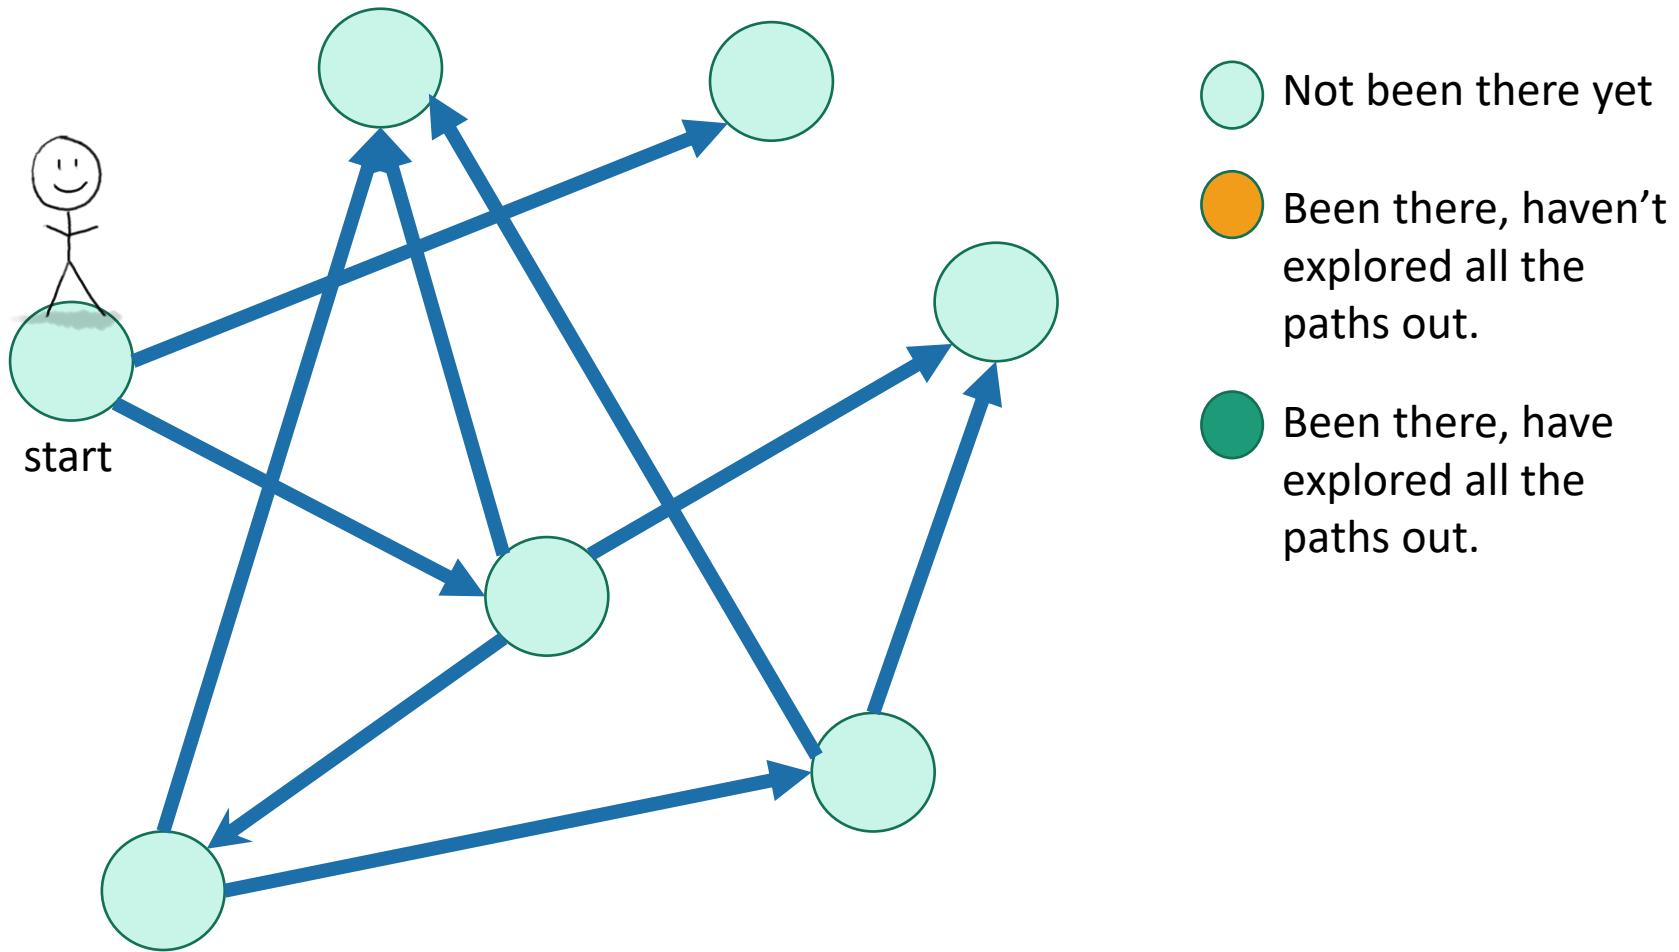
How do we explore a graph?

At each node, you can get a list of neighbors,
and choose to go there if you want.



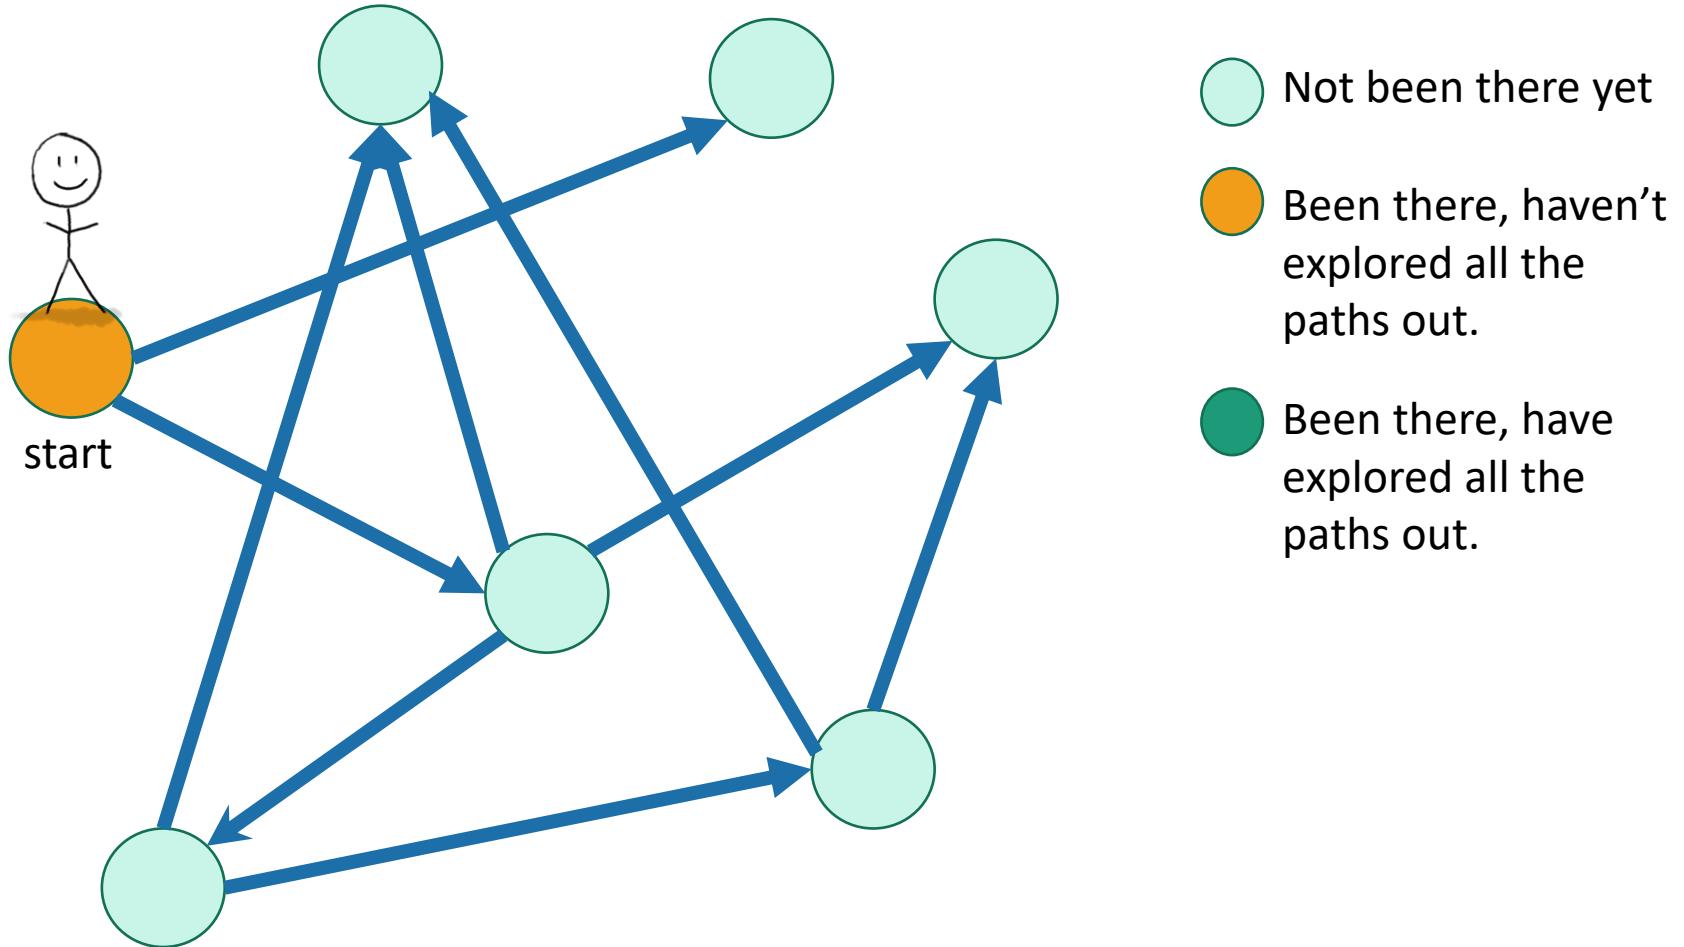
Depth First Search

Exploring a labyrinth with chalk and a piece of string



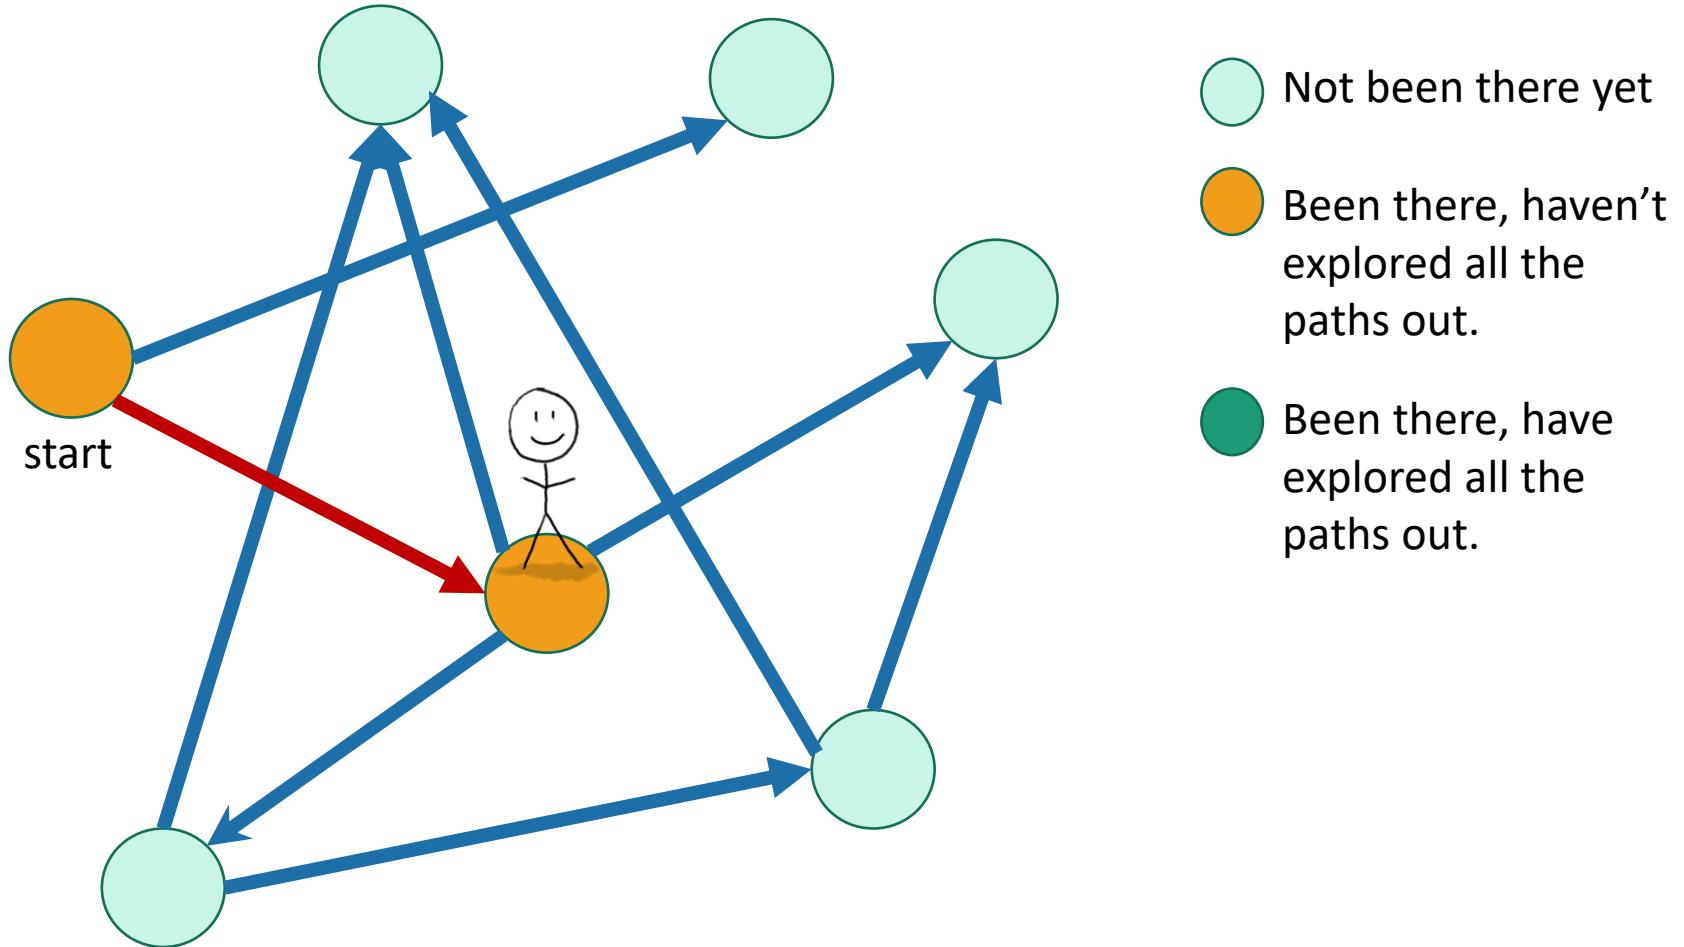
Depth First Search

Exploring a labyrinth with chalk and a piece of string



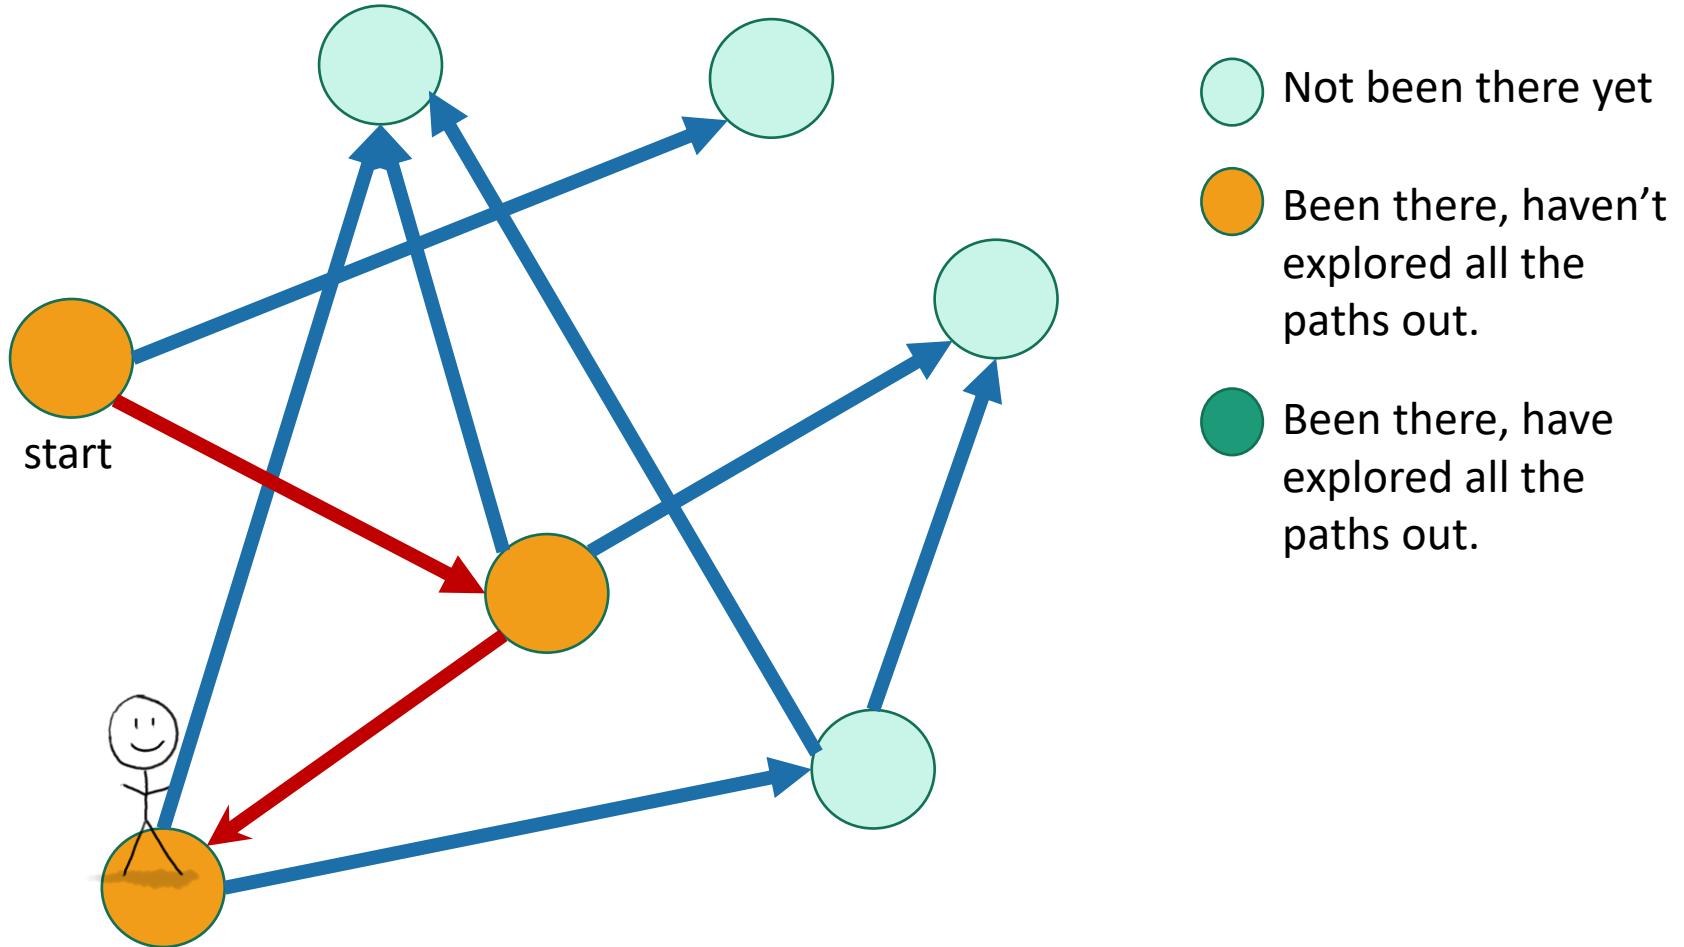
Depth First Search

Exploring a labyrinth with chalk and a piece of string



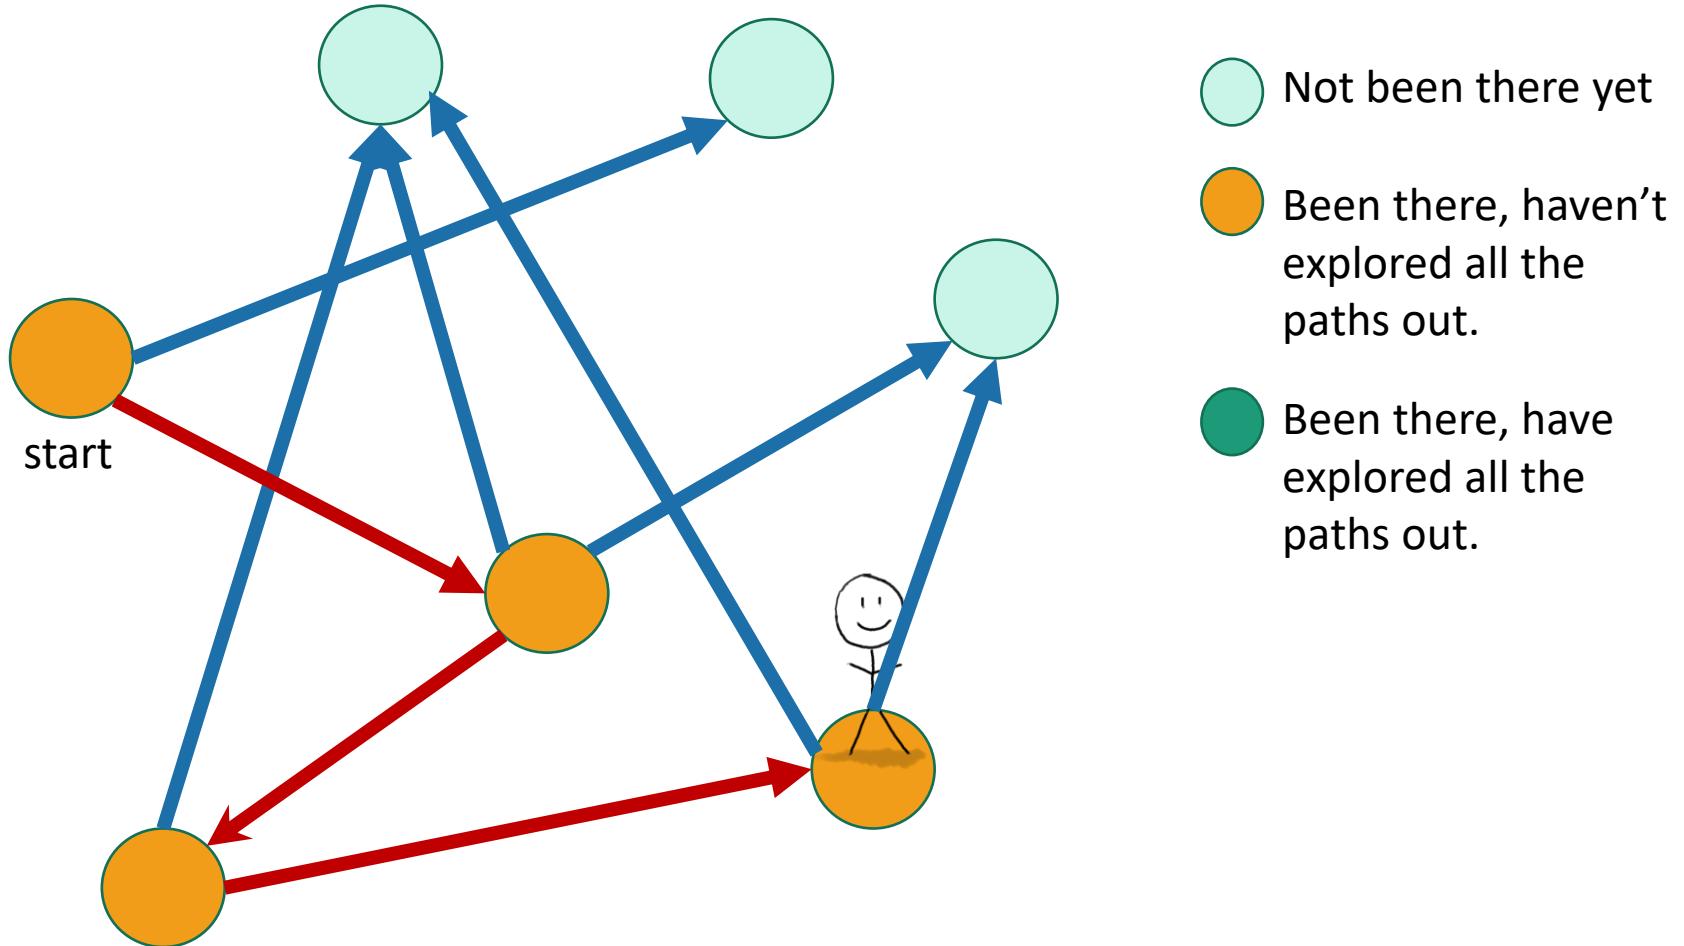
Depth First Search

Exploring a labyrinth with chalk and a piece of string



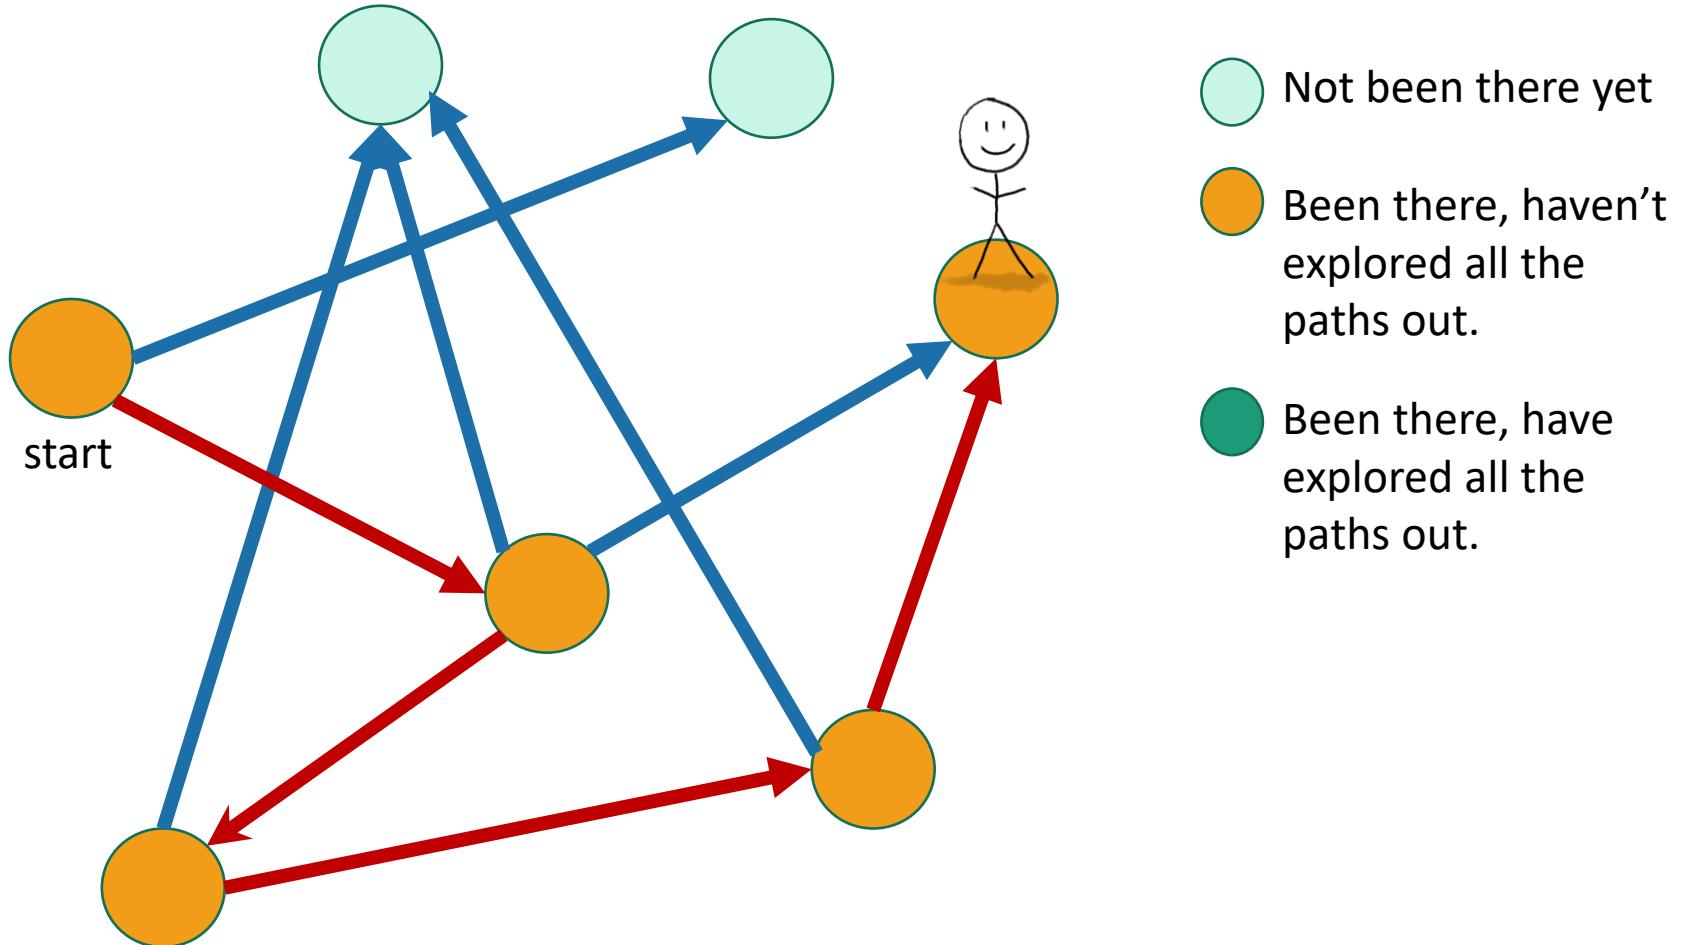
Depth First Search

Exploring a labyrinth with chalk and a piece of string



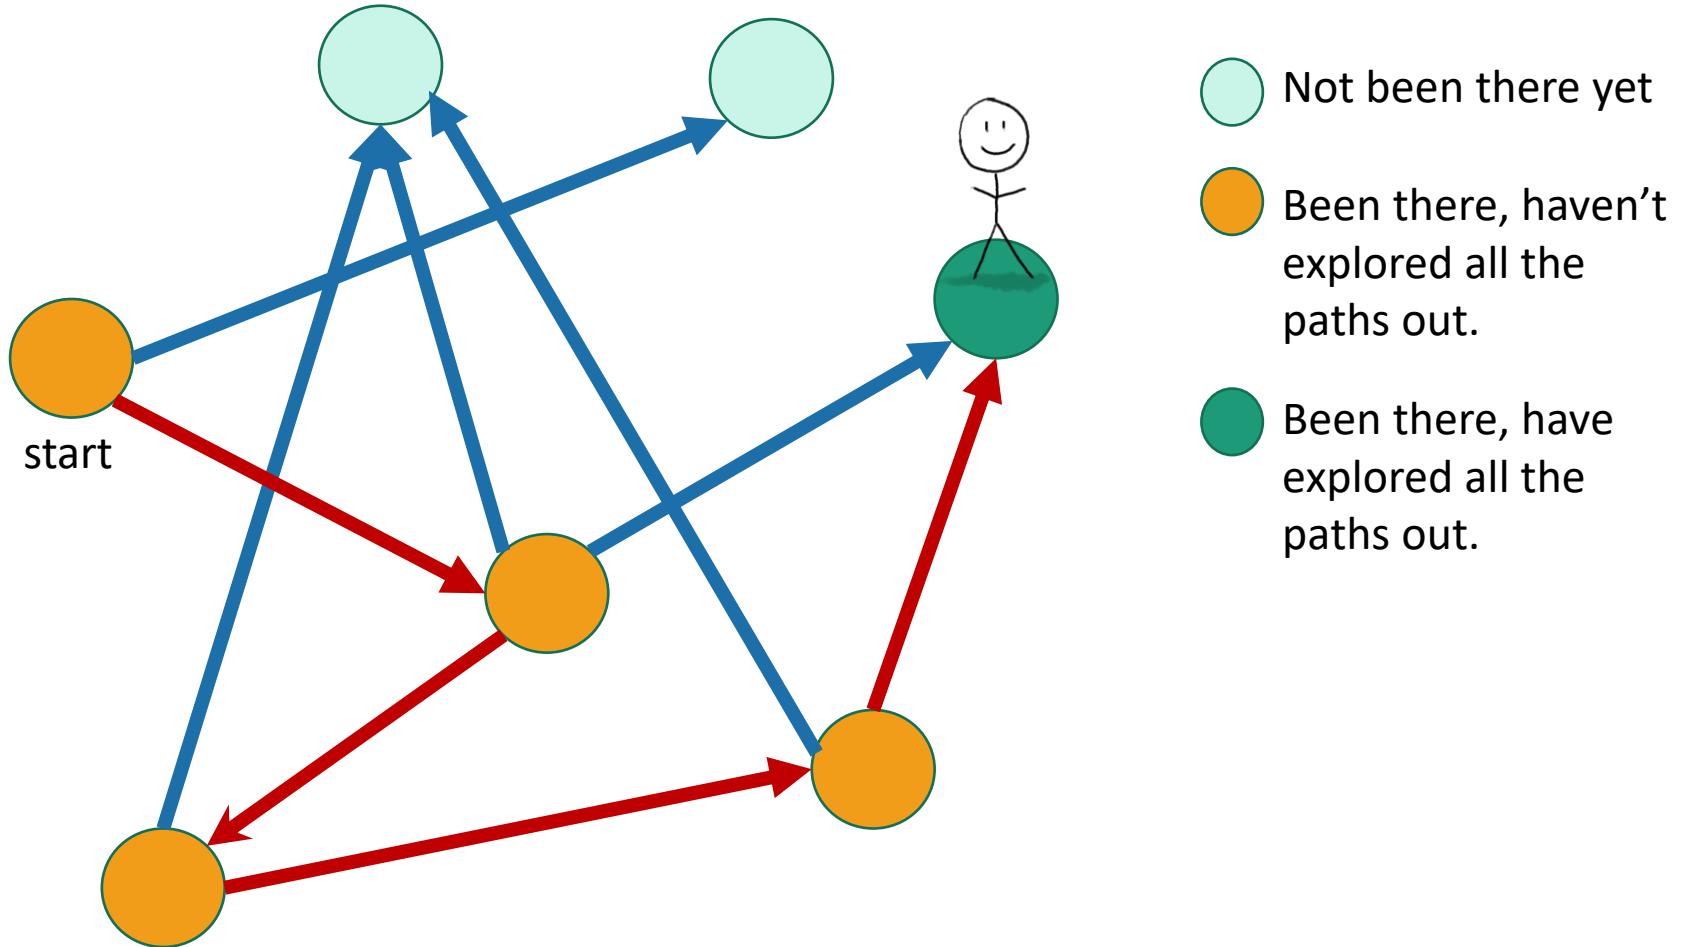
Depth First Search

Exploring a labyrinth with chalk and a piece of string



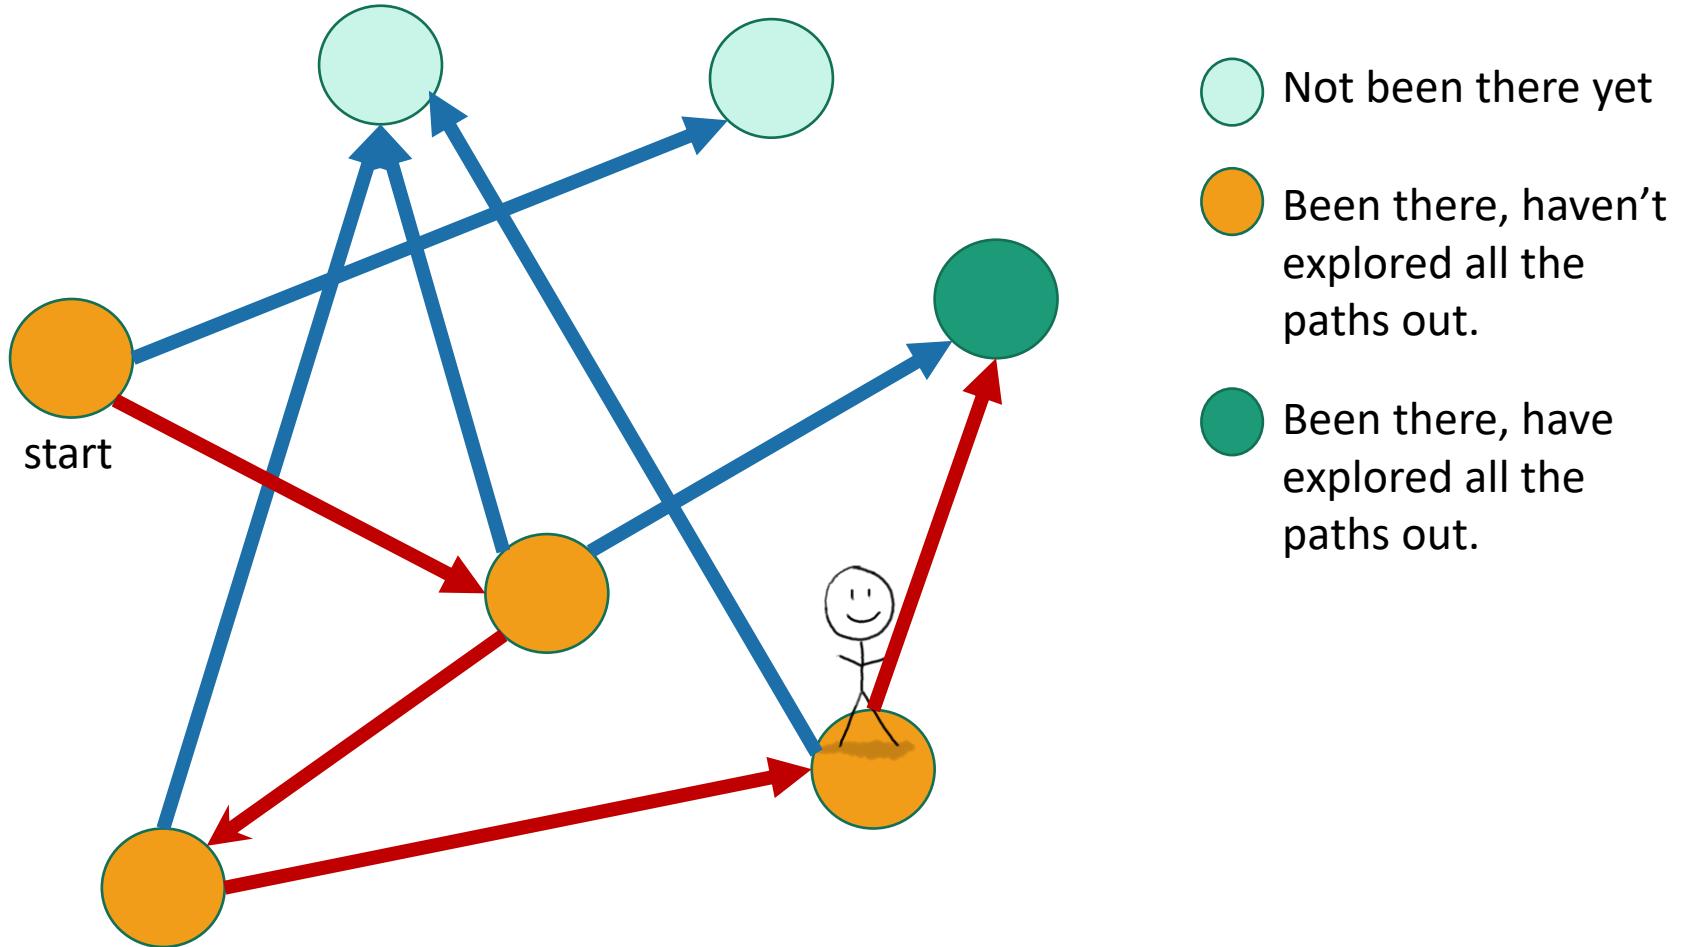
Depth First Search

Exploring a labyrinth with chalk and a piece of string



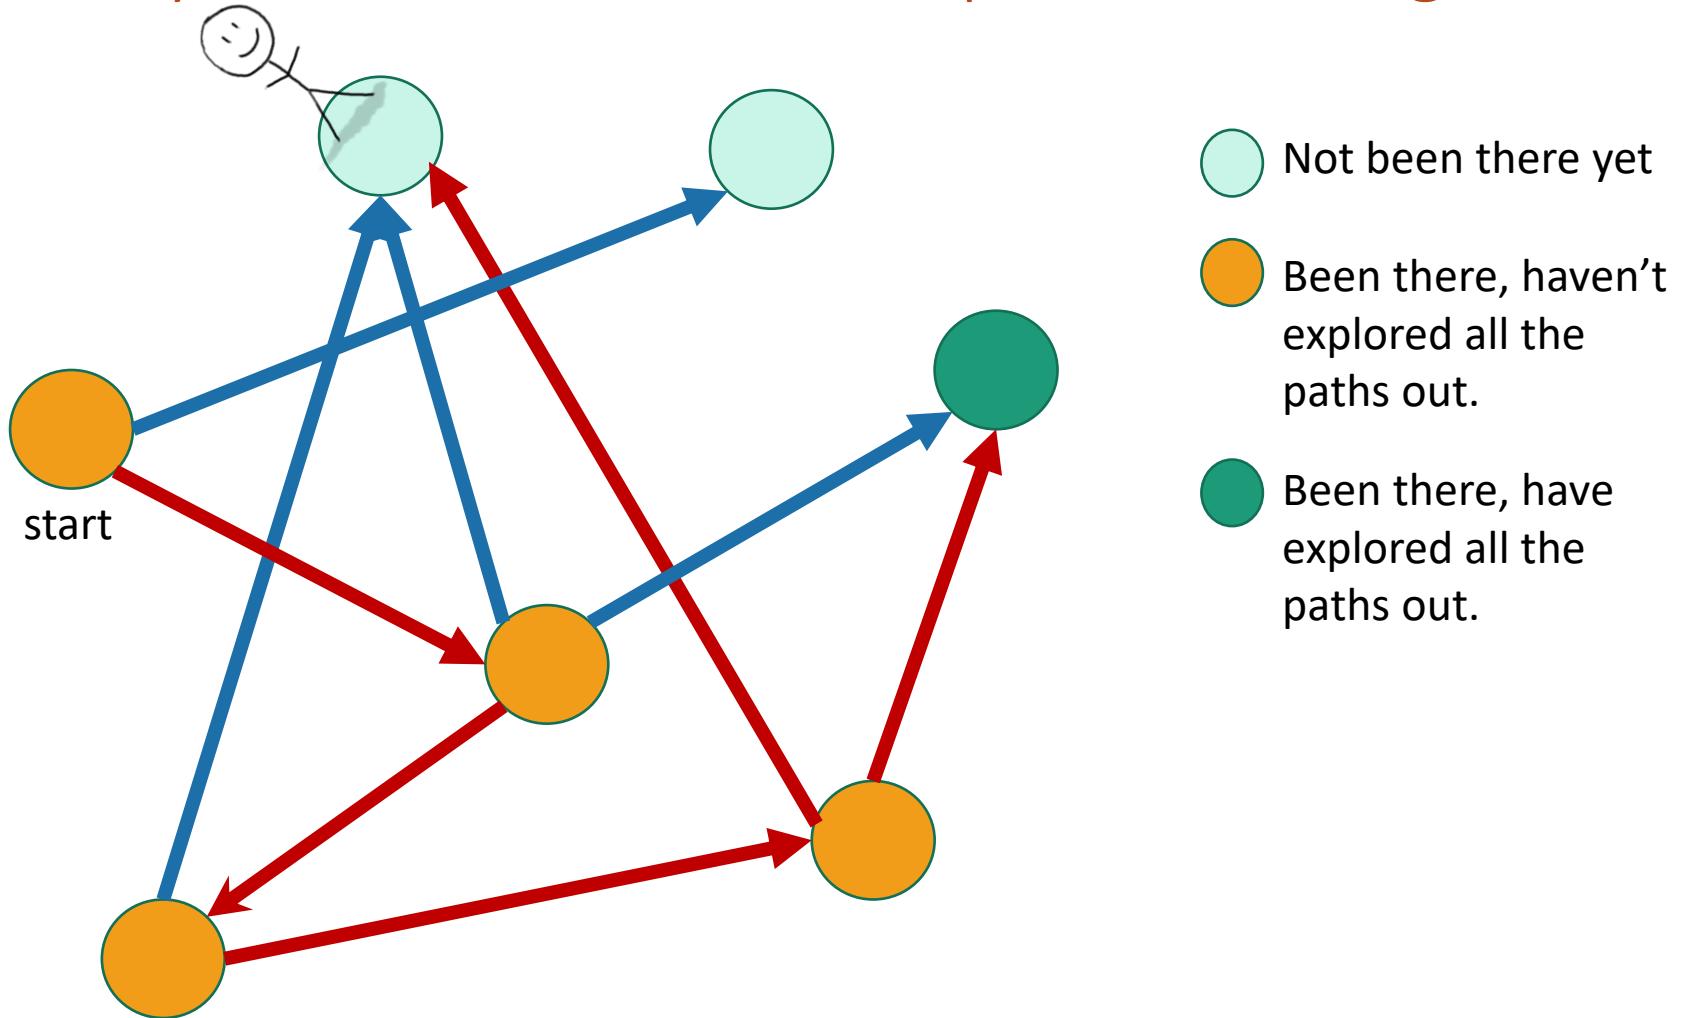
Depth First Search

Exploring a labyrinth with chalk and a piece of string



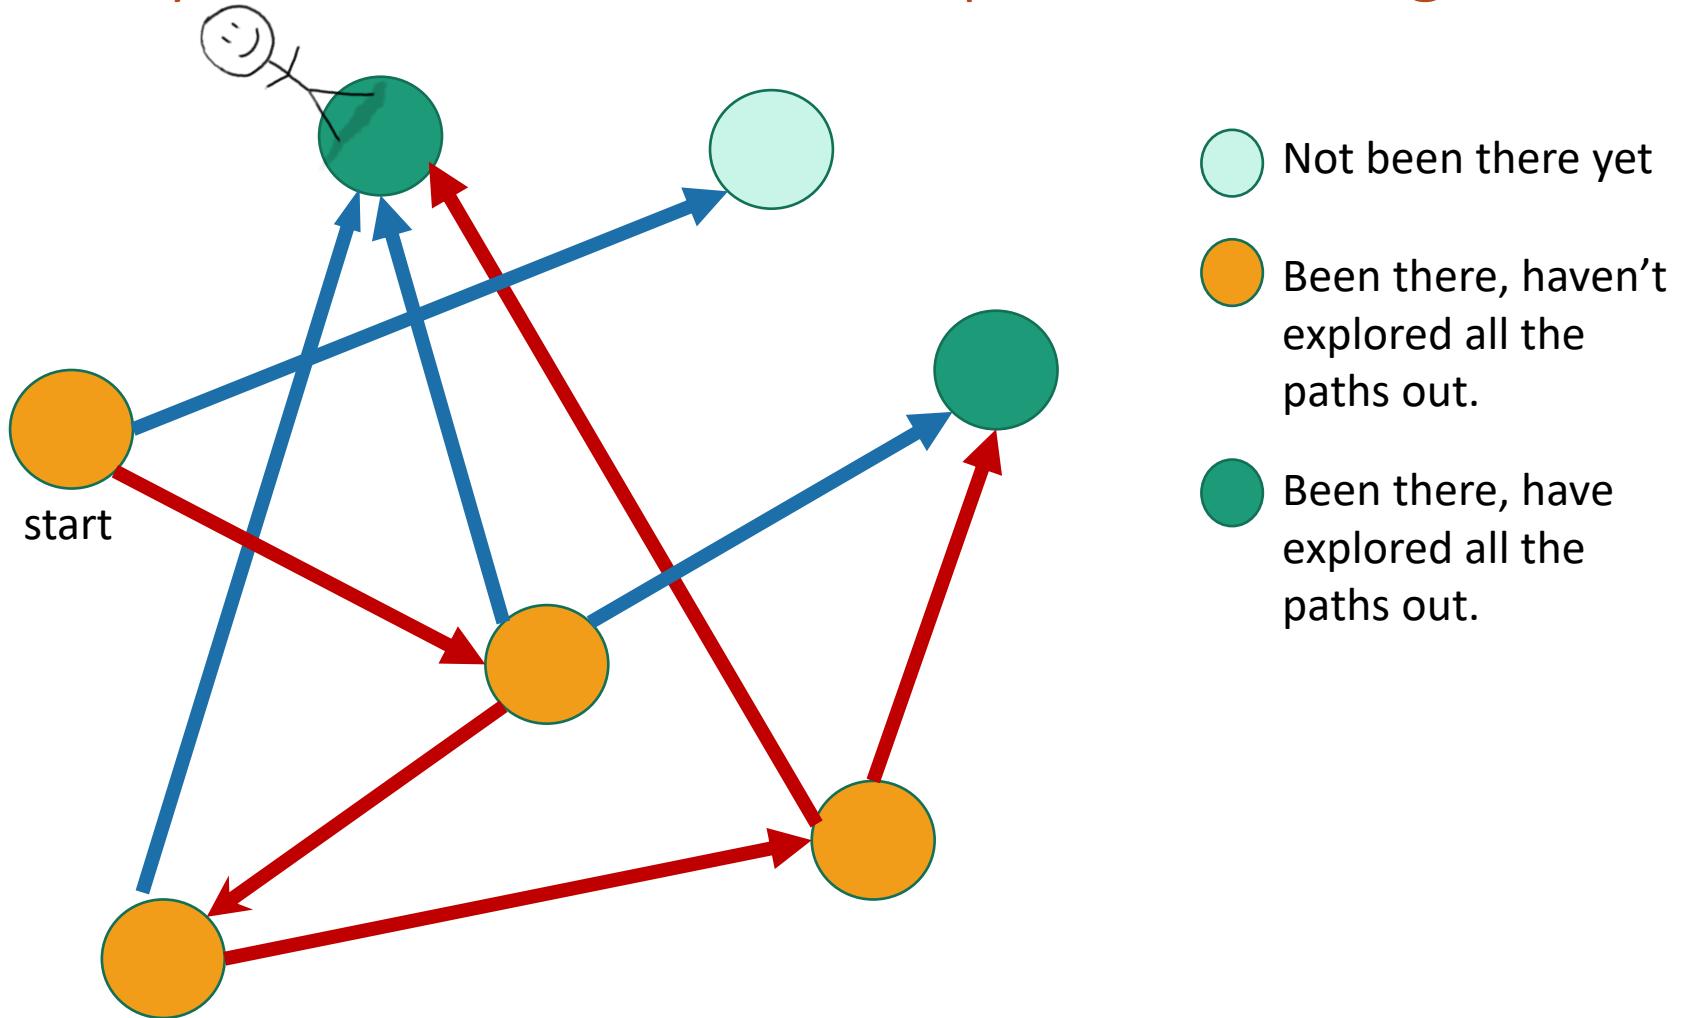
Depth First Search

Exploring a labyrinth with chalk and a piece of string



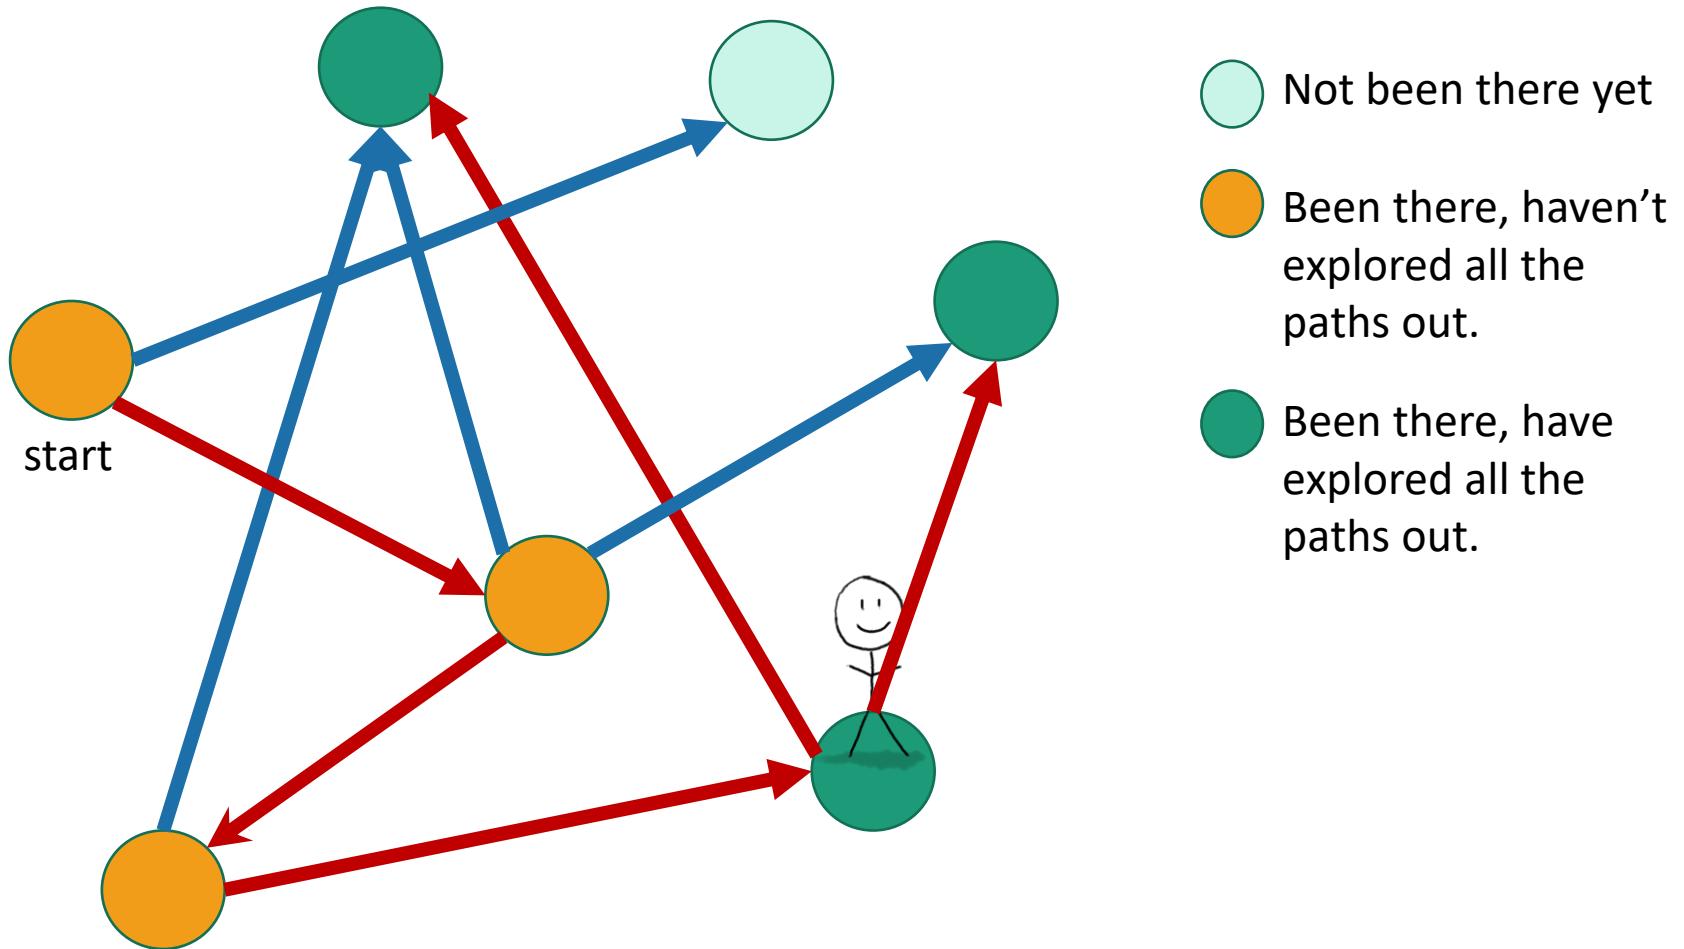
Depth First Search

Exploring a labyrinth with chalk and a piece of string



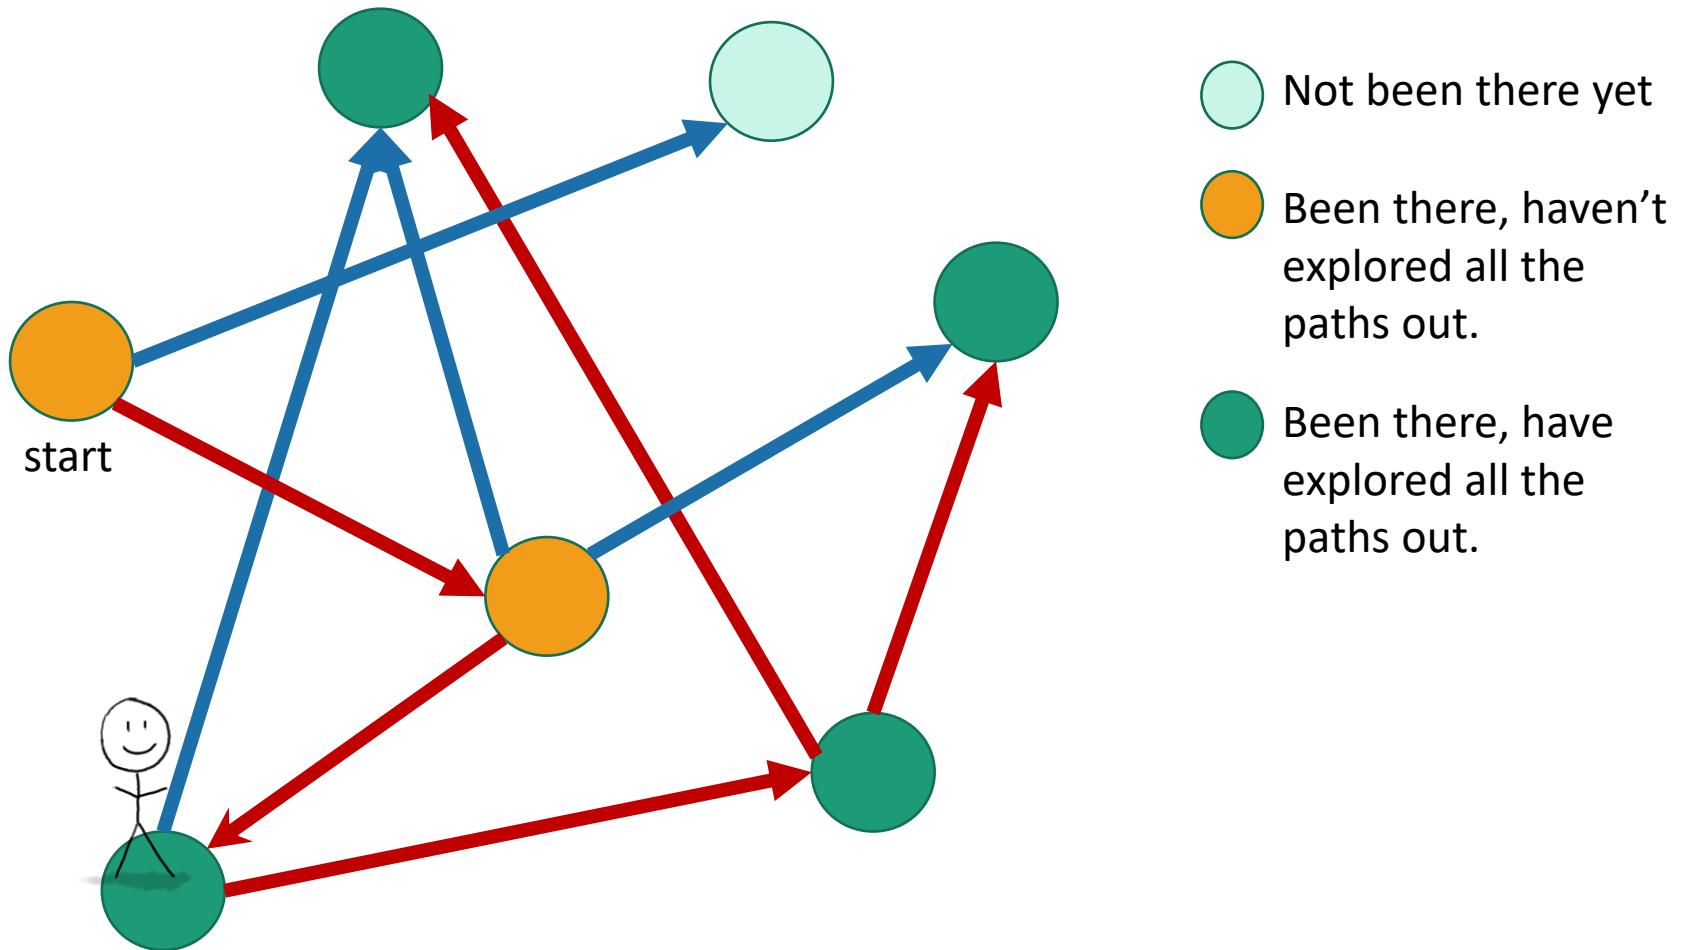
Depth First Search

Exploring a labyrinth with chalk and a piece of string



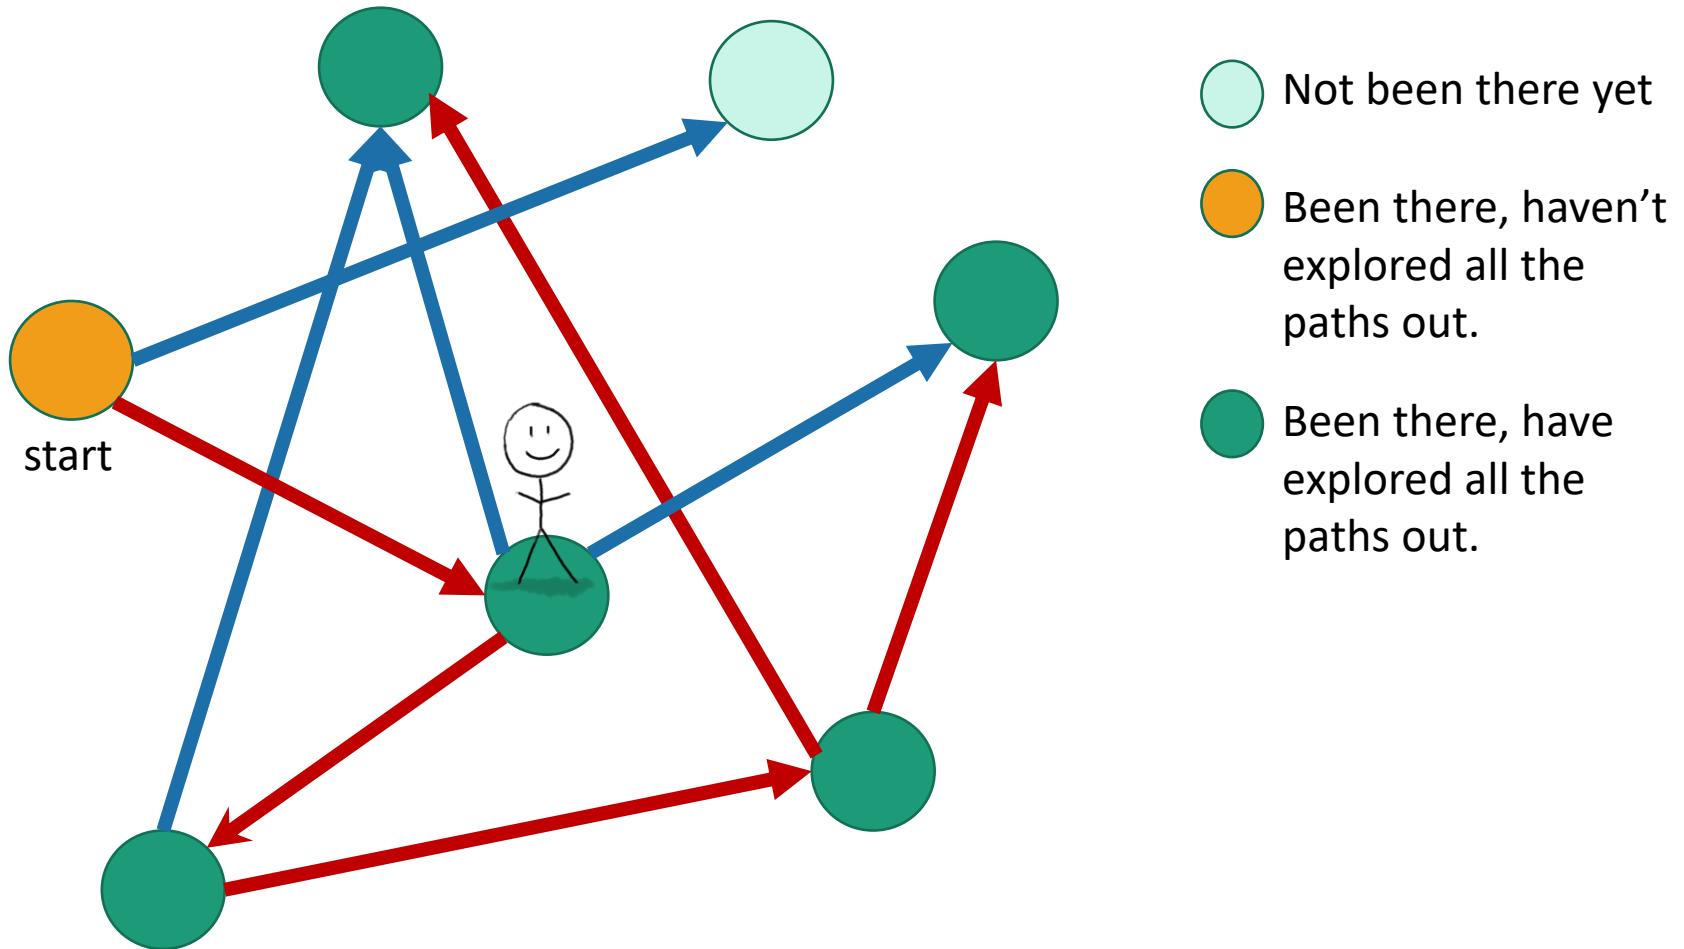
Depth First Search

Exploring a labyrinth with chalk and a piece of string



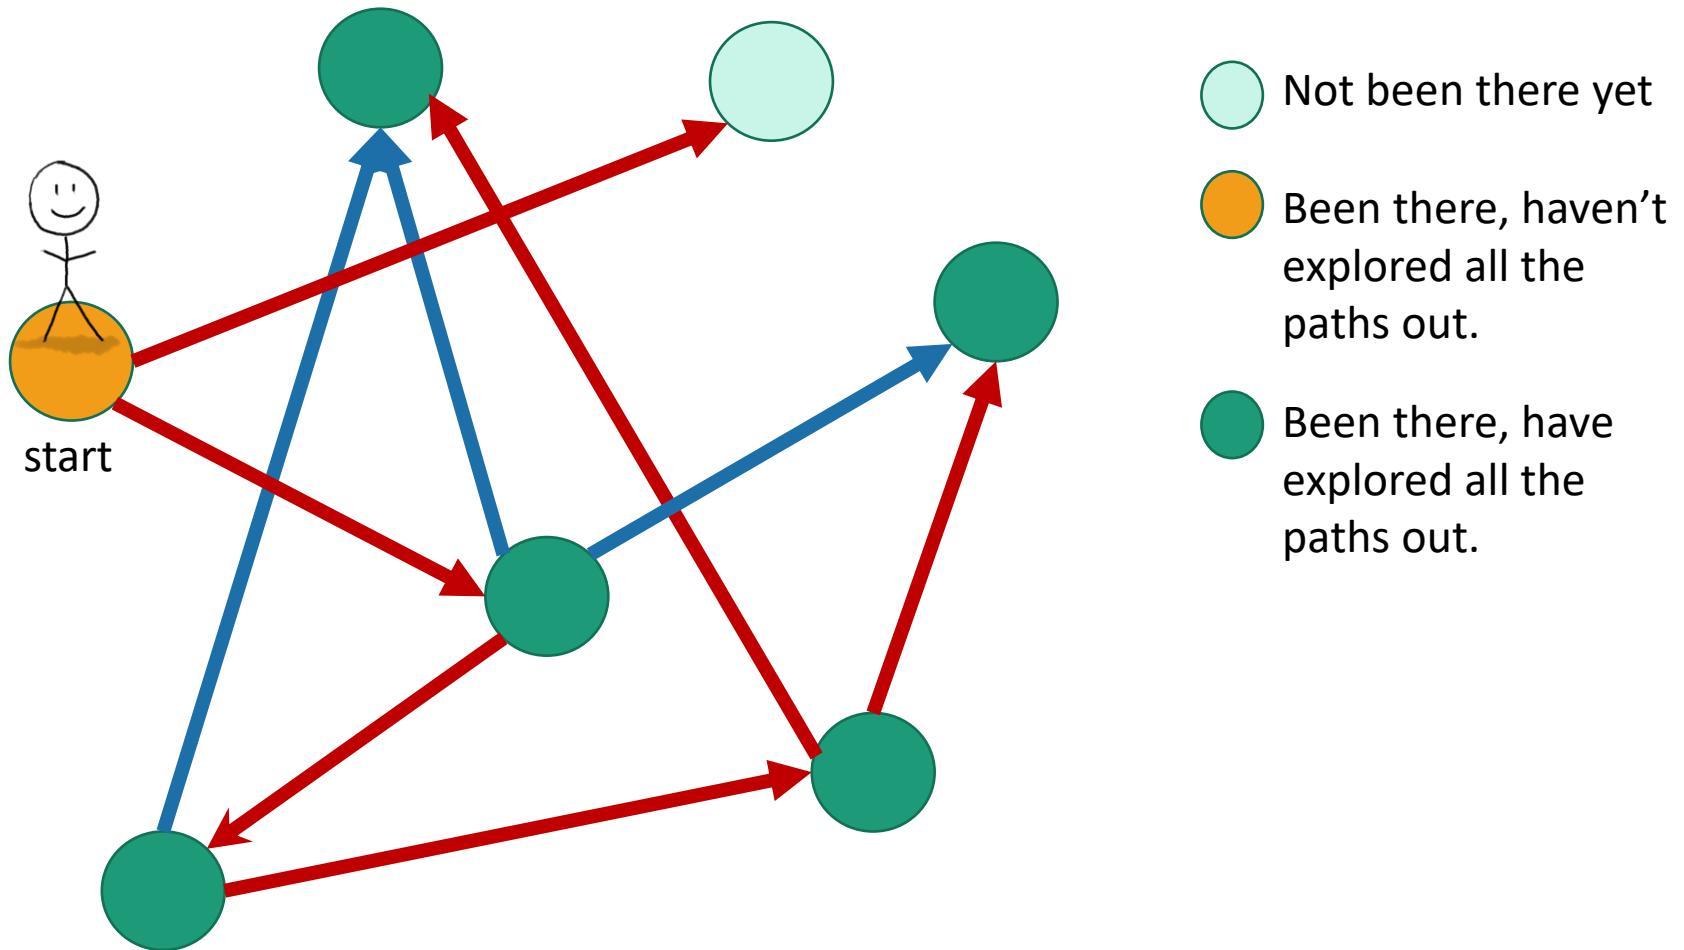
Depth First Search

Exploring a labyrinth with chalk and a piece of string



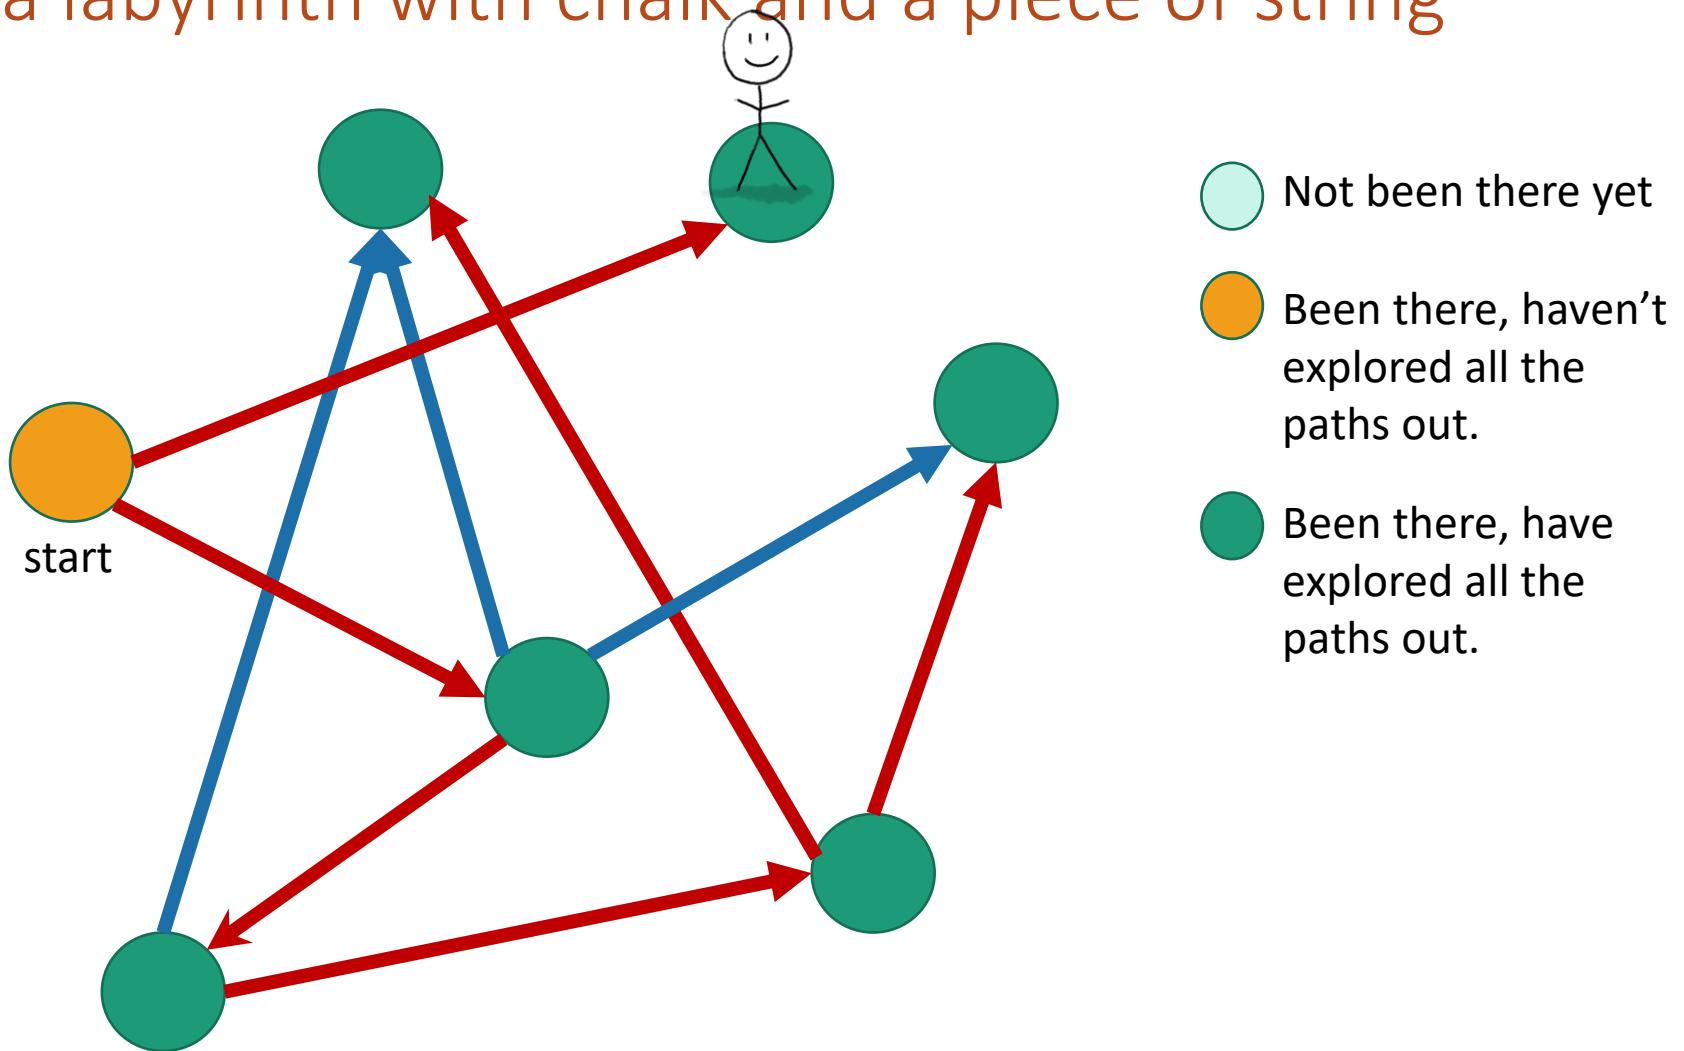
Depth First Search

Exploring a labyrinth with chalk and a piece of string



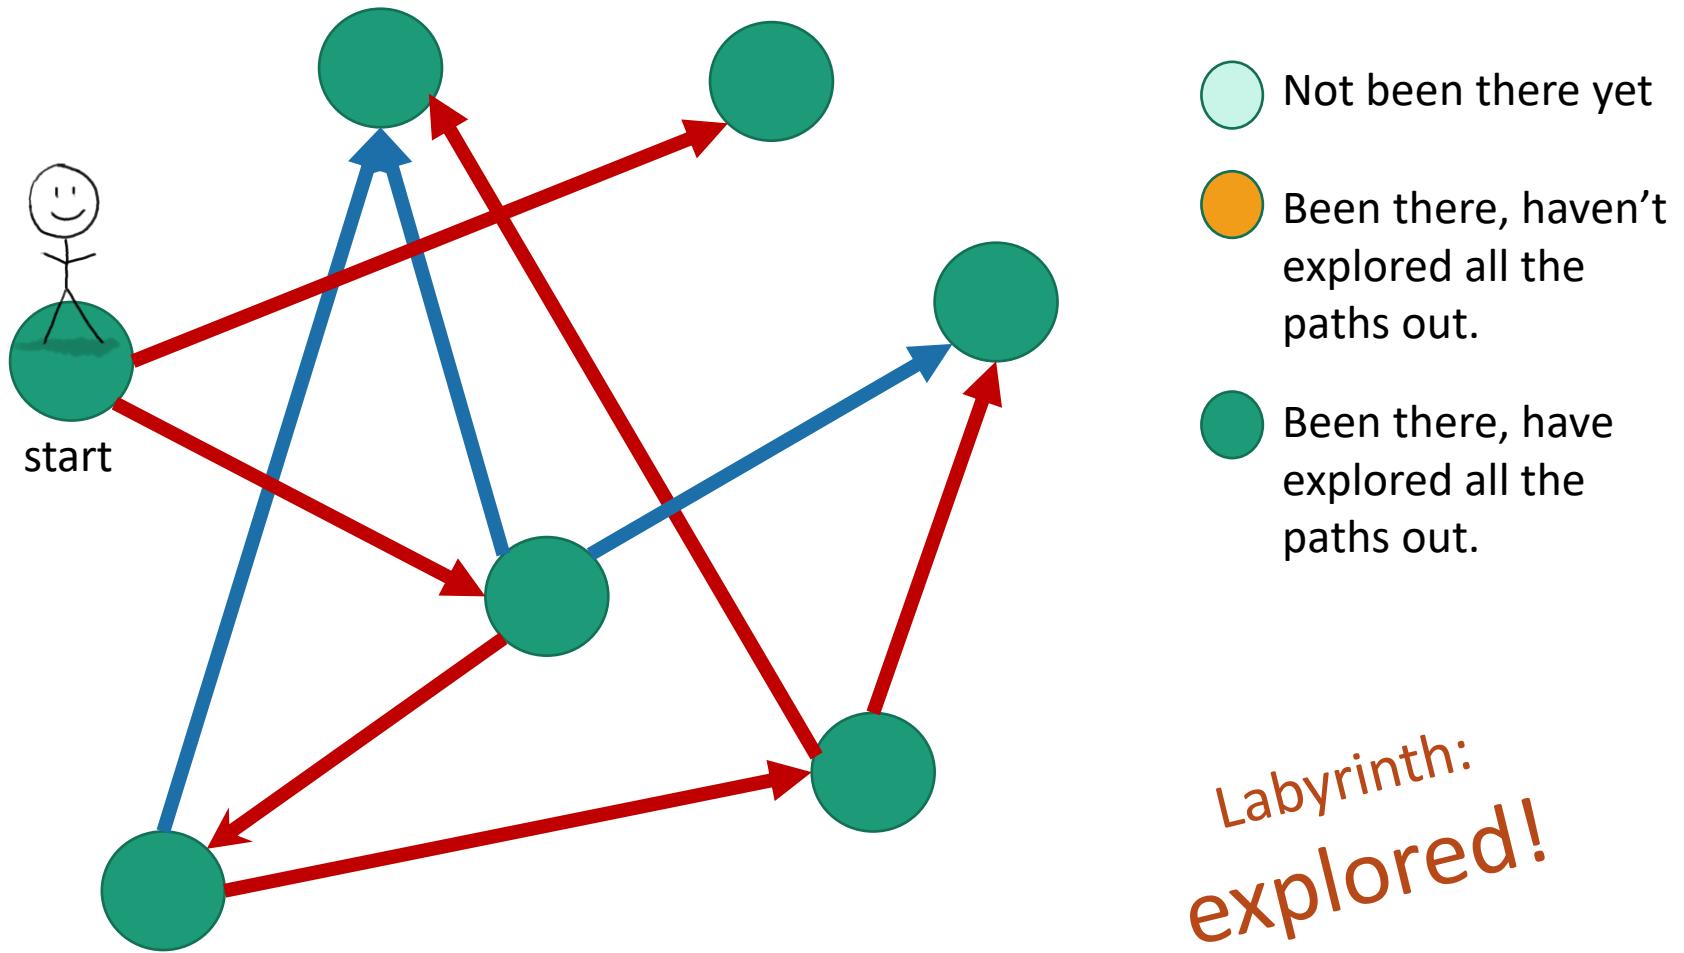
Depth First Search

Exploring a labyrinth with chalk and a piece of string



Depth First Search

Exploring a labyrinth with chalk and a piece of string



Depth First Search

Exploring a labyrinth with pseudocode

- Each vertex keeps track of whether it is:

- Unvisited
- In progress
- All done



- Each vertex will also keep track of:

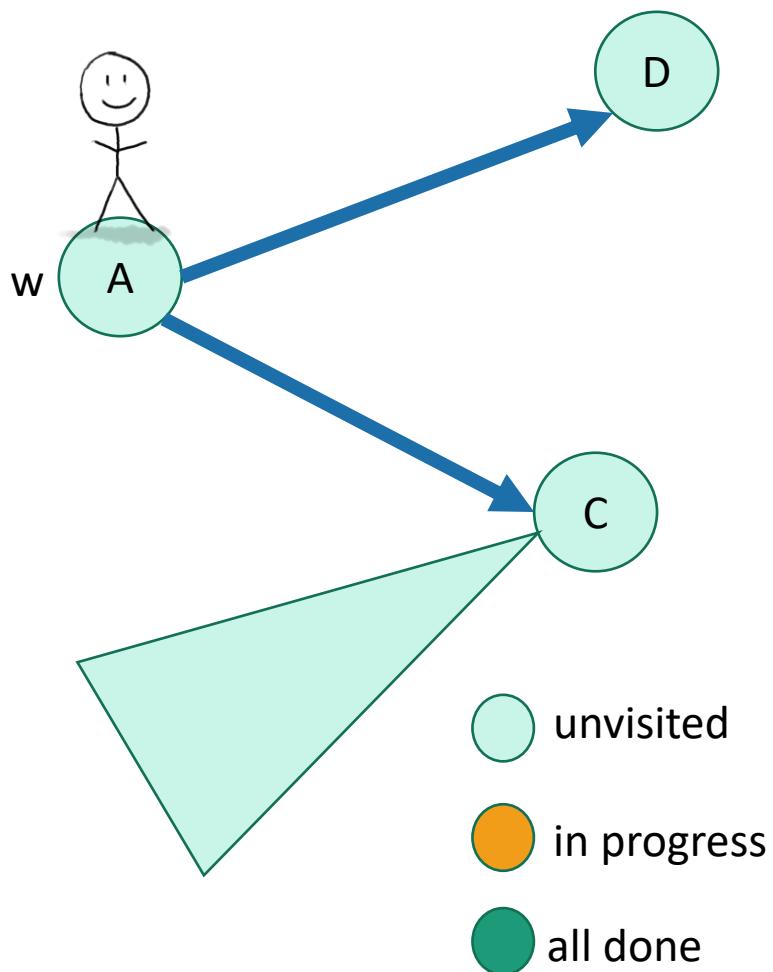
- The time we **first enter it**.
- The time we finish with it and mark it **all done**.



You might have seen other ways to implement DFS than what we are about to go through. This way has more bookkeeping, but more intuition – also, the bookkeeping will be useful later!

Depth First Search

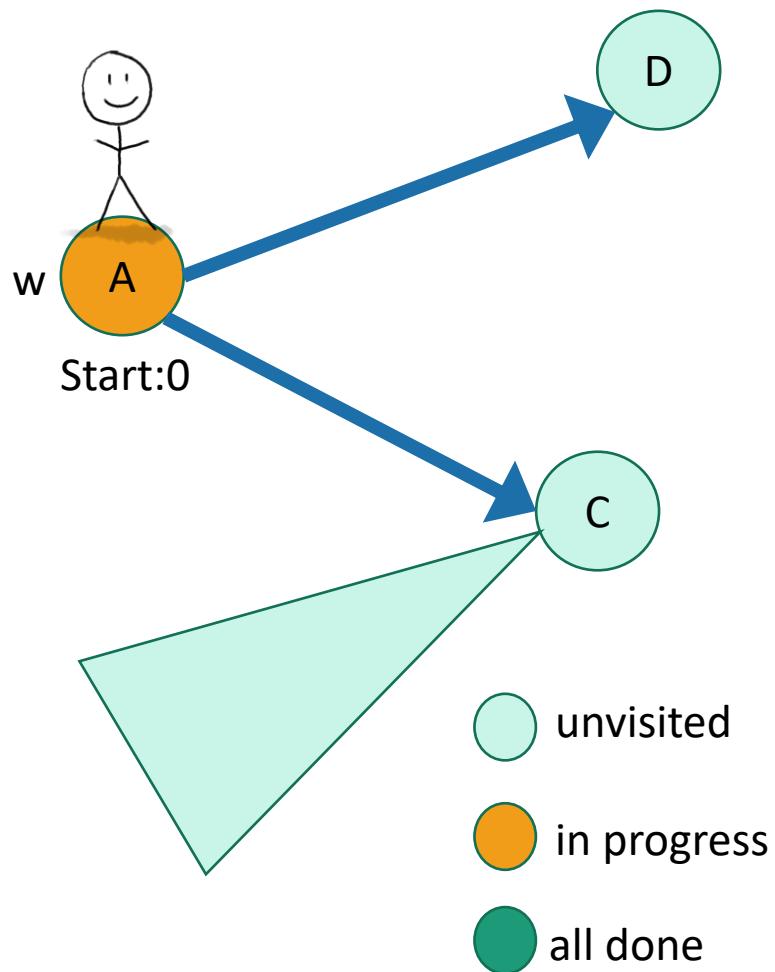
currentTime = 0



- **DFS(w, currentTime):**
 - w.entryTime = currentTime
 - currentTime ++
 - Mark w as **in progress**.
 - **for** v in w.neighbors:
 - **if** v is **unvisited**:
 - currentTime
= **DFS(v, currentTime)**
 - currentTime ++
 - w.finishTime = currentTime
 - Mark w as **all done**
 - **return** currentTime

Depth First Search

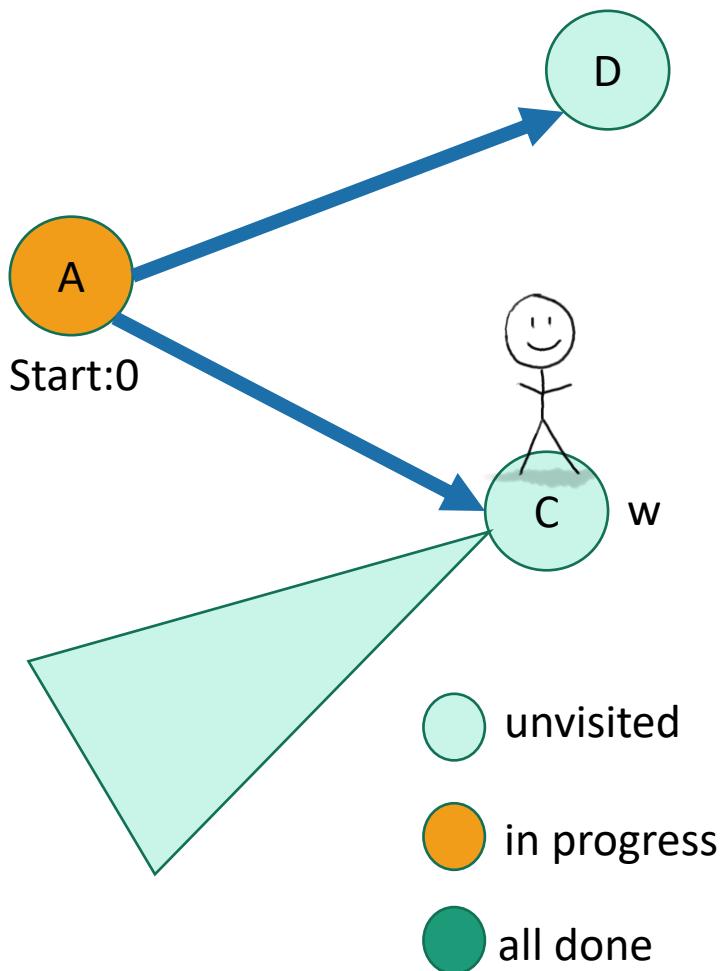
currentTime = 1



- **DFS(w, currentTime):**
 - `w.entryTime = currentTime`
 - `currentTime ++`
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - `currentTime`
= **DFS(v, currentTime)**
 - `currentTime ++`
 - `w.finishTime = currentTime`
 - Mark w as **all done**
 - **return** `currentTime`

Depth First Search

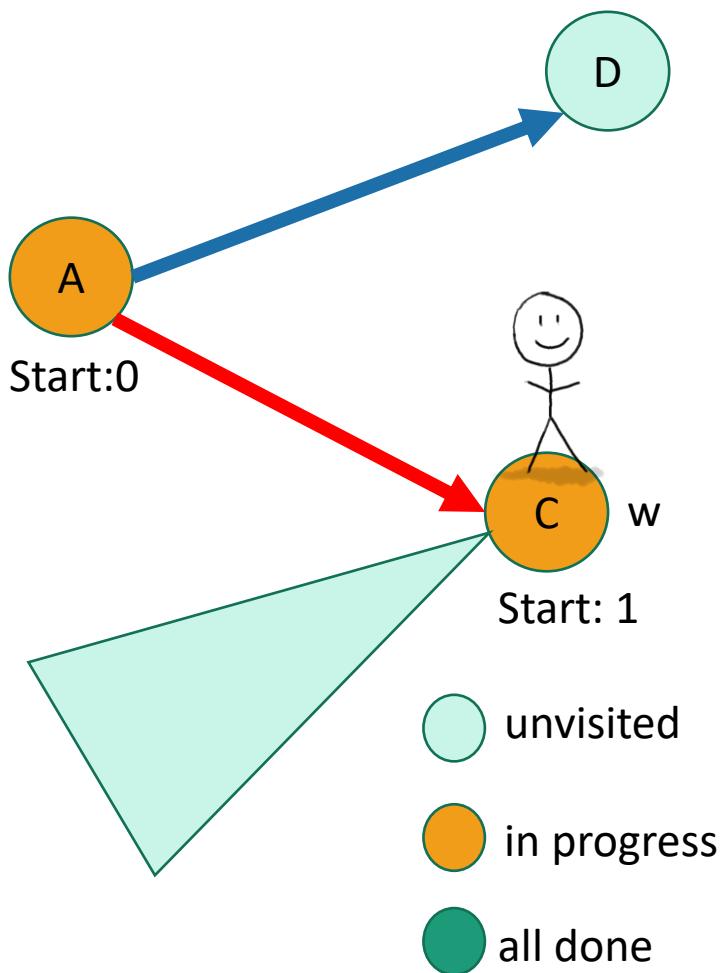
currentTime = 1



- **DFS(w, currentTime):**
 - `w.entryTime = currentTime`
 - `currentTime ++`
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - `currentTime`
= **DFS(v, currentTime)**
 - `currentTime ++`
 - `w.finishTime = currentTime`
 - Mark w as **all done**
 - **return currentTime**

Depth First Search

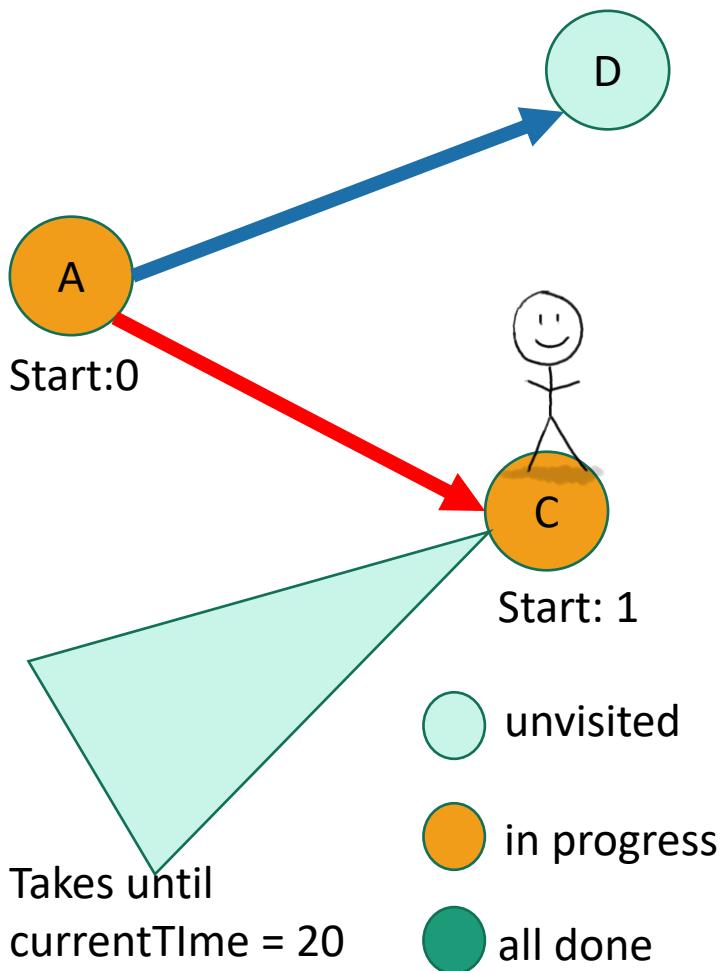
currentTime = 2



- **DFS(w, currentTime):**
 - w.entryTime = currentTime
 - currentTime ++
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - currentTime
= **DFS(v, currentTime)**
 - currentTime ++
 - w.finishTime = currentTime
 - Mark w as **all done**
 - **return** currentTime

Depth First Search

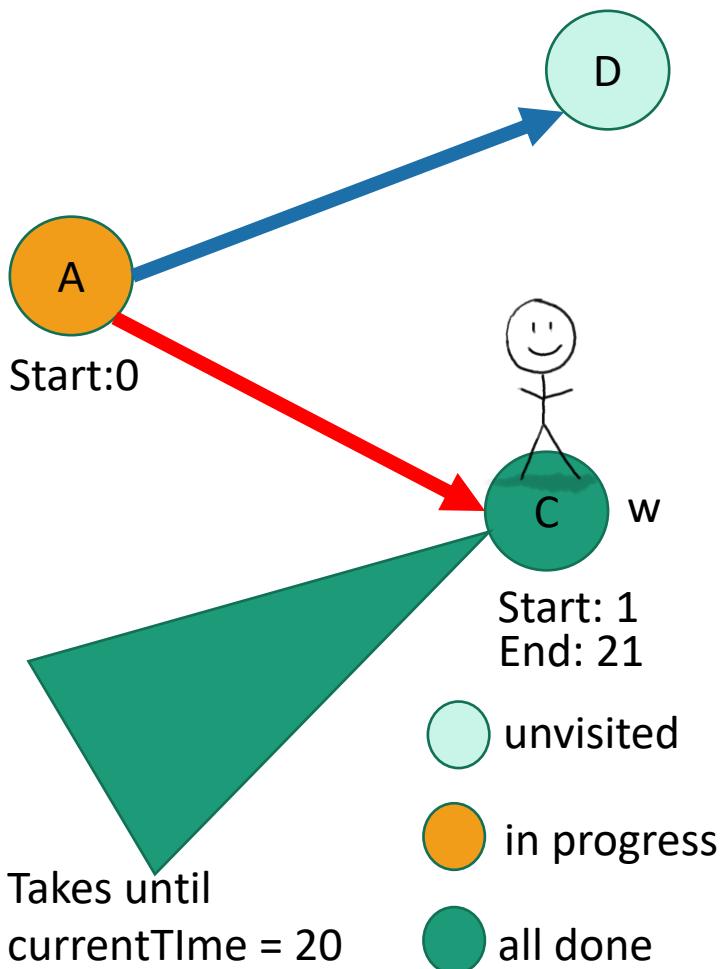
currentTime = 20



- **DFS(w, currentTime):**
 - `w.entryTime = currentTime`
 - `currentTime ++`
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - `currentTime`
`= DFS(v, currentTime)`
 - `currentTime ++`
 - `w.finishTime = currentTime`
 - Mark w as **all done**
 - **return currentTime**

Depth First Search

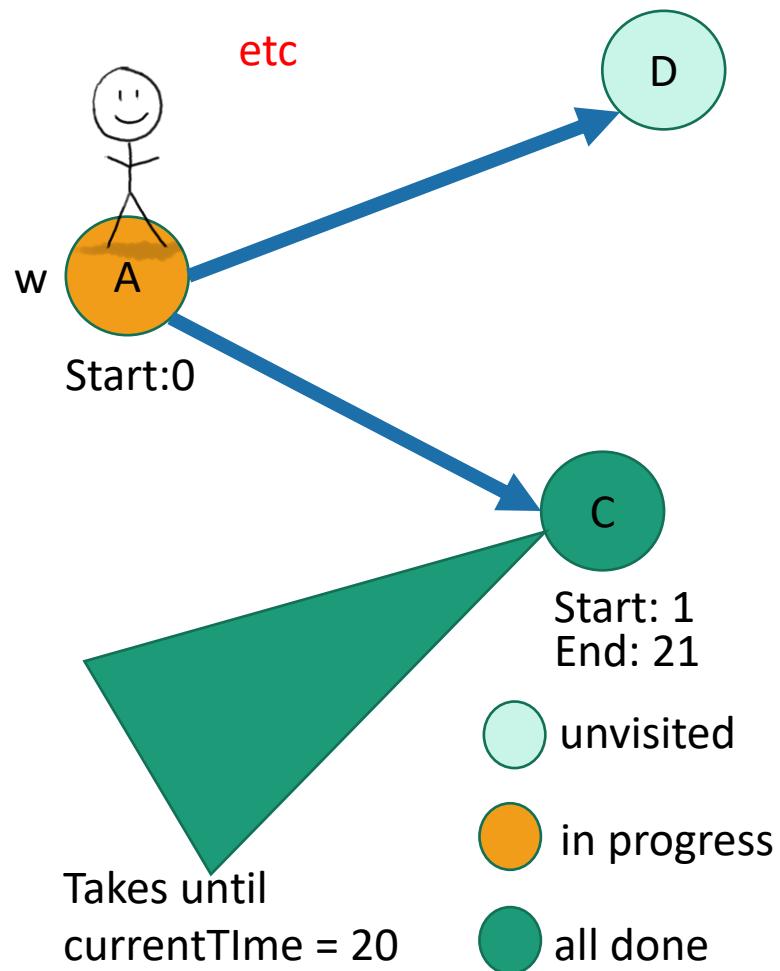
currentTime = 21



- **DFS(w, currentTime):**
 - w.startTime = currentTime
 - currentTime ++
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - currentTime
= **DFS(v, currentTime)**
 - currentTime ++
 - w.finishTime = currentTime
 - Mark w as **all done**
 - **return** currentTime

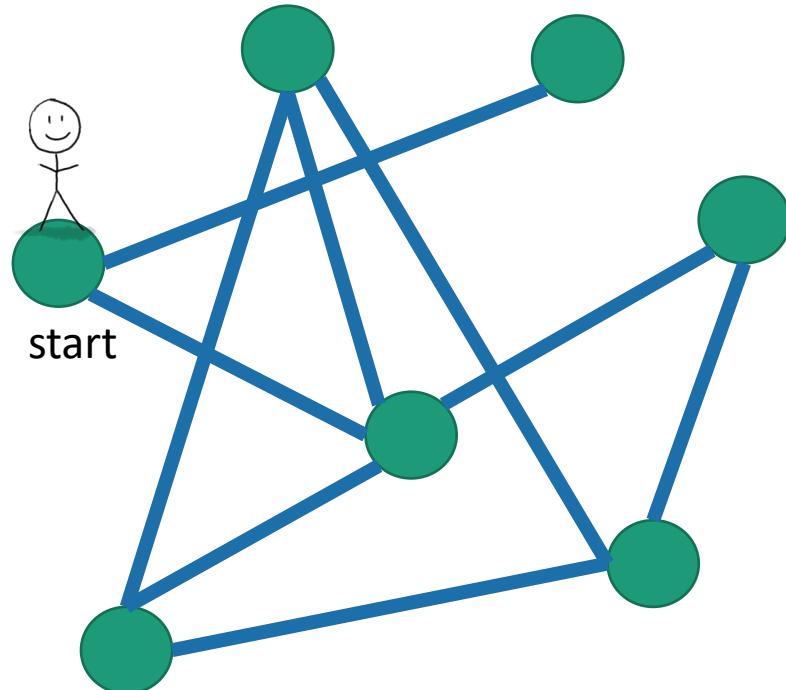
Depth First Search

currentTime = 21

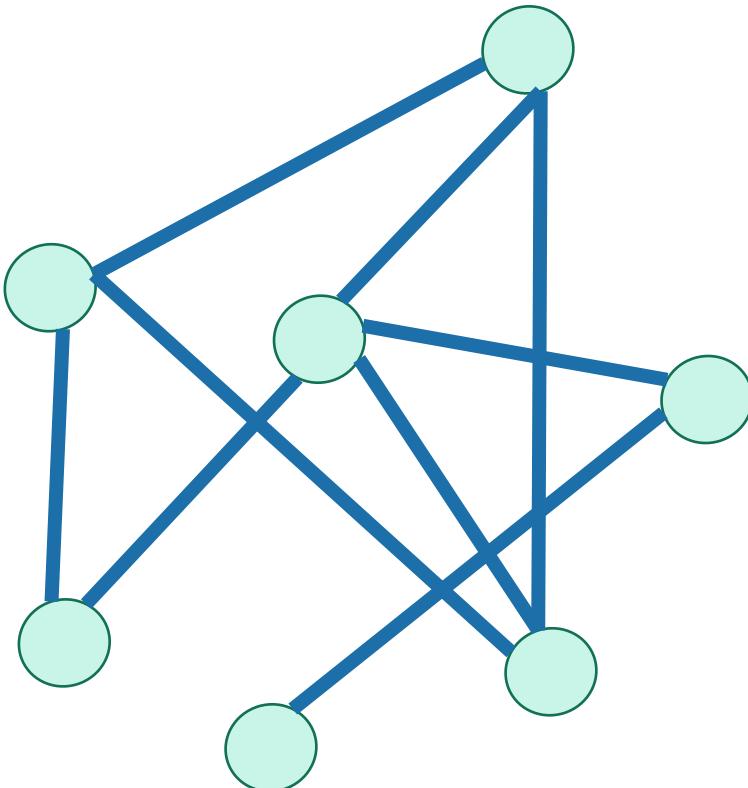


- **DFS(w, currentTime):**
 - `w.startTime = currentTime`
 - `currentTime ++`
 - Mark w as **in progress**.
 - **for v in w.neighbors:**
 - **if v is unvisited:**
 - `currentTime`
= **DFS(v, currentTime)**
 - `currentTime ++`
 - `w.finishTime = currentTime`
 - Mark w as **all done**
 - **return currentTime**

DFS finds all the nodes reachable from the starting point



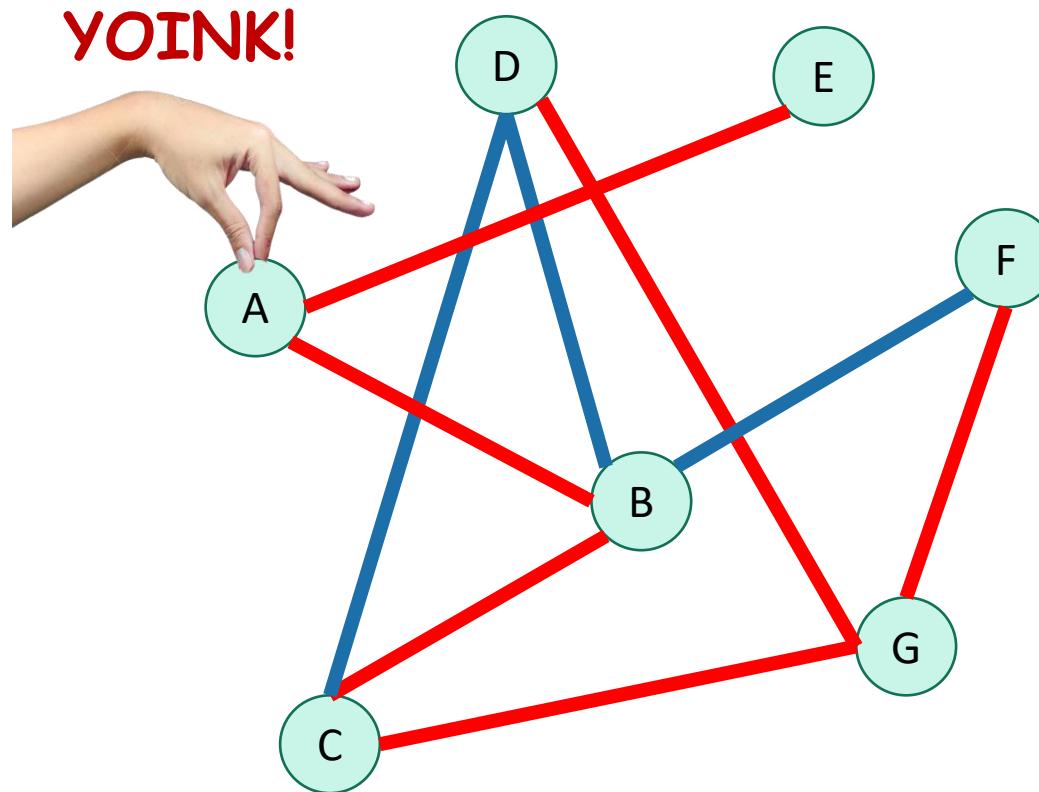
In an undirected graph, this is called a **connected component**.



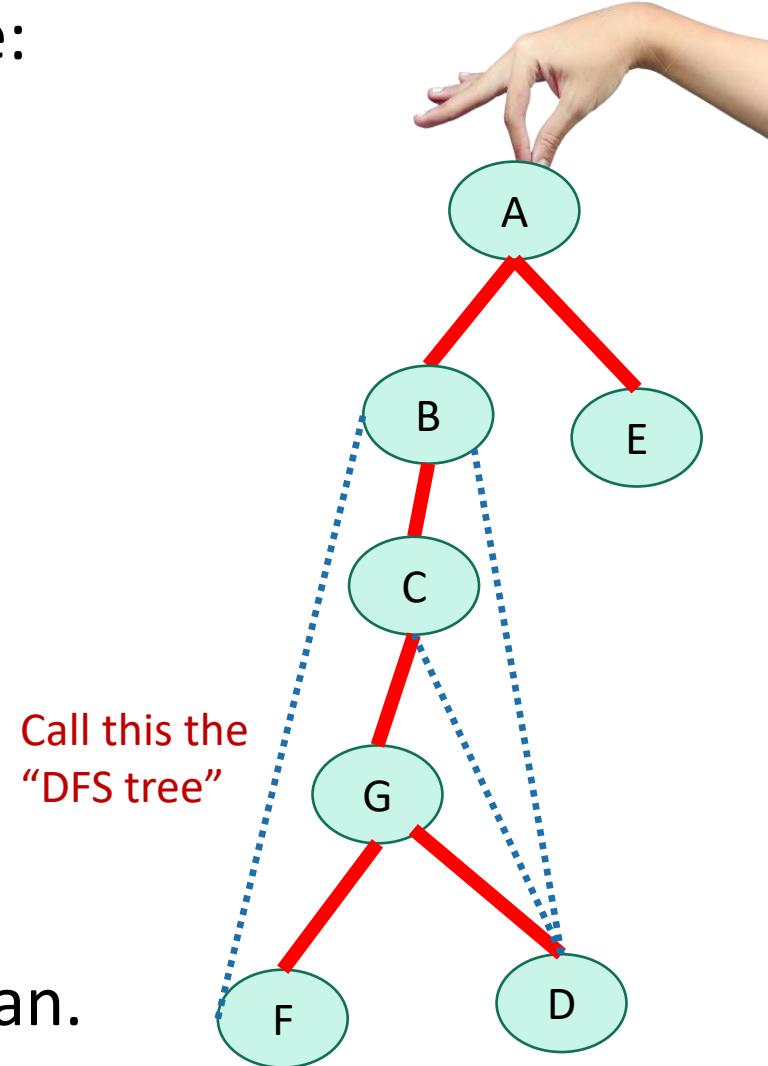
One application: finding connected components.

Why is it called depth-first?

- We are implicitly building a tree:



- And first we go as deep as we can.



Running time

To explore just the connected component we started in

- We look at each edge only once.
- And basically don't do anything else.
- So...

$O(m)$



- (Assuming we are using the linked-list representation)

Verify this
formally!

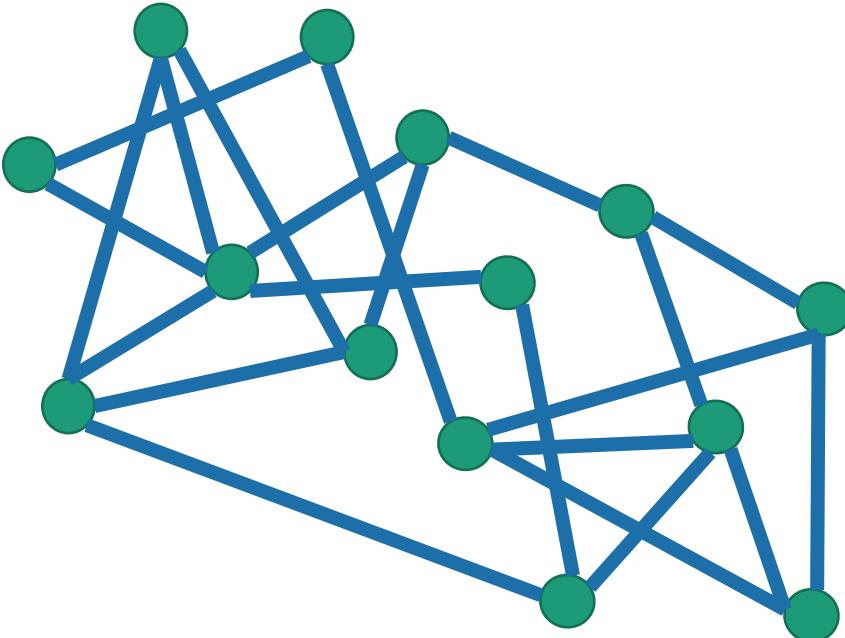


Ollie the over-achieving ostrich

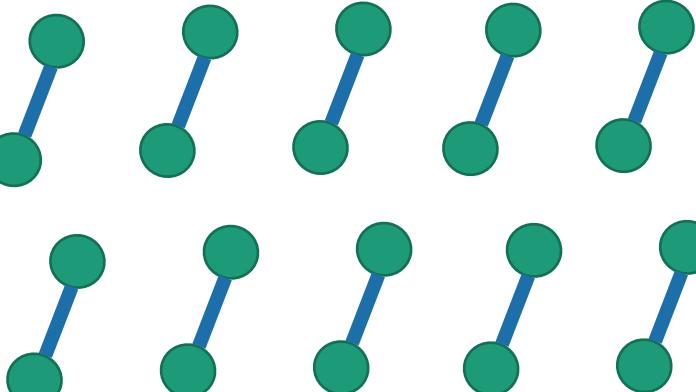
Running time

To explore the whole thing

- Explore the connected components one-by-one.
- This takes time

$$O(n + m)$$


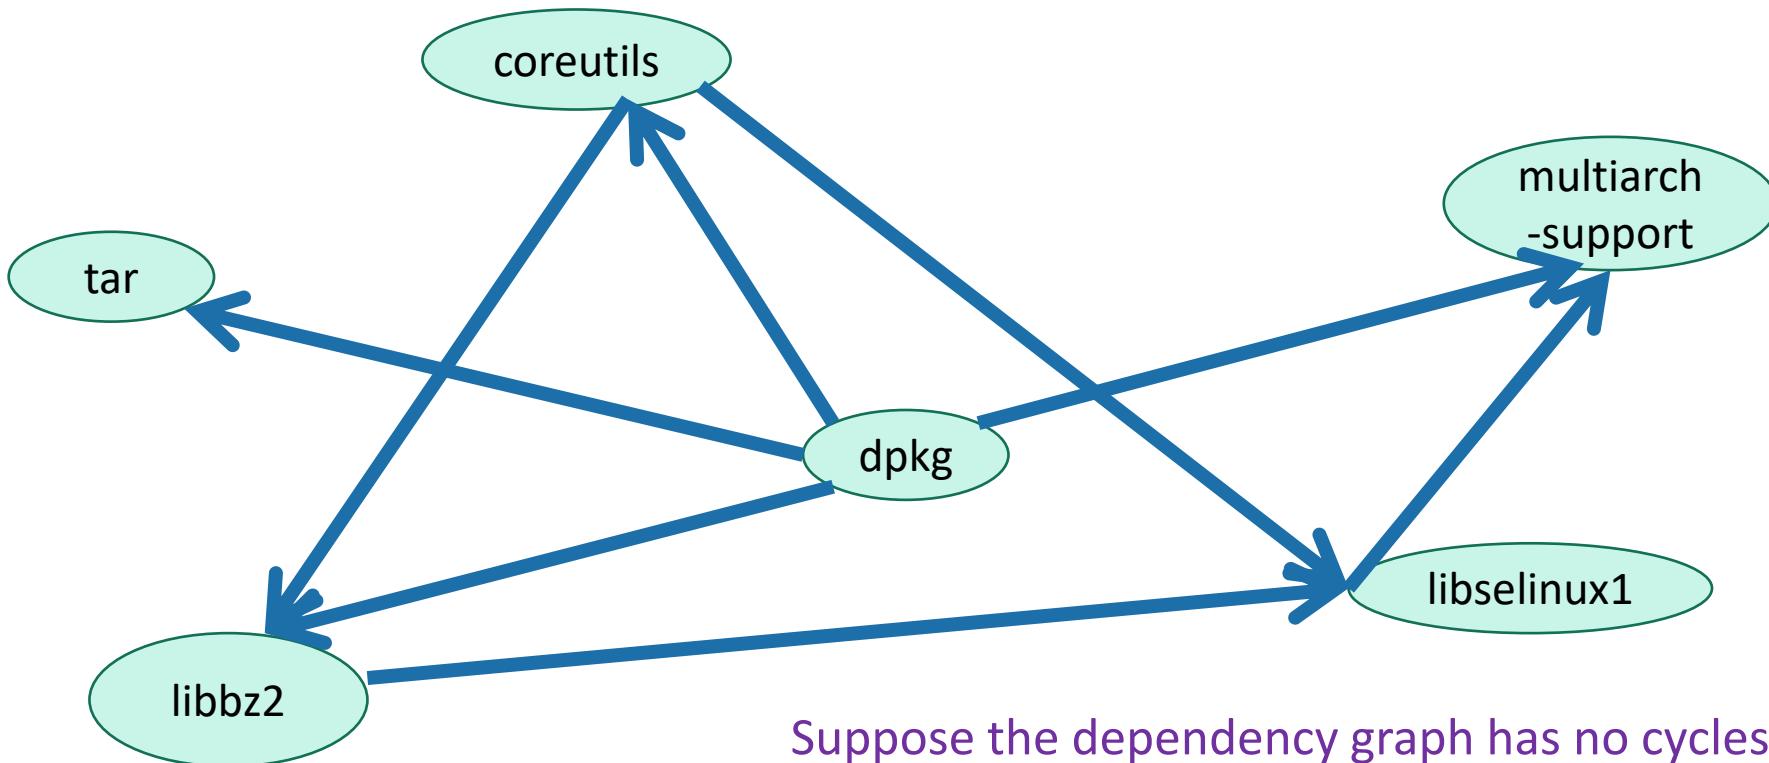
or



Applications of DFS

topological sorting: Application of DFS

- Example: package dependency graph
- Question: in what order should I install packages?

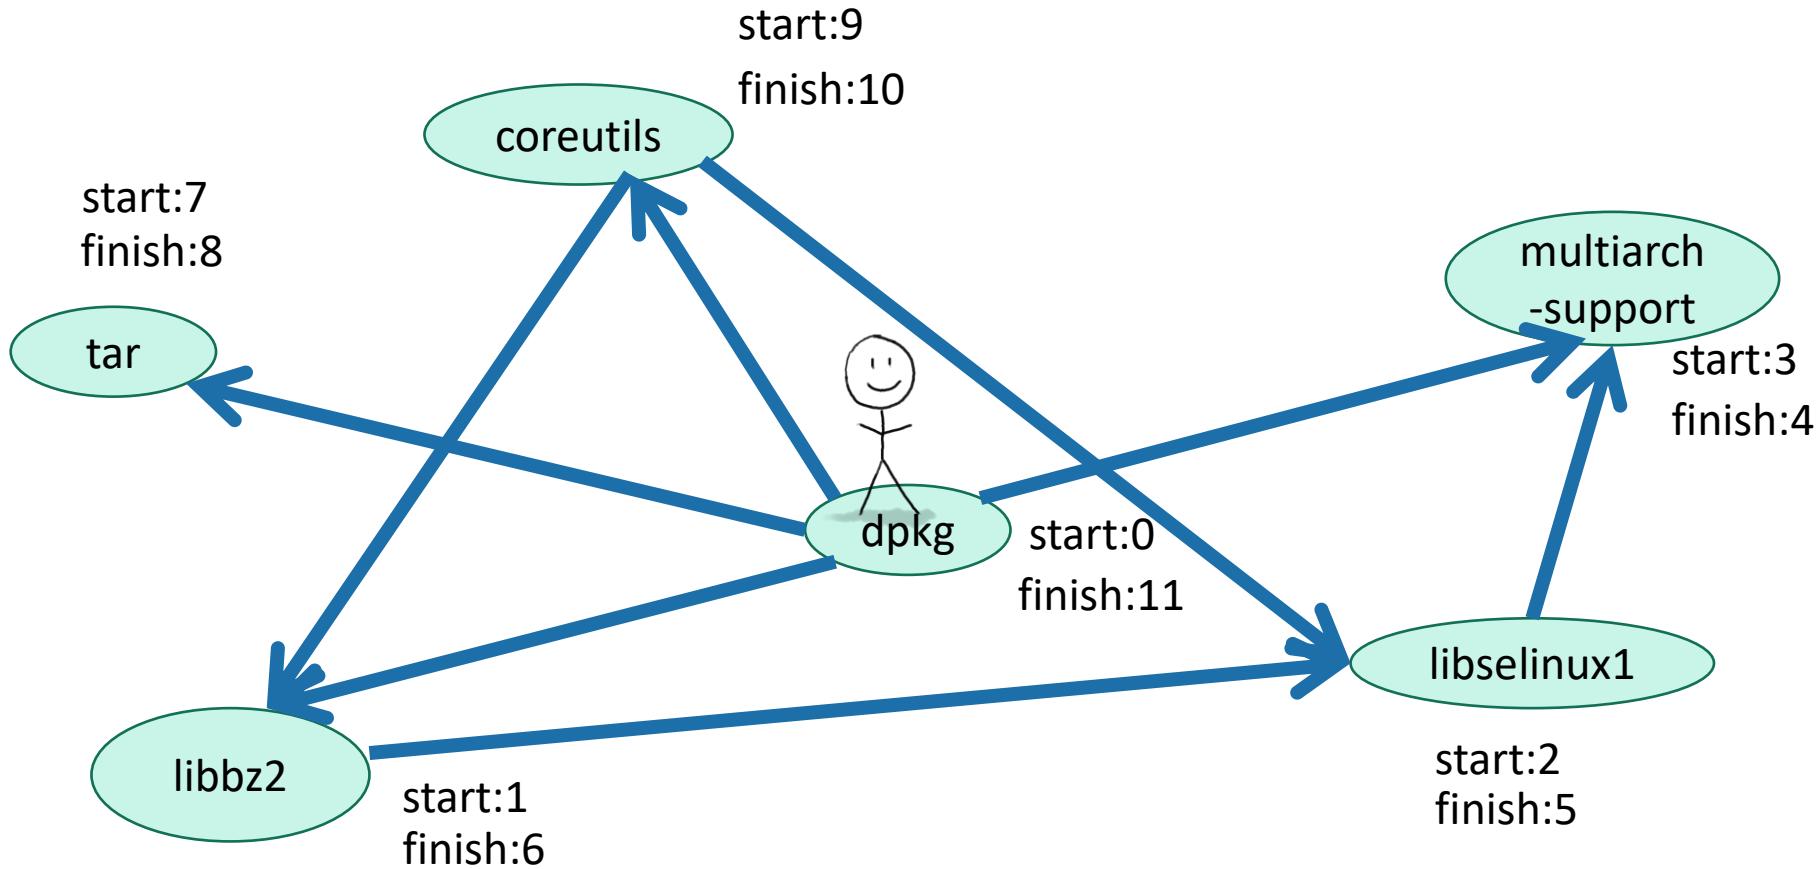


Suppose the dependency graph has no cycles:
it is a **Directed Acyclic Graph (DAG)**

Let's do DFS

Observations:

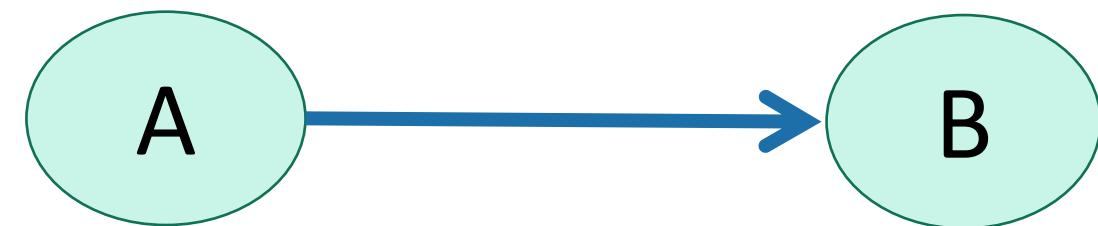
- The start times don't seem that useful.
- But the packages we should include **earlier** have **larger finish times**.



This is not an accident

Suppose the underlying
graph has no cycles

Claim: In general, we'll always have:



finish: [larger]

finish: [smaller]

To understand why, let's go back to that DFS tree.

A more general statement

(this holds even if there are cycles)

This is called the “parentheses theorem” in CLRS



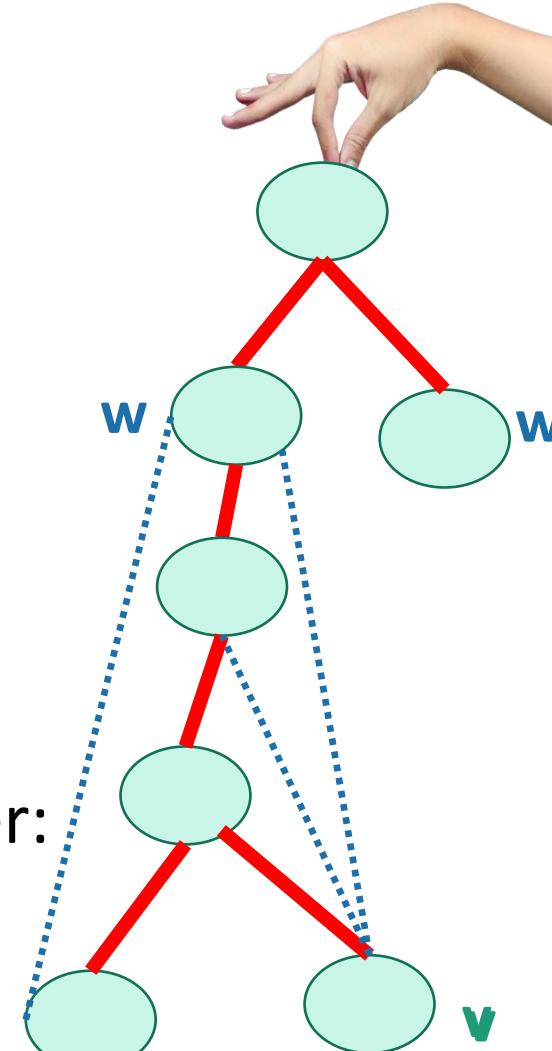
- If v is a descendent of w in this tree:



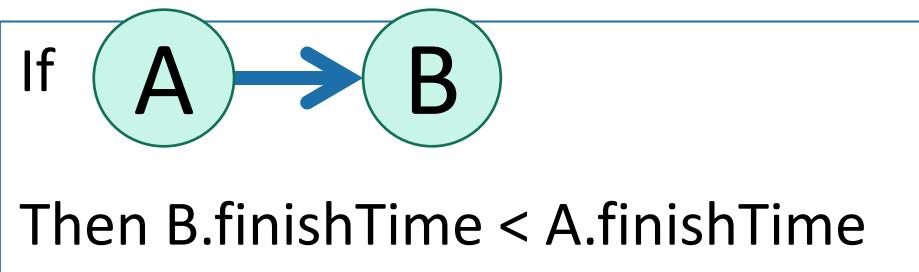
- If w is a descendent of v in this tree:



- If neither are descendants of each other:



So to prove this ->



Suppose the underlying graph has no cycles

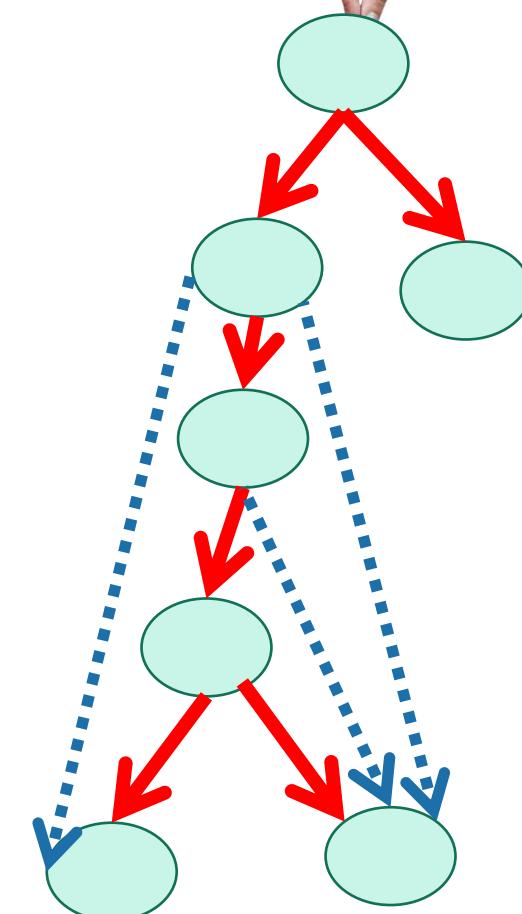
- Since the graph has no cycles, B must be a descendent of A in that tree.

- All edges go down the tree.

- Then

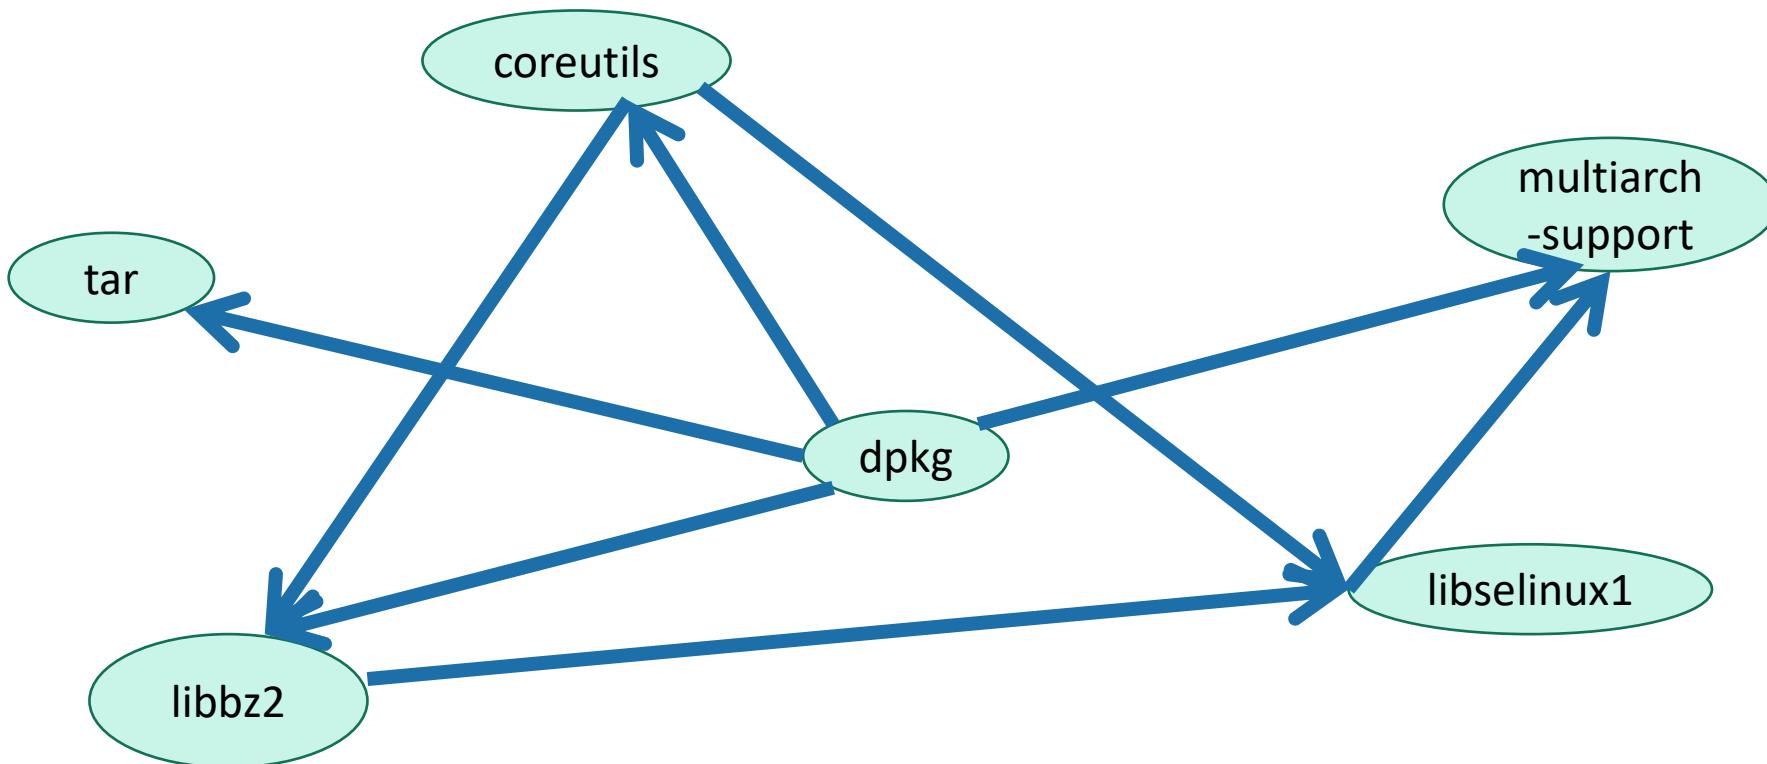


- aka, $B.\text{finishTime} < A.\text{finishTime}$.



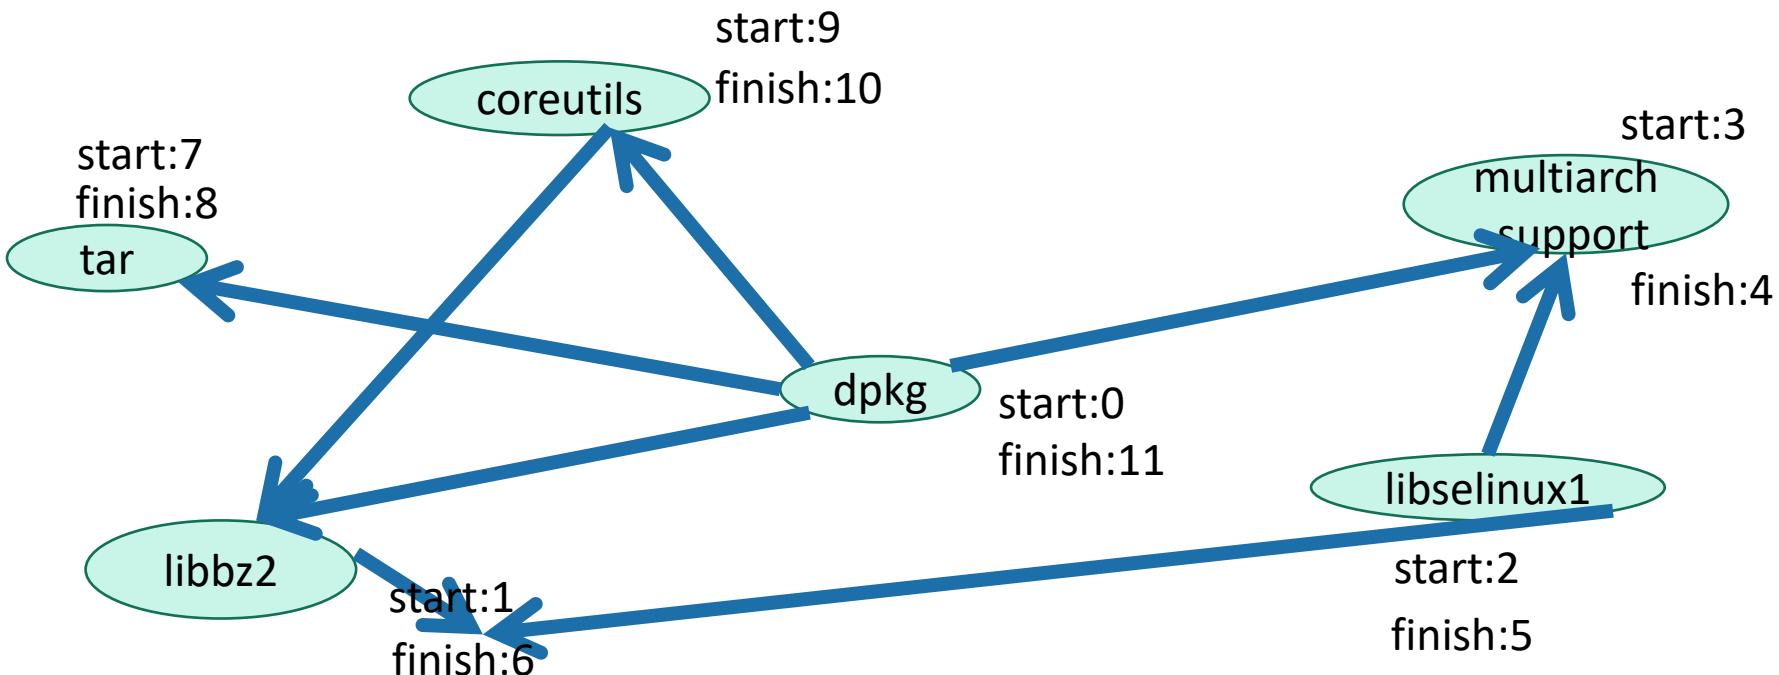
Back to this problem

- Example: package dependency graph
- Question: in what order should I install packages?



In reverse order of finishing time

- Do DFS
 - Maintain a list of packages, in the order you want to install them.
 - When you mark a vertex as **all done**, put it at the **beginning** of the list.
- dpkg
 - coreutils
 - tar
 - libbz2
 - libselinux1
 - multiarch_support

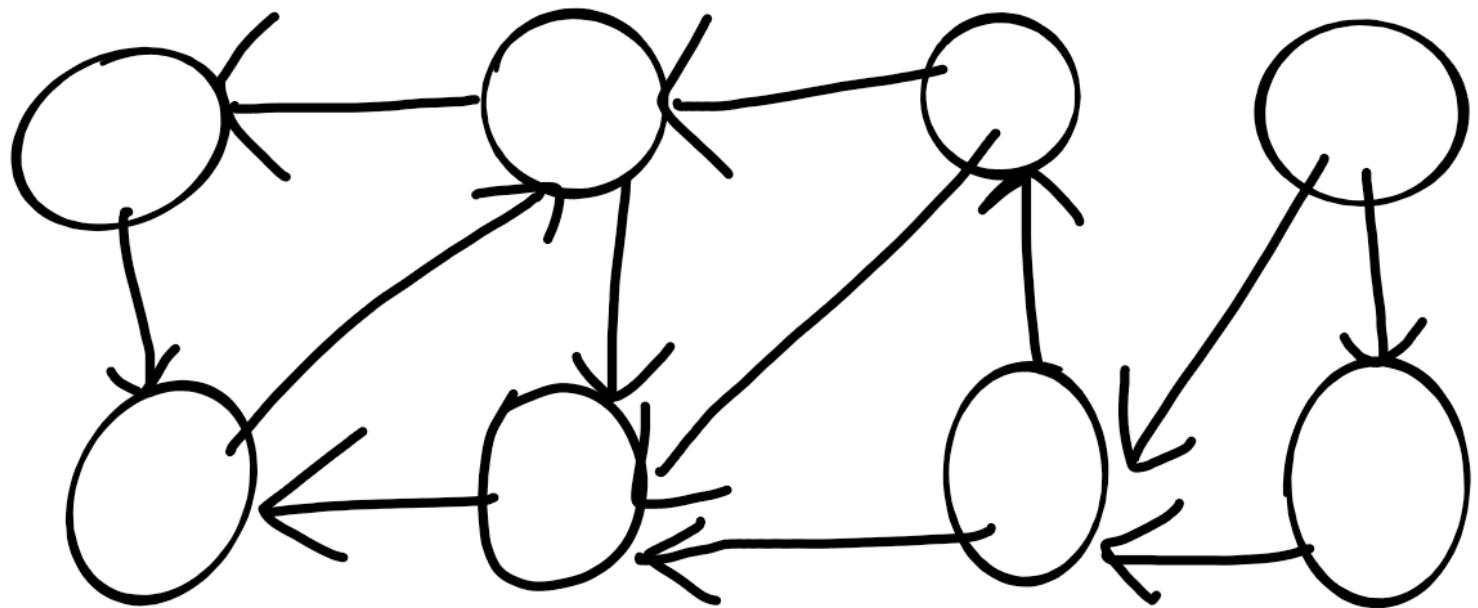


What did we just learn?

- DFS can help you solve the **TOPOLOGICAL SORTING PROBLEM**
 - That's the fancy name for the problem of finding an ordering that respects all the dependencies
- Thinking about the DFS tree is helpful.

Classification of Edges

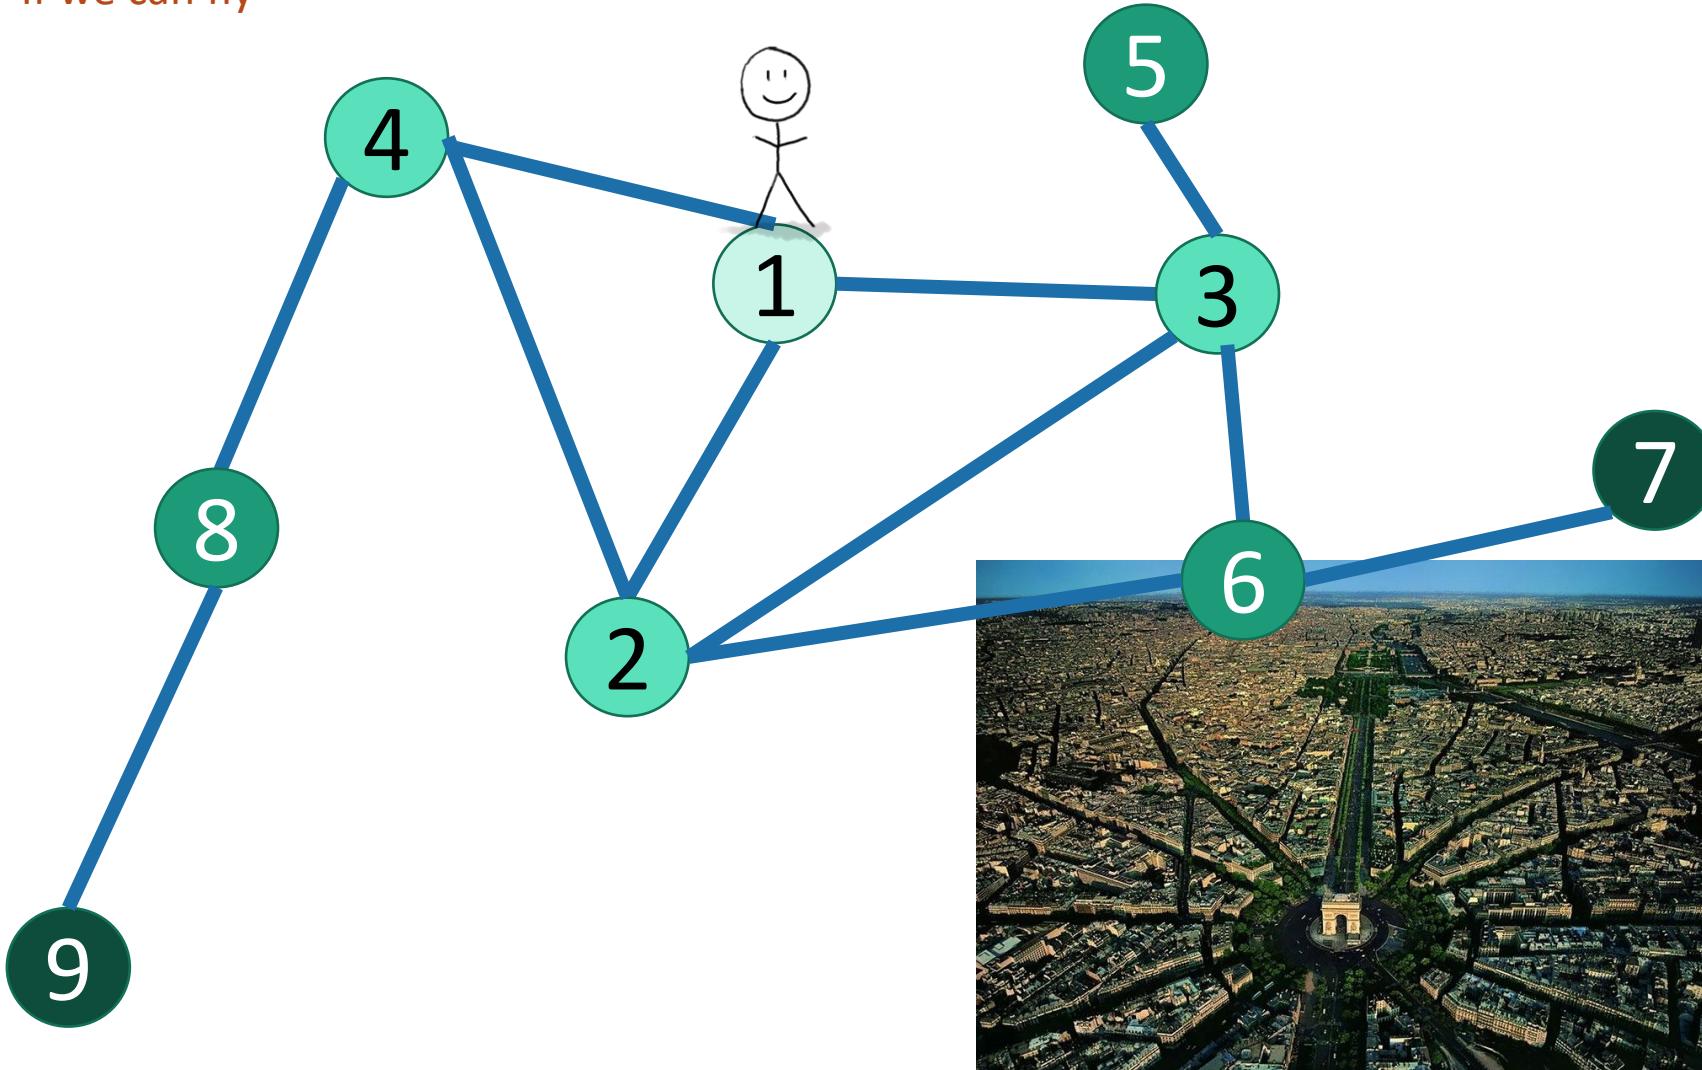
- Tree Edge
- Back Edge
- Forward Edge
- Cross Edge



Breadth-First Search

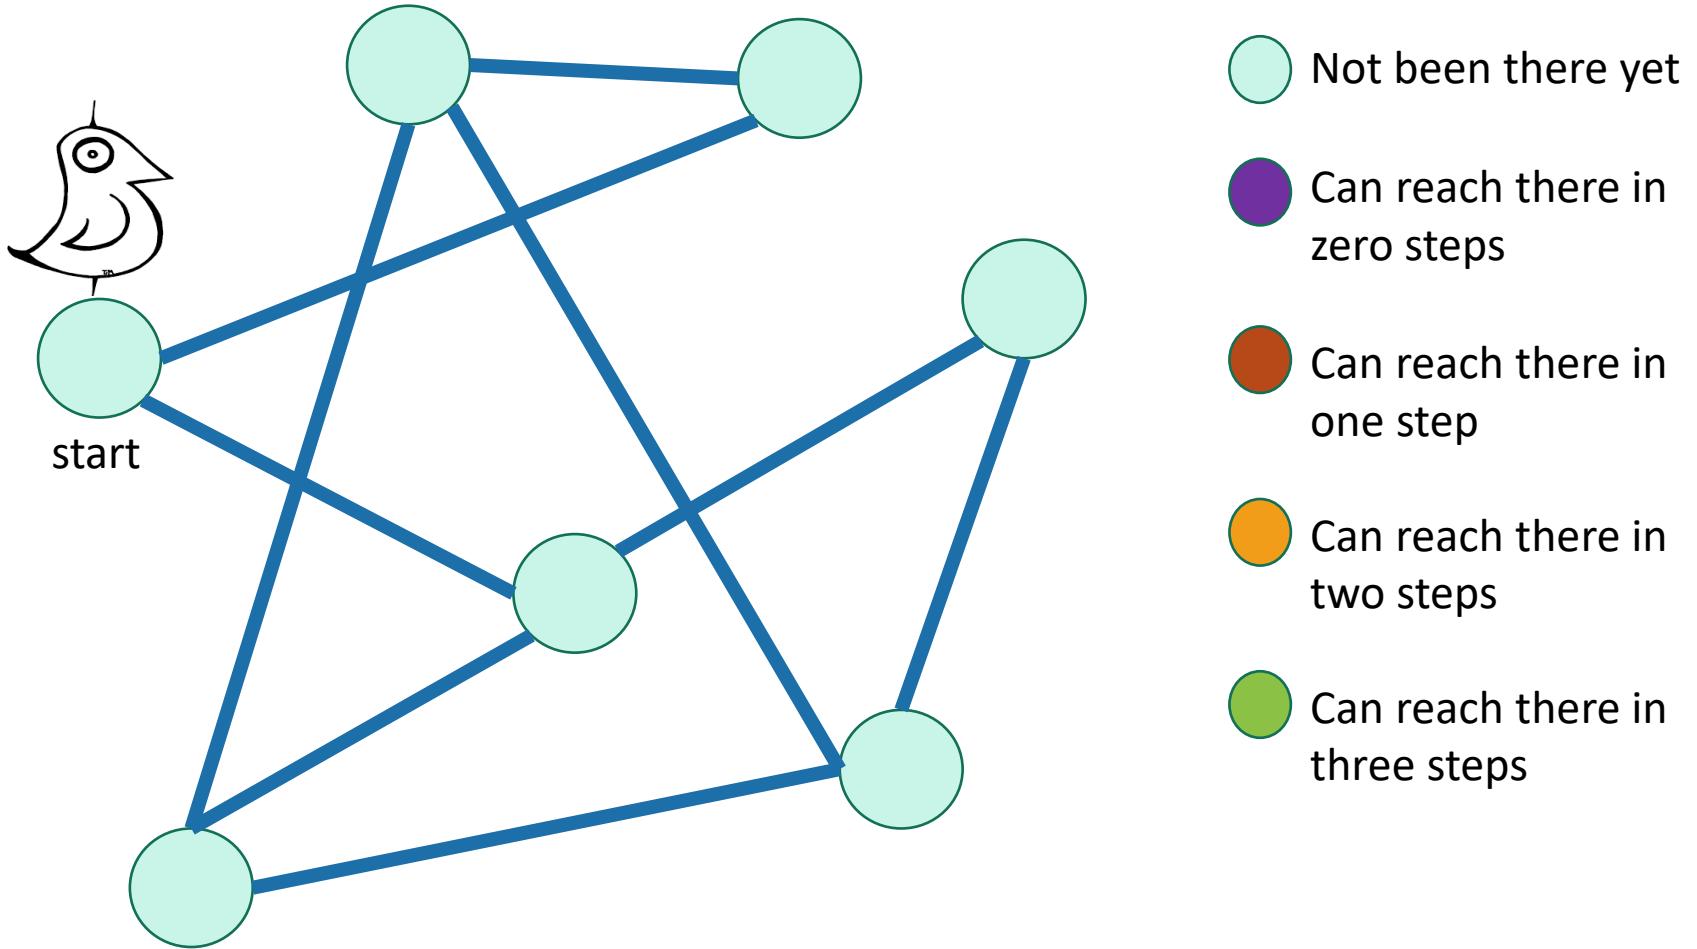
How do we explore a graph?

If we can fly



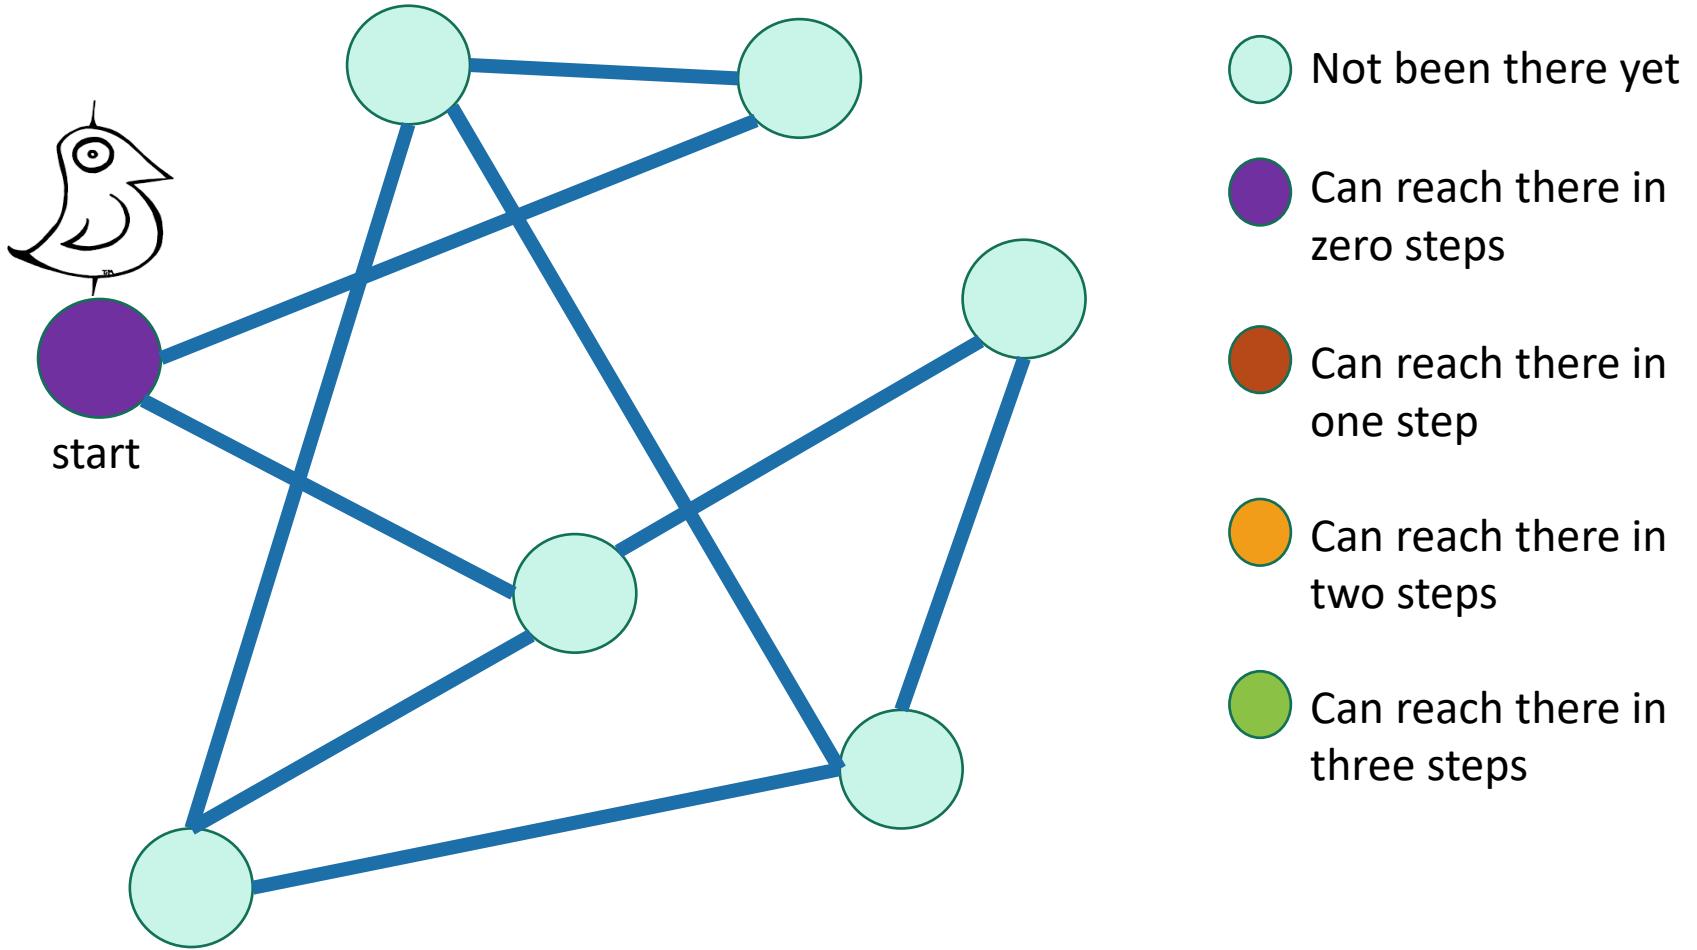
Breadth-First Search

Exploring the world with a bird's-eye view



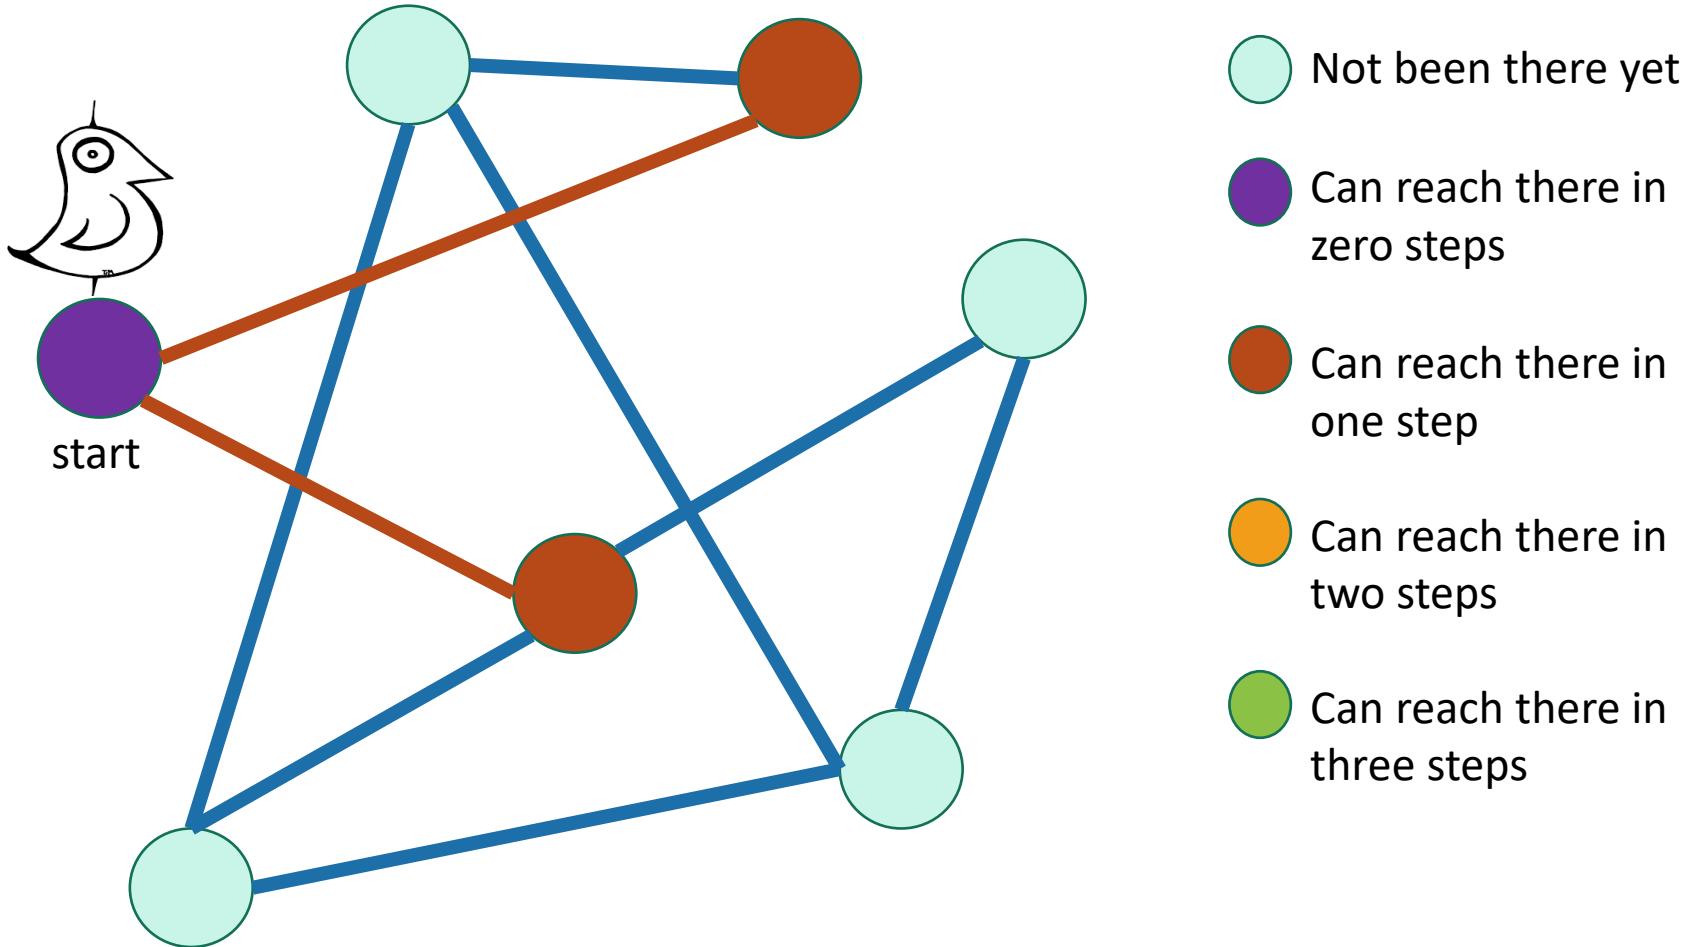
Breadth-First Search

Exploring the world with a bird's-eye view



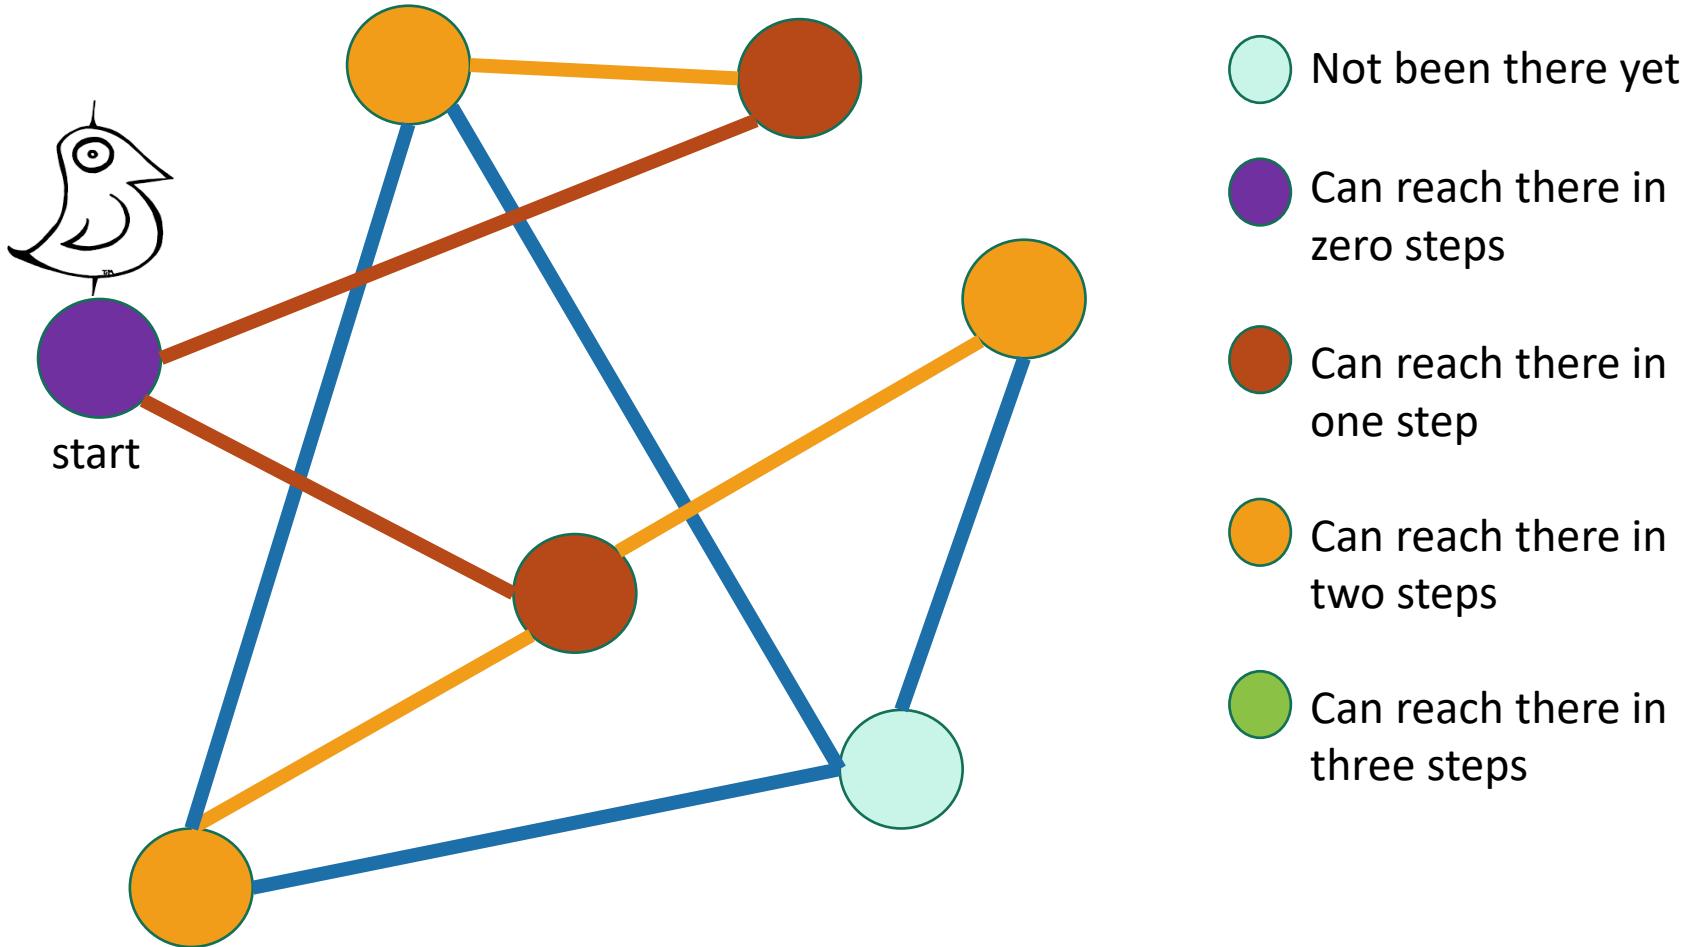
Breadth-First Search

Exploring the world with a bird's-eye view



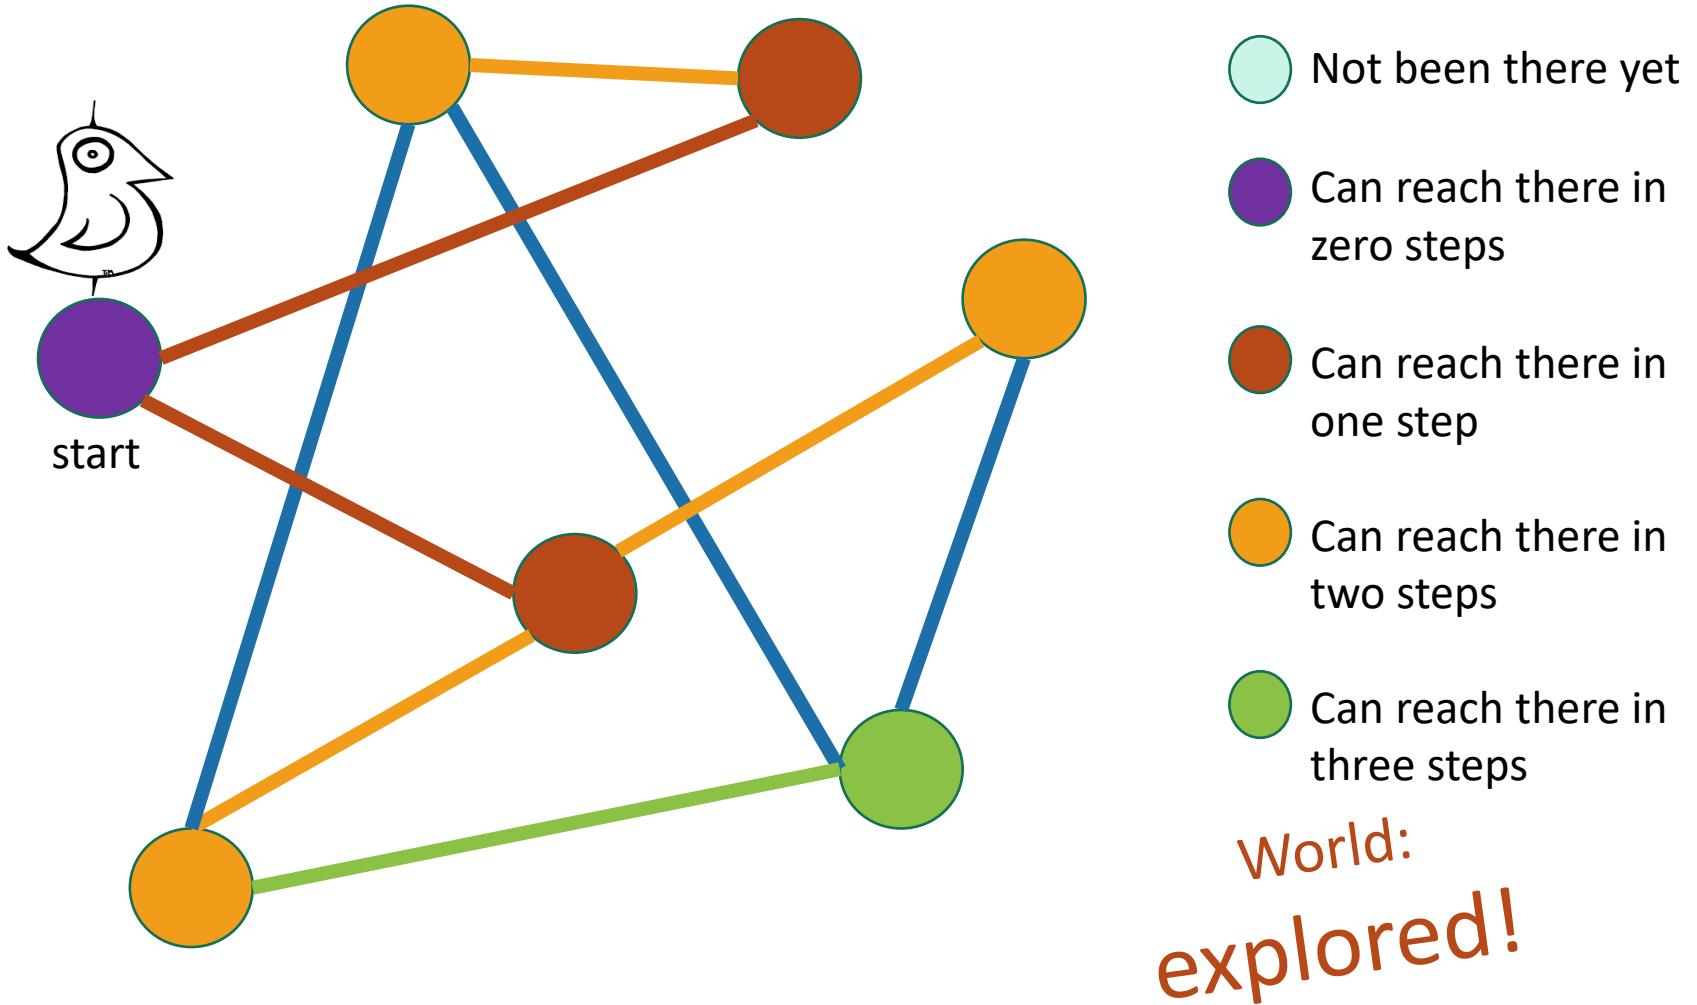
Breadth-First Search

Exploring the world with a bird's-eye view



Breadth-First Search

Exploring the world with a bird's-eye view



Same disclaimer as for DFS: you may have seen other ways to implement this,
this will be convenient for us.

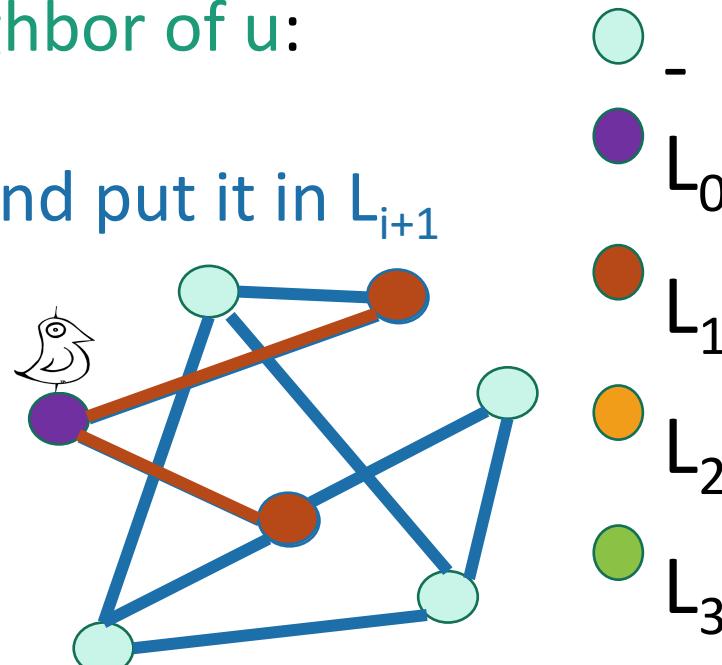
Breadth-First Search

Exploring the world with pseudocode

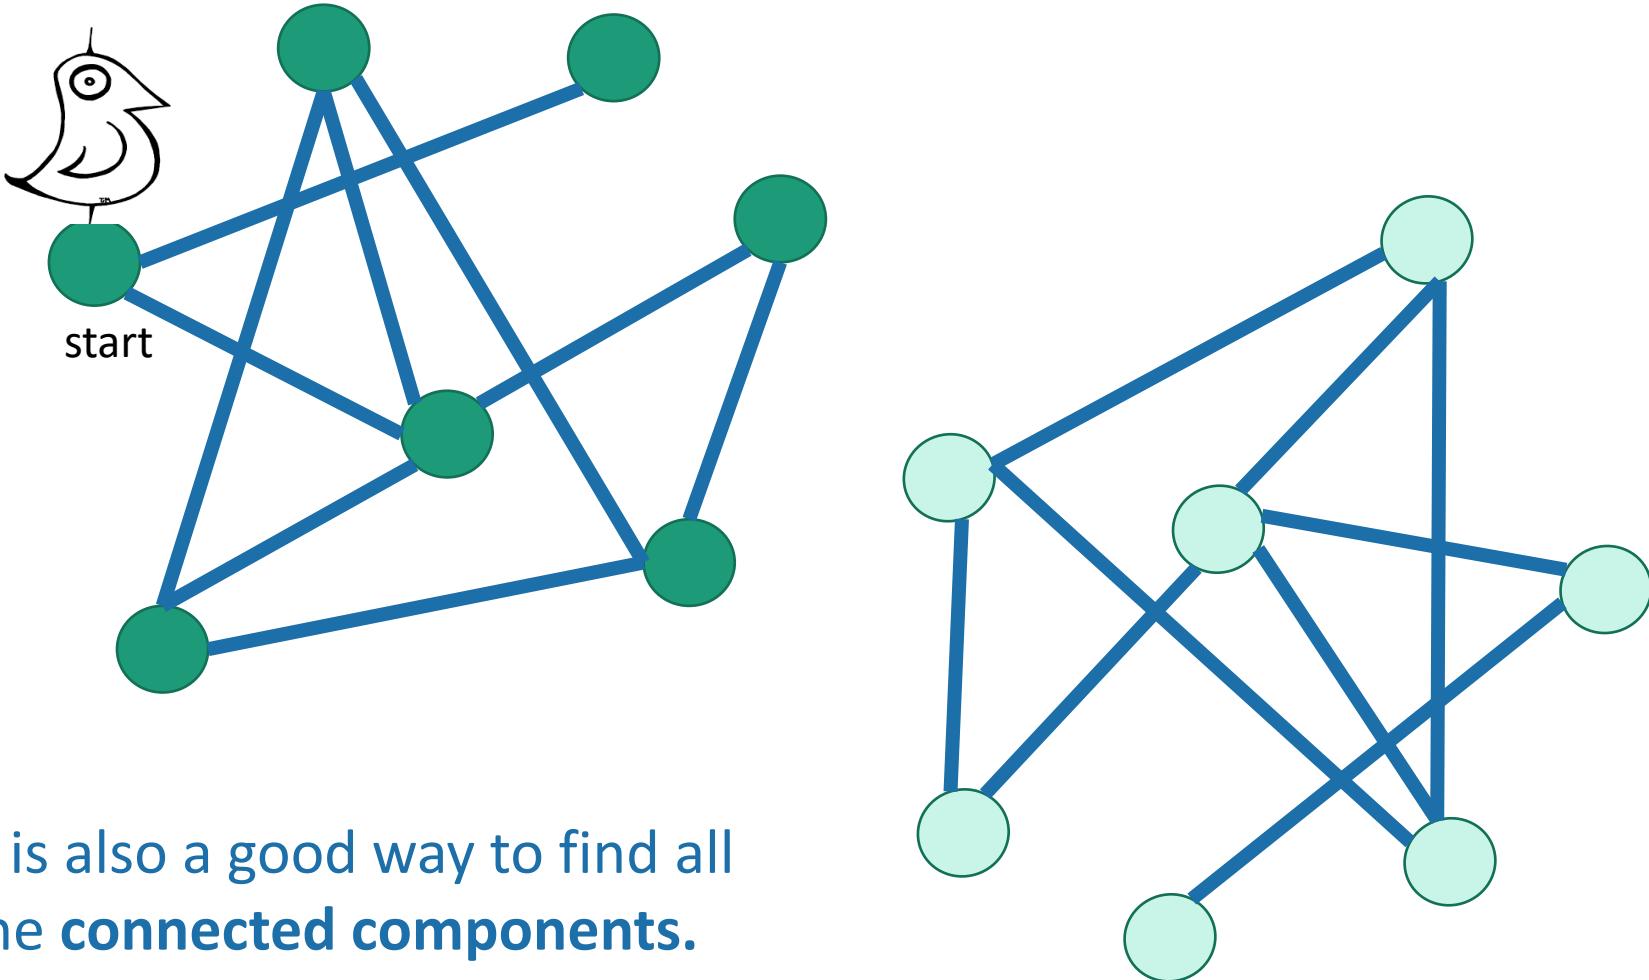
- Set $L_i = \{\}$ for $i=1, \dots, n$
- $L_0 = \{w\}$, where w is the start node
- **For** $i = 0, \dots, n-1$:
 - **For** u in L_i :
 - **For** each v which is a neighbor of u :
 - **If** v isn't yet visited:
 - mark v as visited, and put it in L_{i+1}

L_i is the set of nodes
we can reach in i
steps from w

Go through all the nodes
in L_i and add their
unvisited neighbors to L_{i+1}



BFS also finds all the nodes reachable from the starting point



Running time

To explore the whole thing

- Explore the connected components one-by-one.
- Same argument as DFS: running time is

$$O(n + m)$$

Verify these!

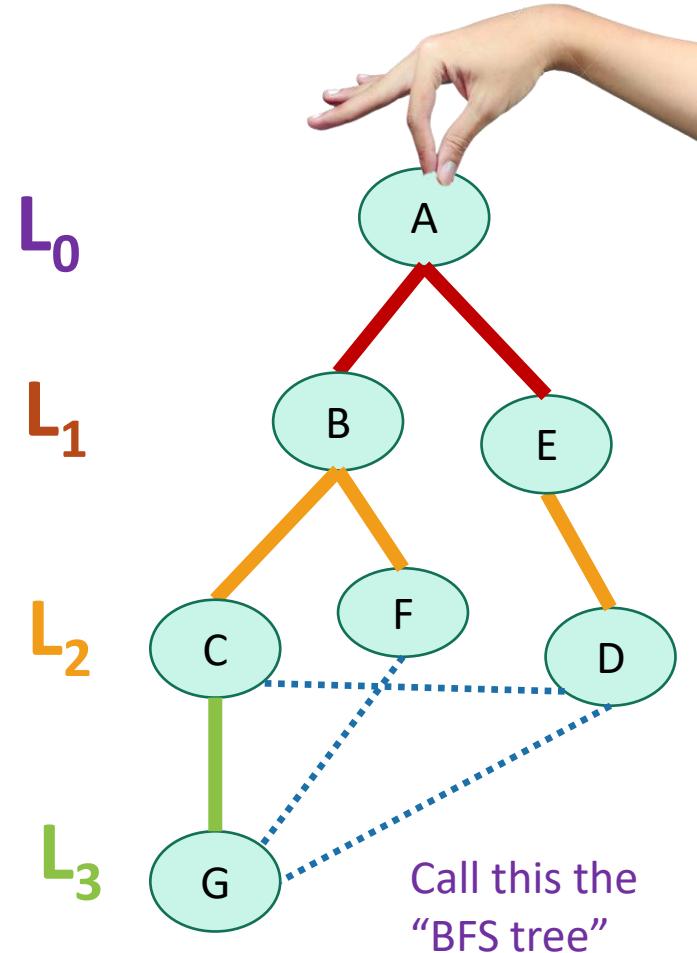
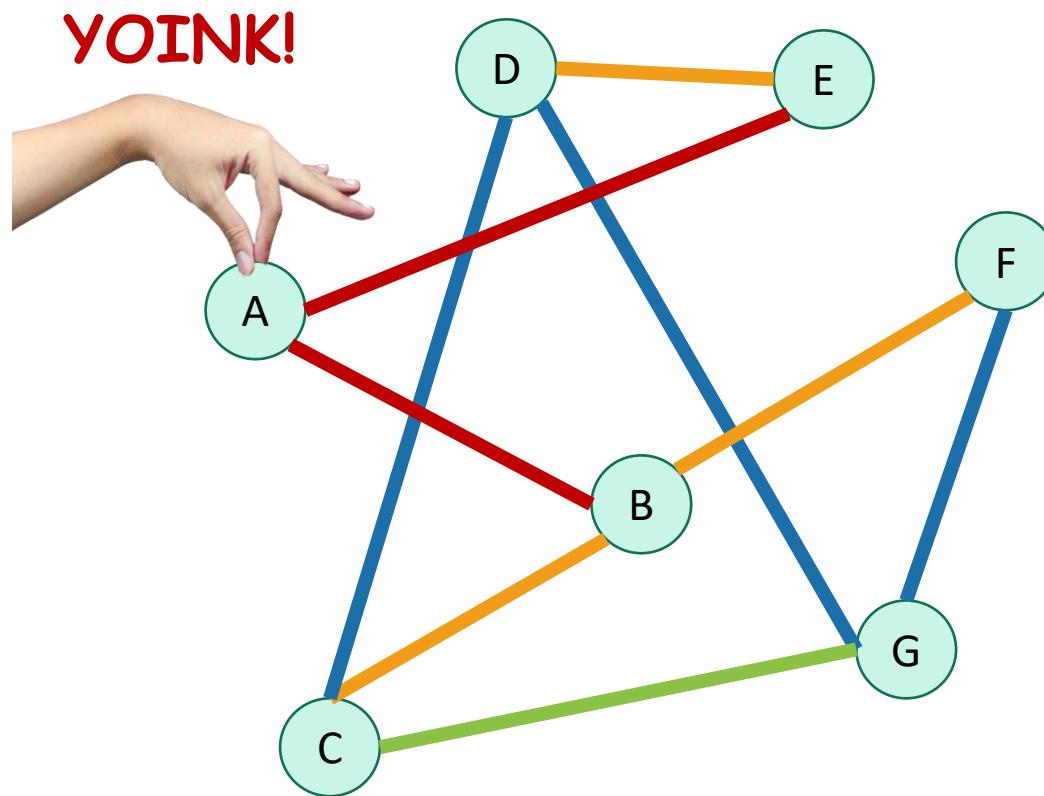


Ollie the over-achieving ostrich

- Like DFS, BFS also works fine on directed graphs.

Why is it called breadth-first?

- We are implicitly building a tree:

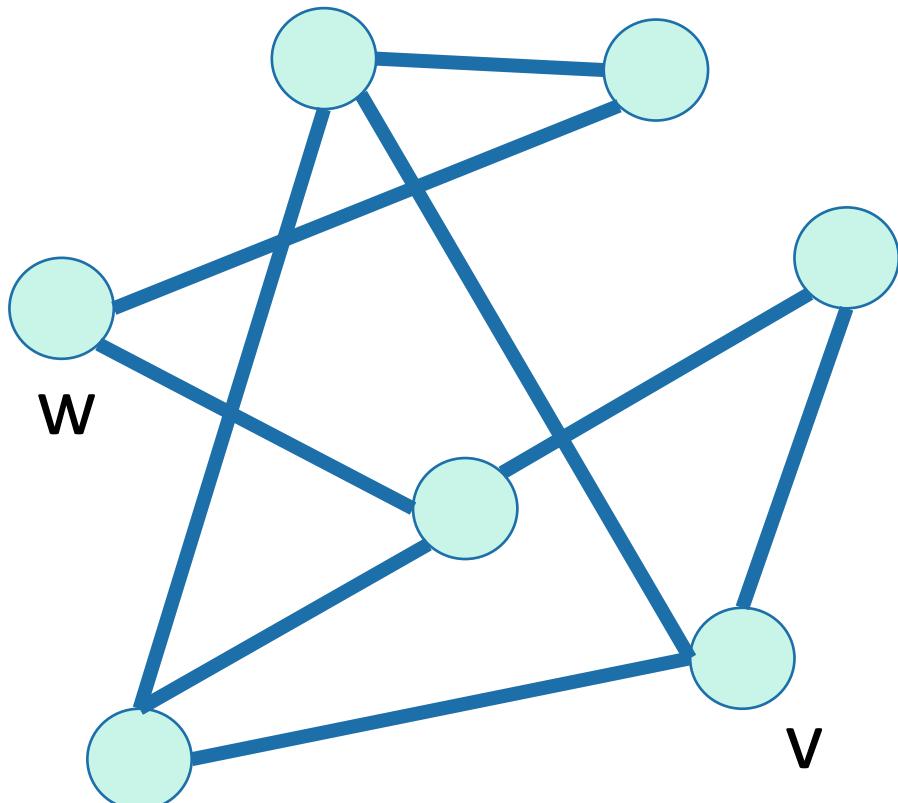


- And first we go as **broadly** as we can.

Applications of BFS

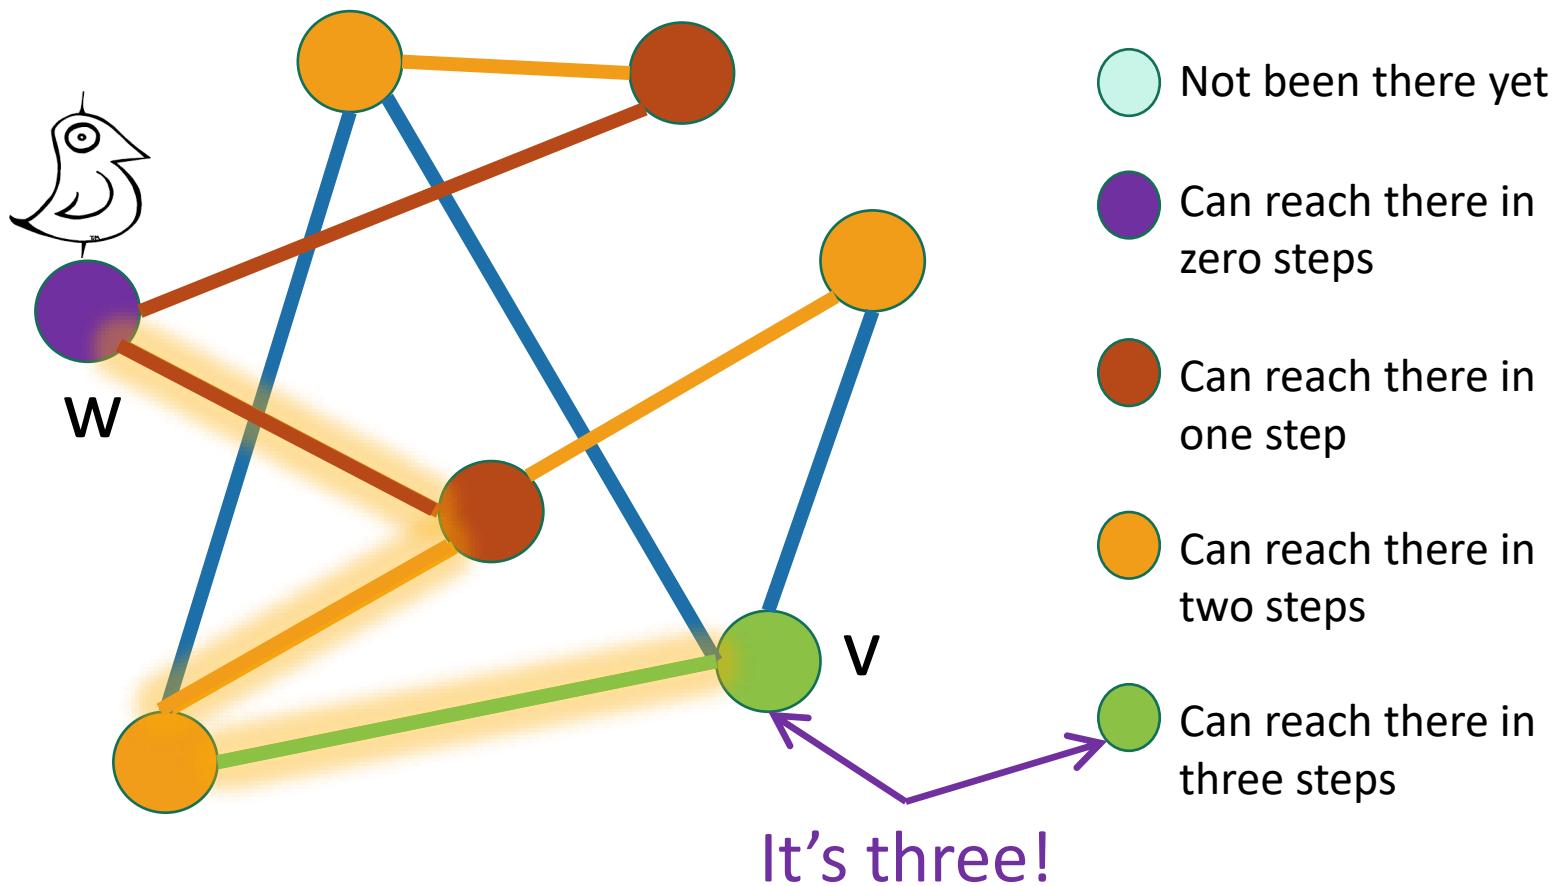
Application: shortest path

- How long is the shortest path between w and v?



Application: shortest path

- How long is the shortest path between w and v?



To find the **distance** between w and all other vertices v

- Do a BFS starting at w
- For all v in L_i (the i 'th level of the BFS tree)
 - The shortest path between w and v has length i
 - A shortest path between w and v is given by the path in the BFS tree.
- If we never found v , the distance is infinite.

The distance between two vertices is the length of the shortest path between them.

What did we just learn?

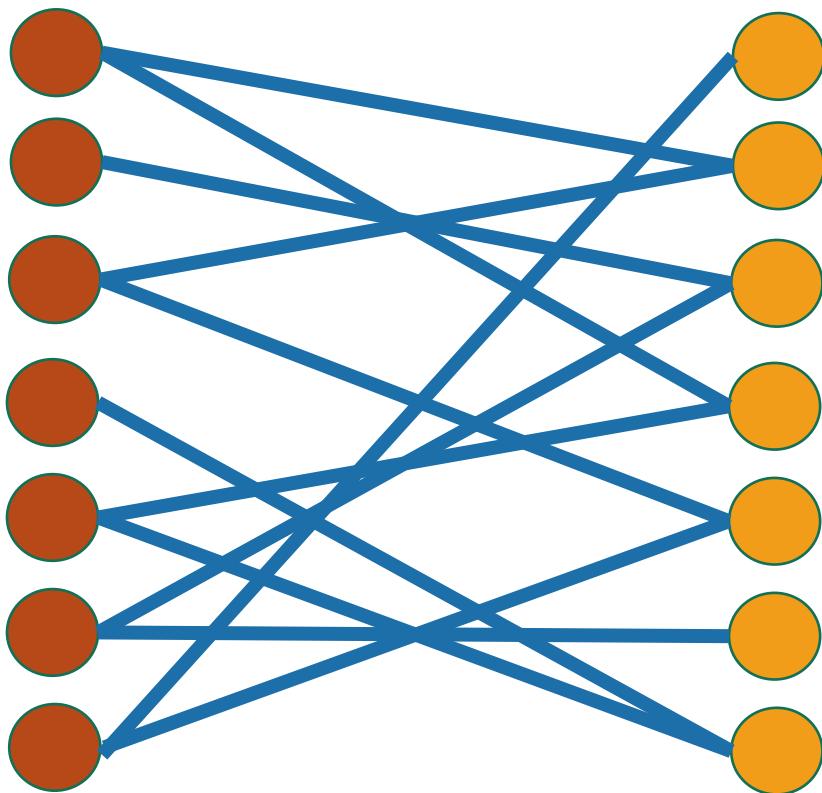
- The BFS tree is useful for computing distances between pairs of vertices.
- We can find the shortest path between u and v in time $O(m)$.

The BSF tree is also helpful for:

- Testing if a graph is bipartite or not.

Application: testing if a graph is bipartite

- Bipartite means it looks like this:

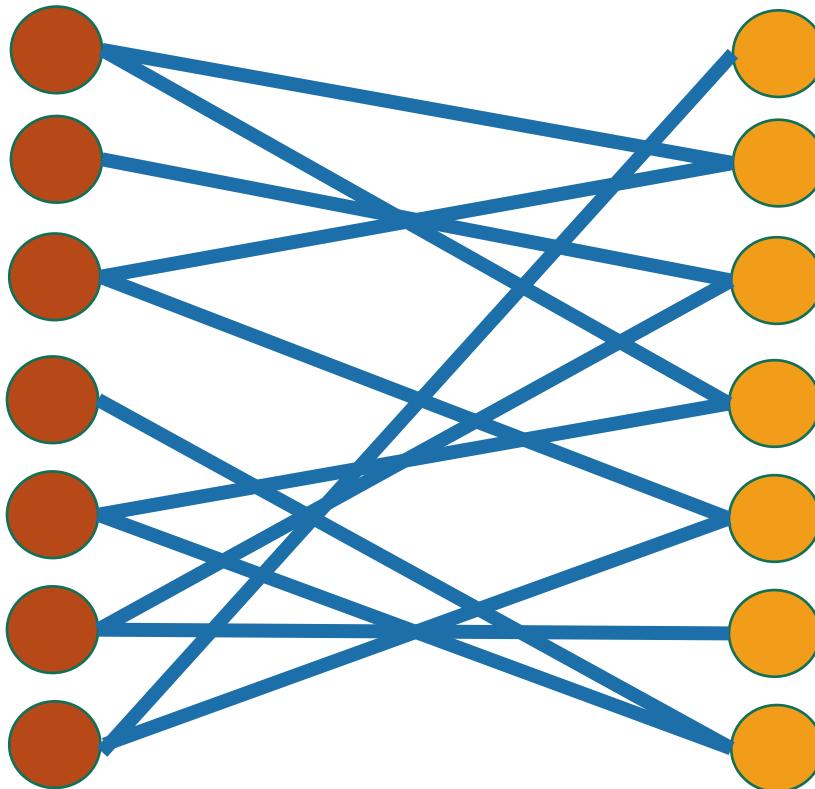


Can color the vertices red and orange so that there are no edges between any same-colored vertices

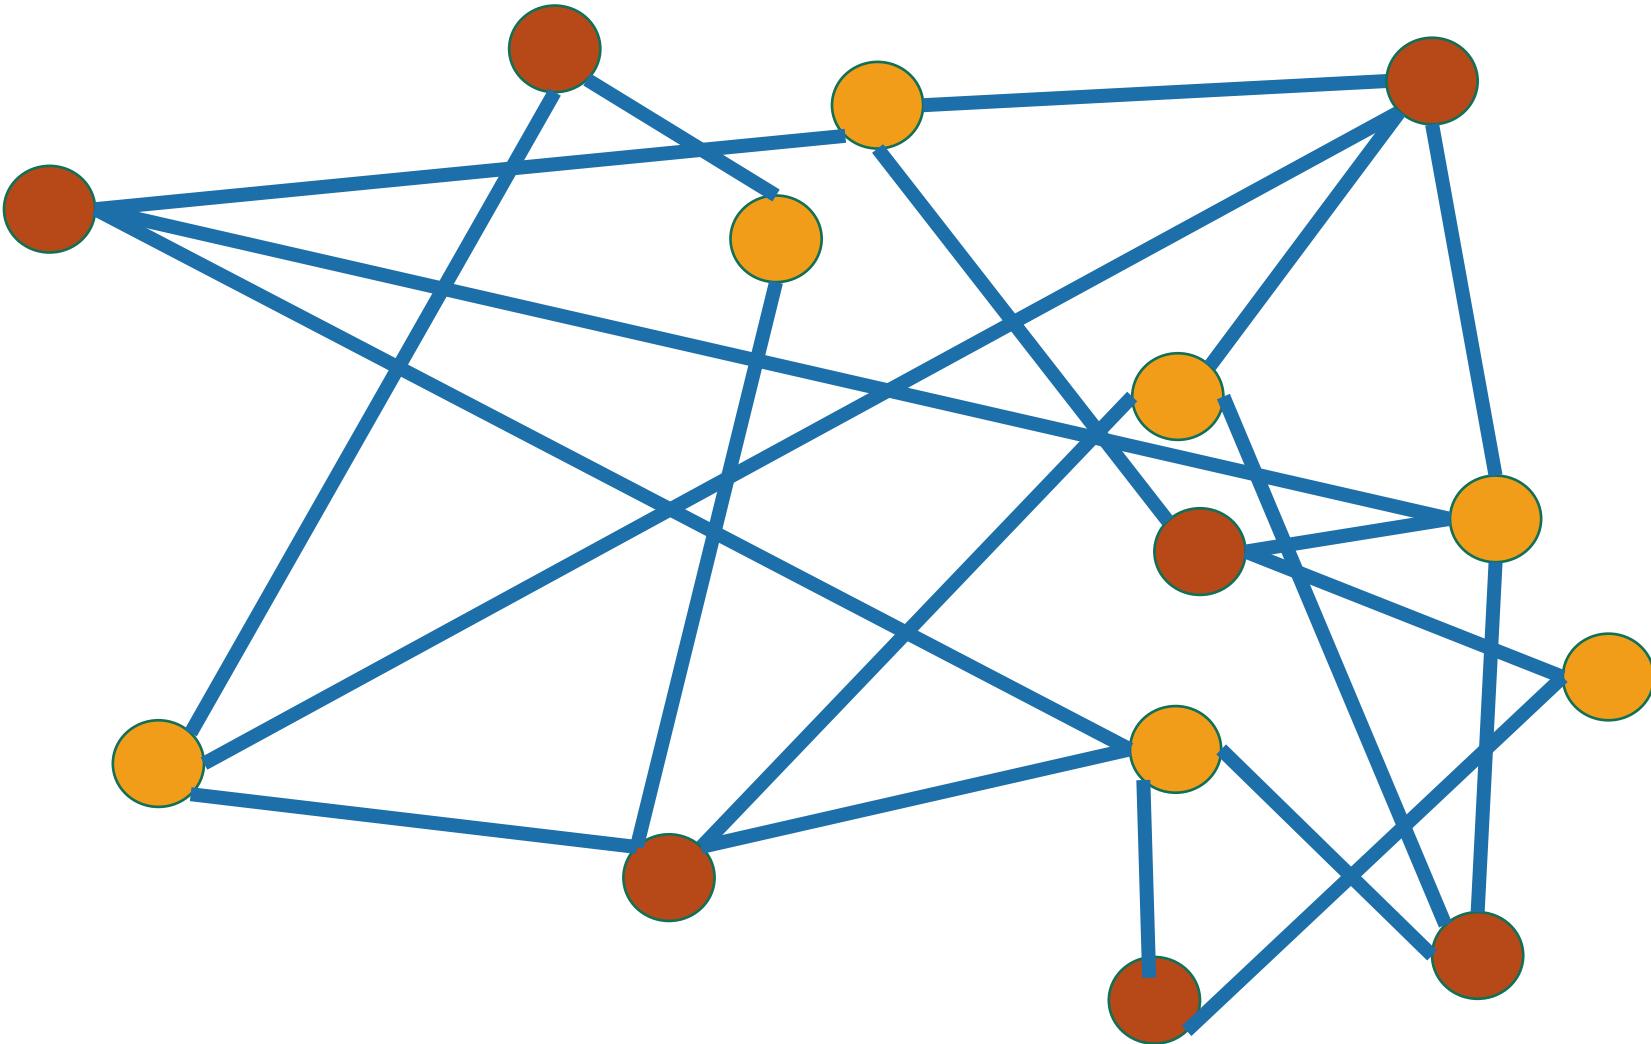
Example:

are students
 are classes
if the student is enrolled in the class

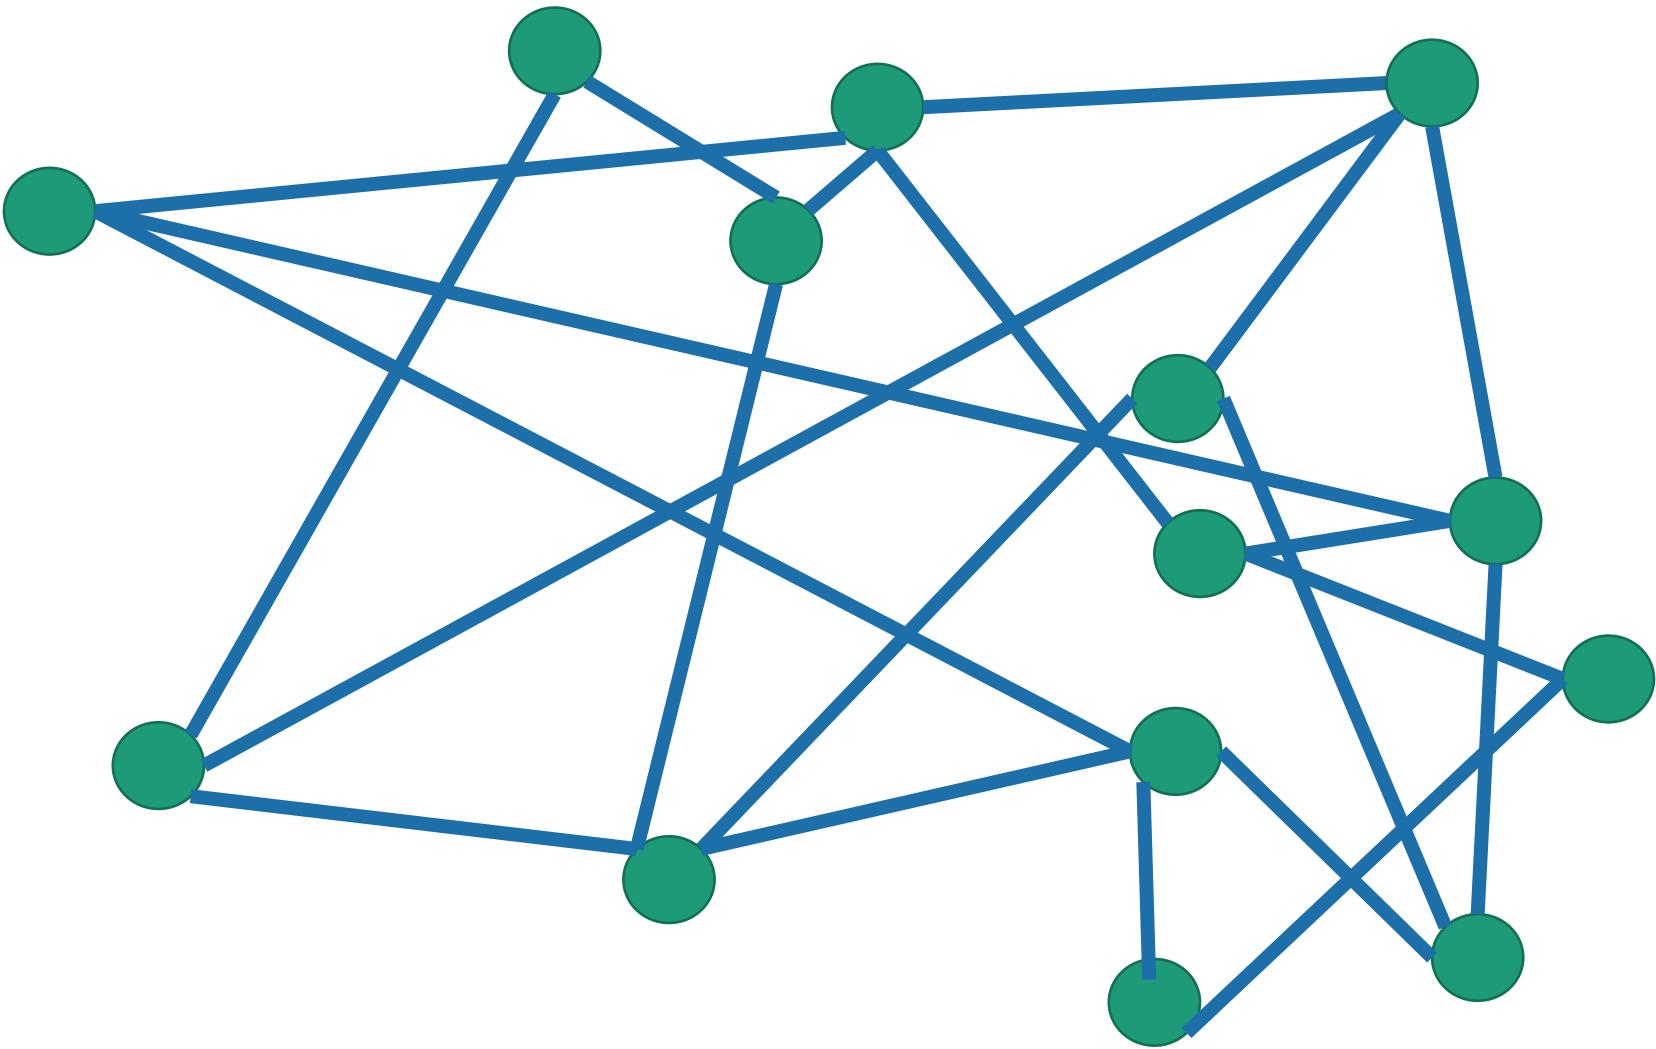
Is this graph bipartite?



How about this one?

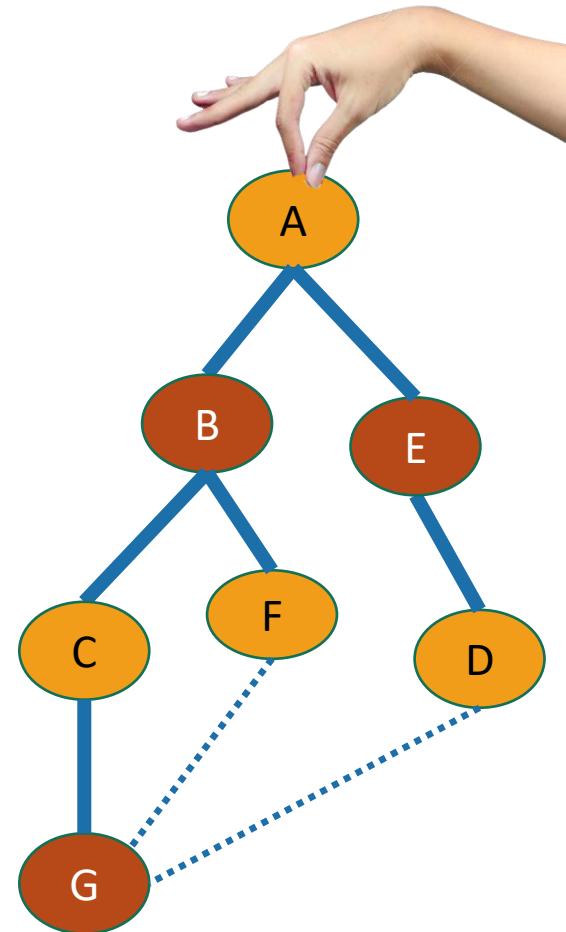


How about this one?



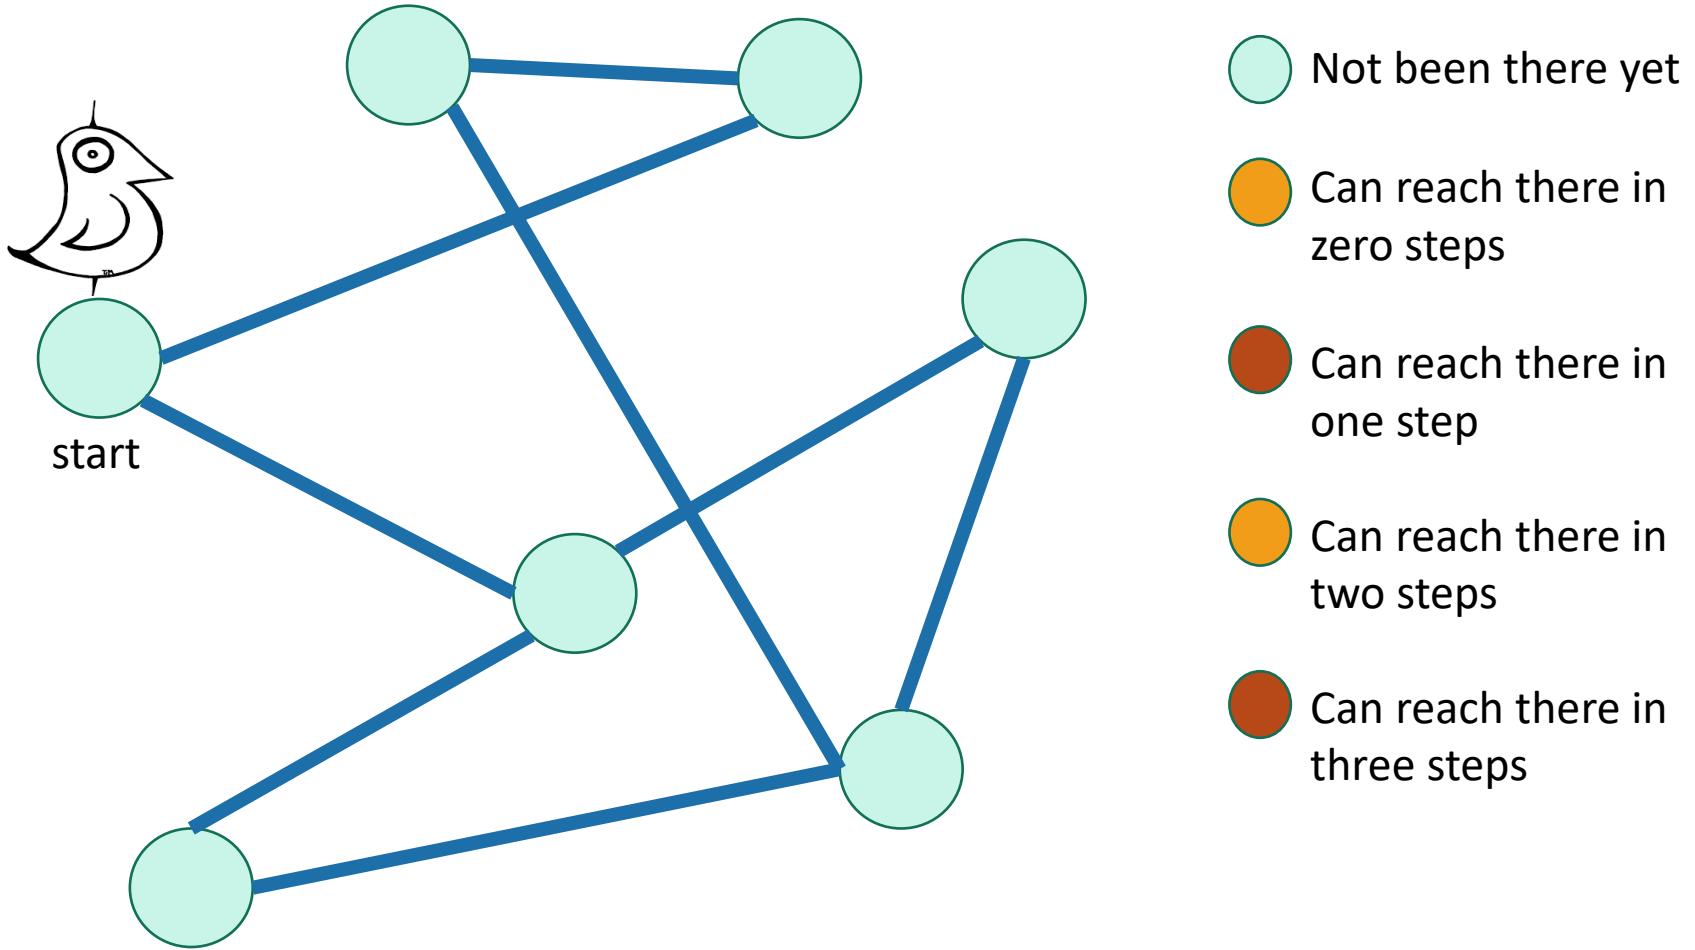
Solution using BFS

- Color the levels of the BFS tree in alternating colors.
- If you ever color a node so that you never color two connected nodes the same, then it is bipartite.
- Otherwise, it's not.



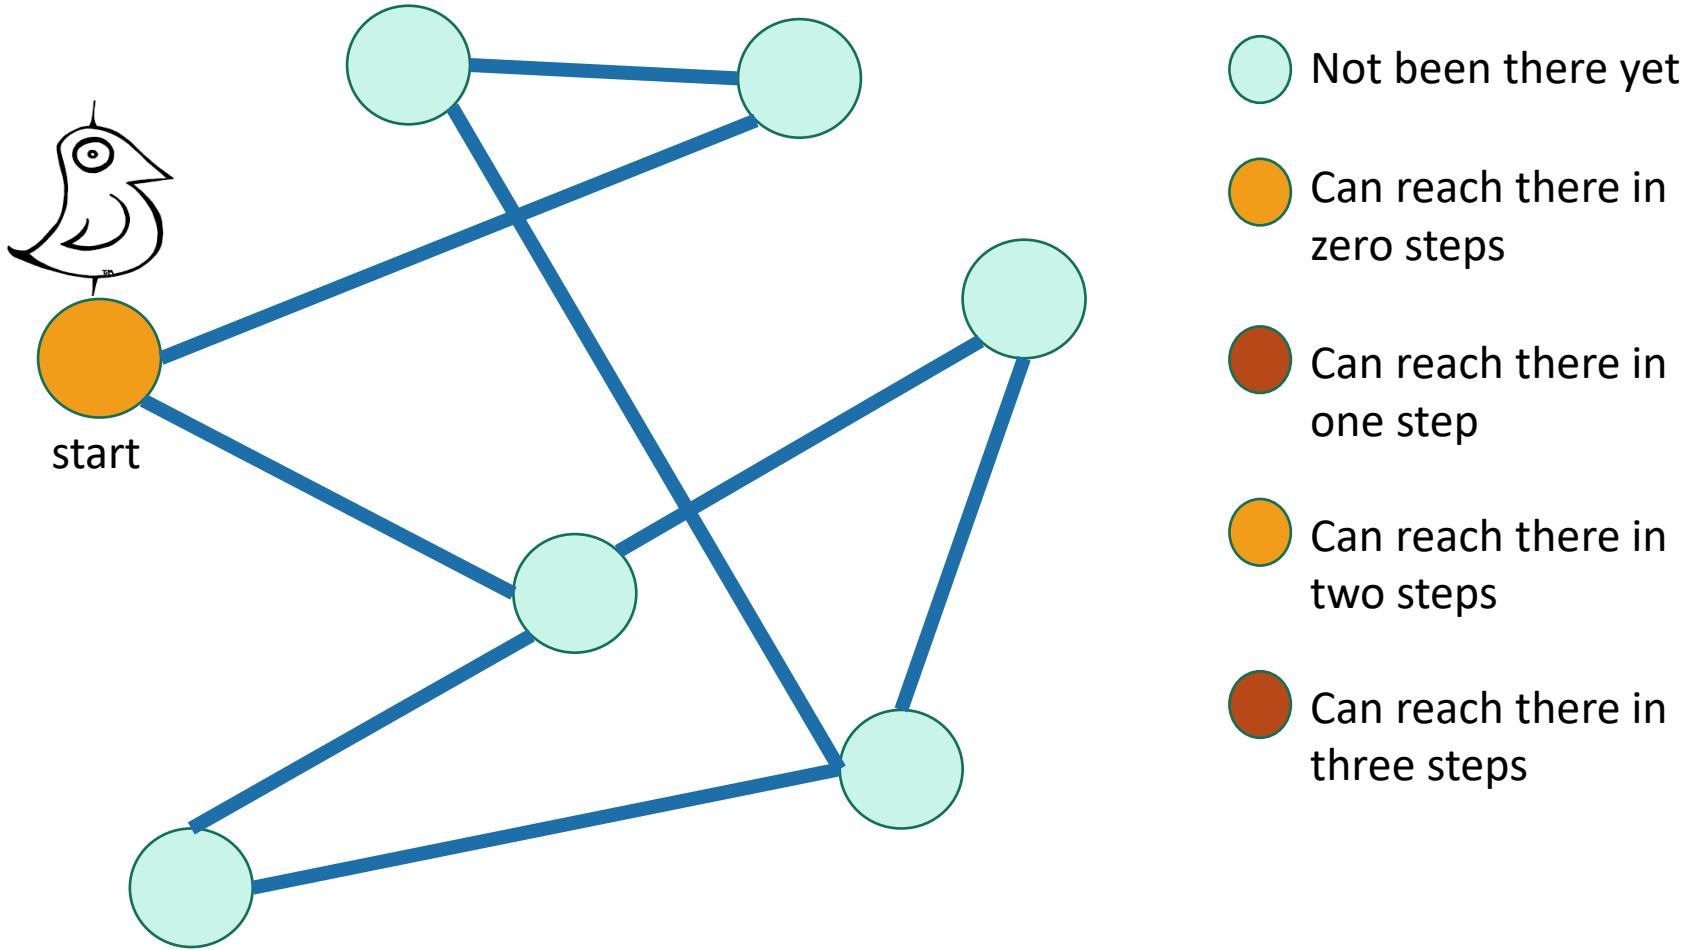
Breadth-First Search

For testing bipartite-ness



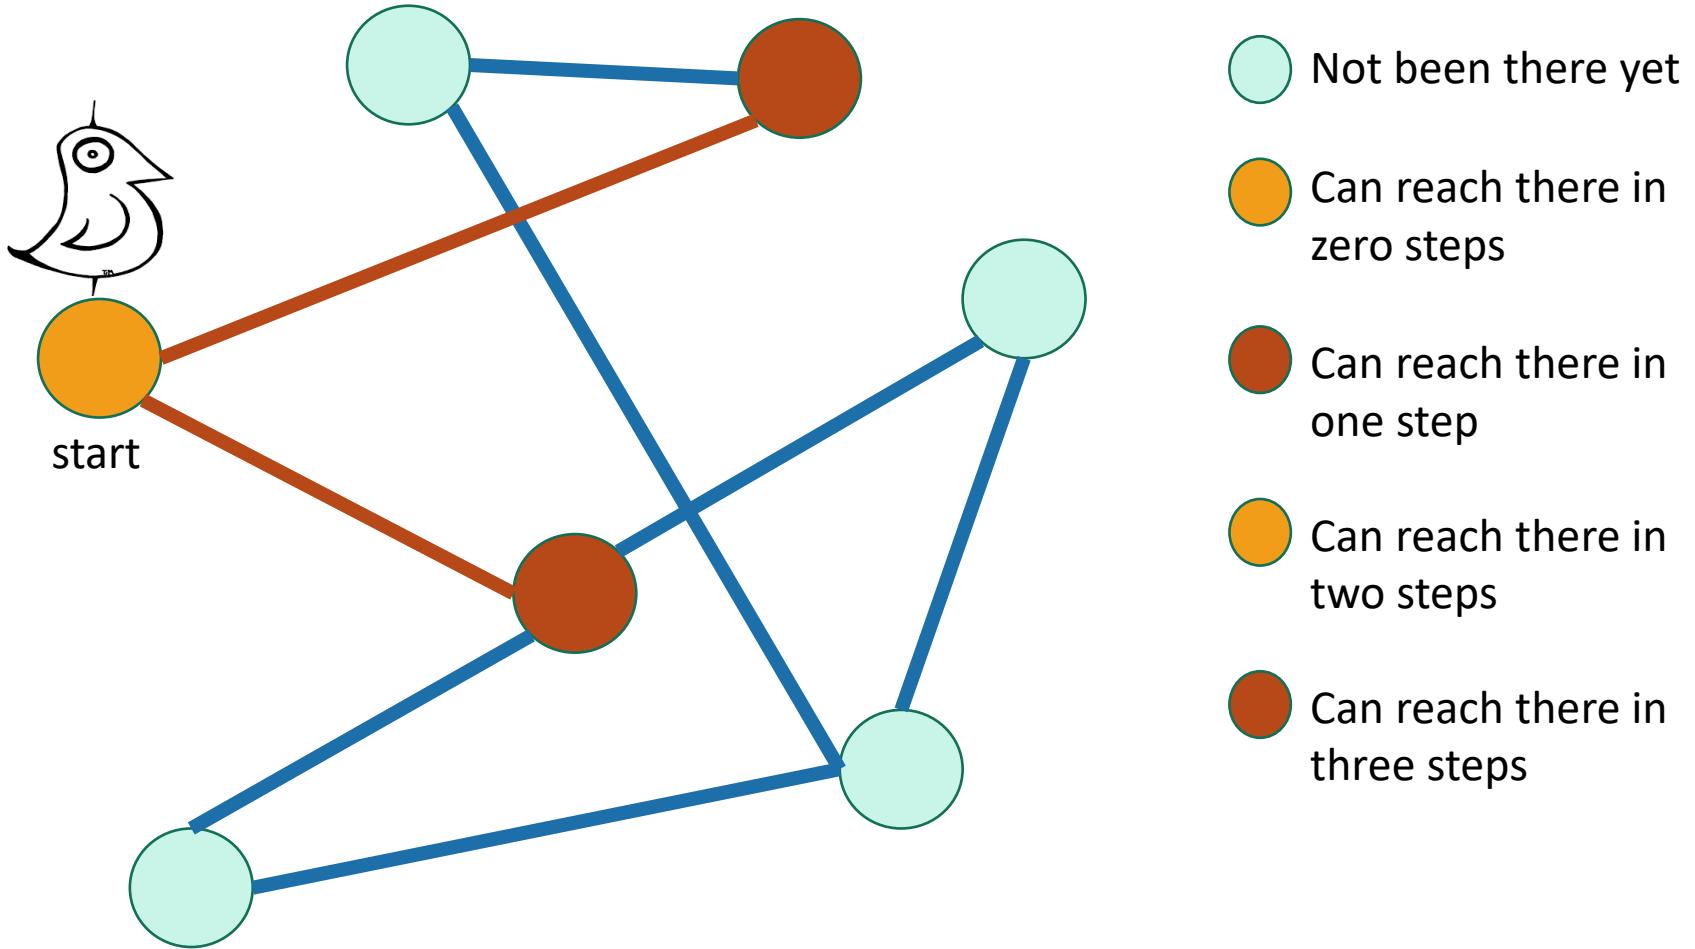
Breadth-First Search

For testing bipartite-ness



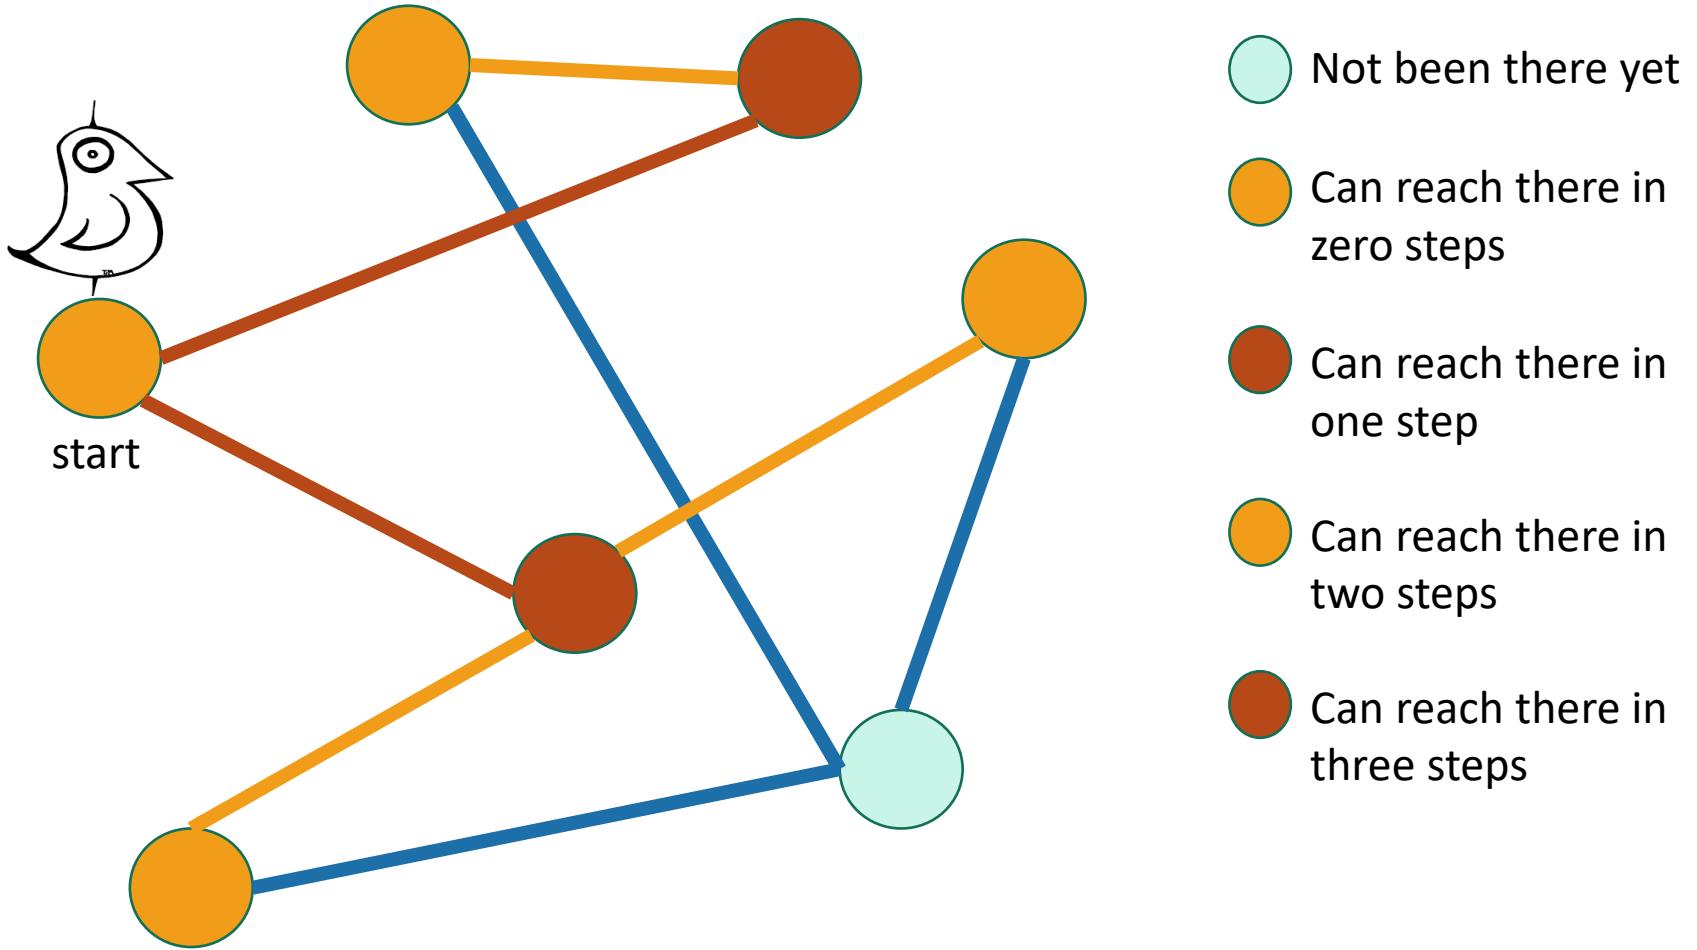
Breadth-First Search

For testing bipartite-ness



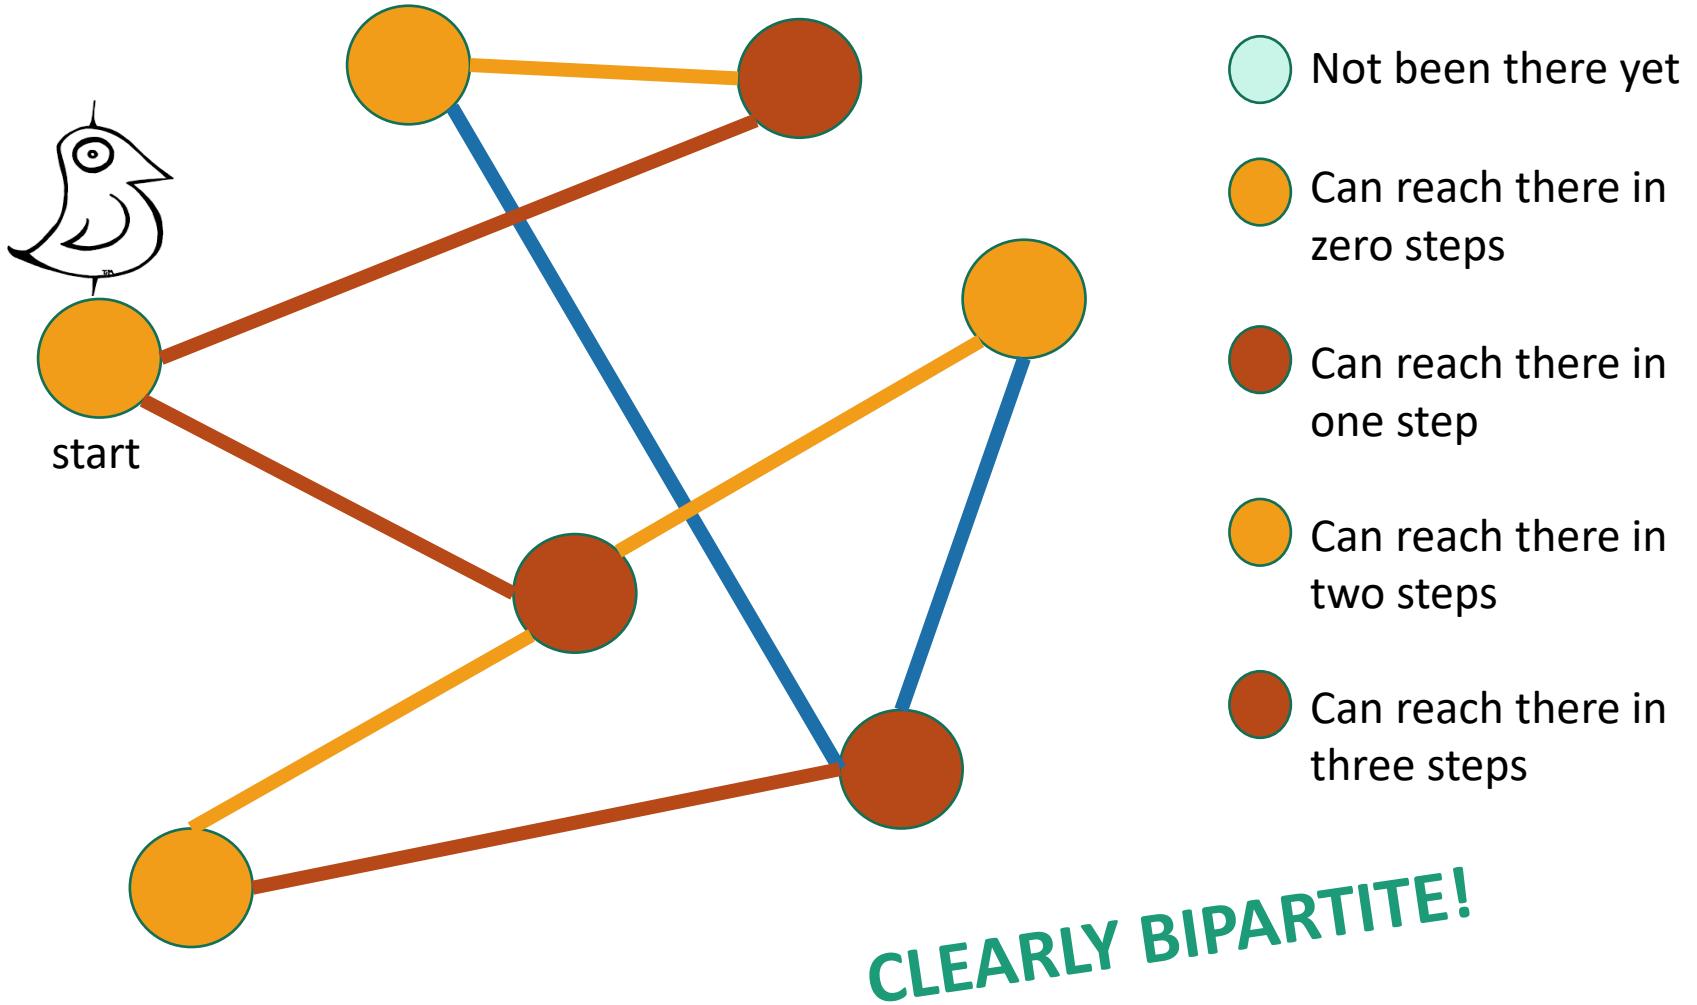
Breadth-First Search

For testing bipartite-ness



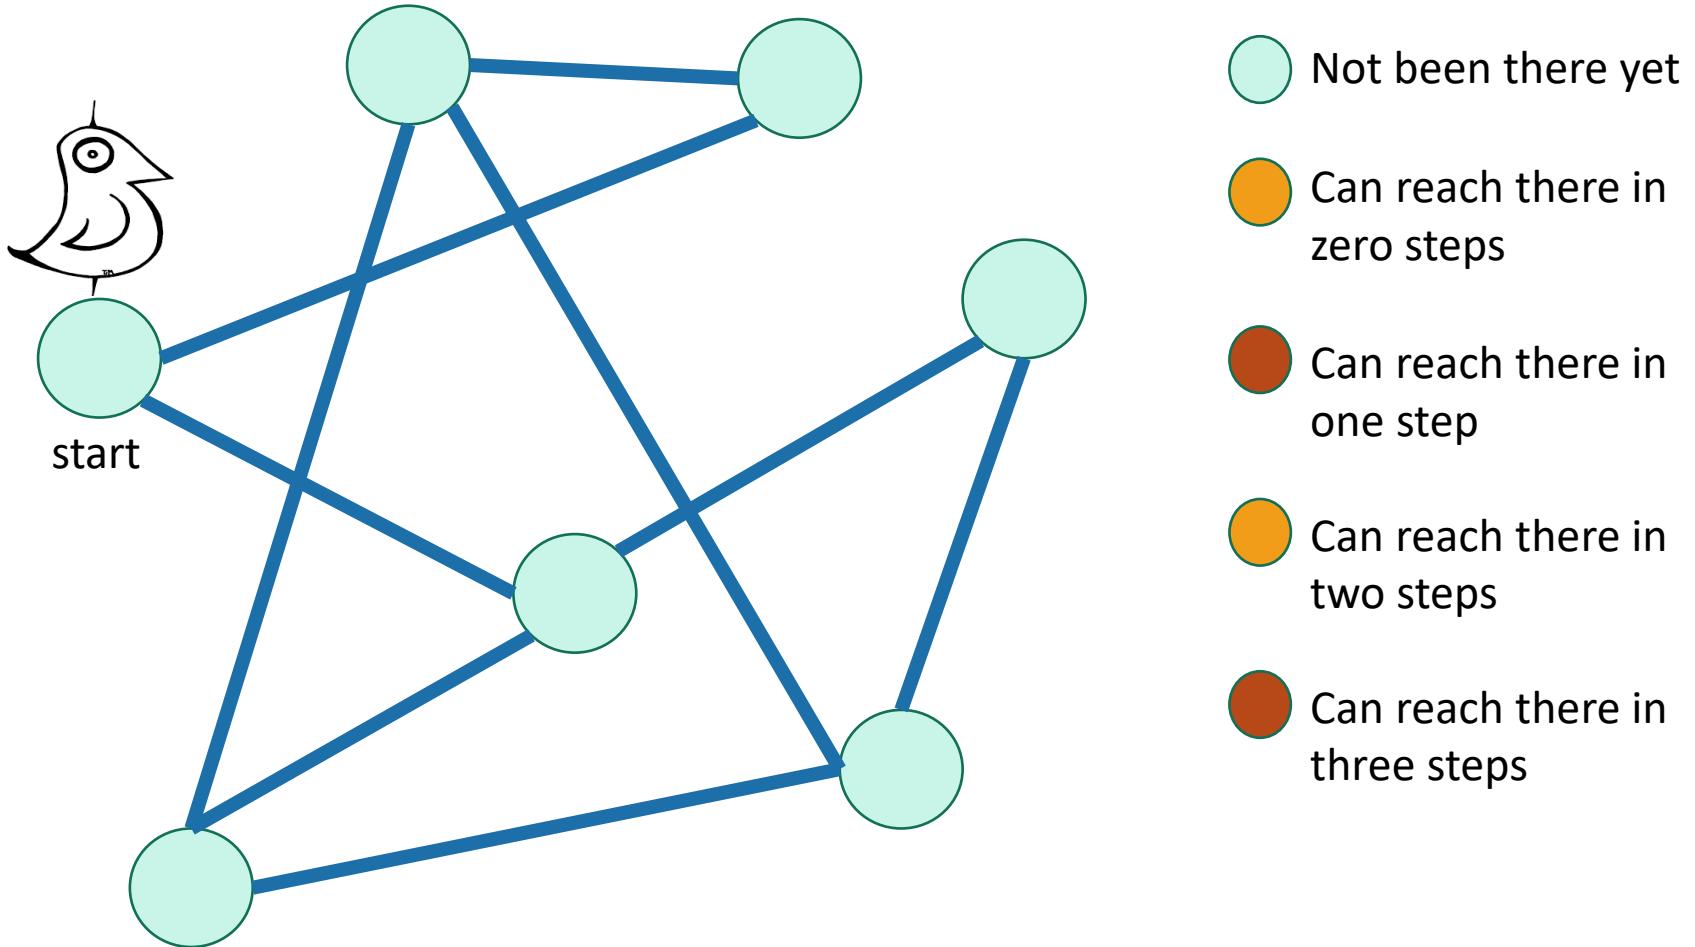
Breadth-First Search

For testing bipartite-ness



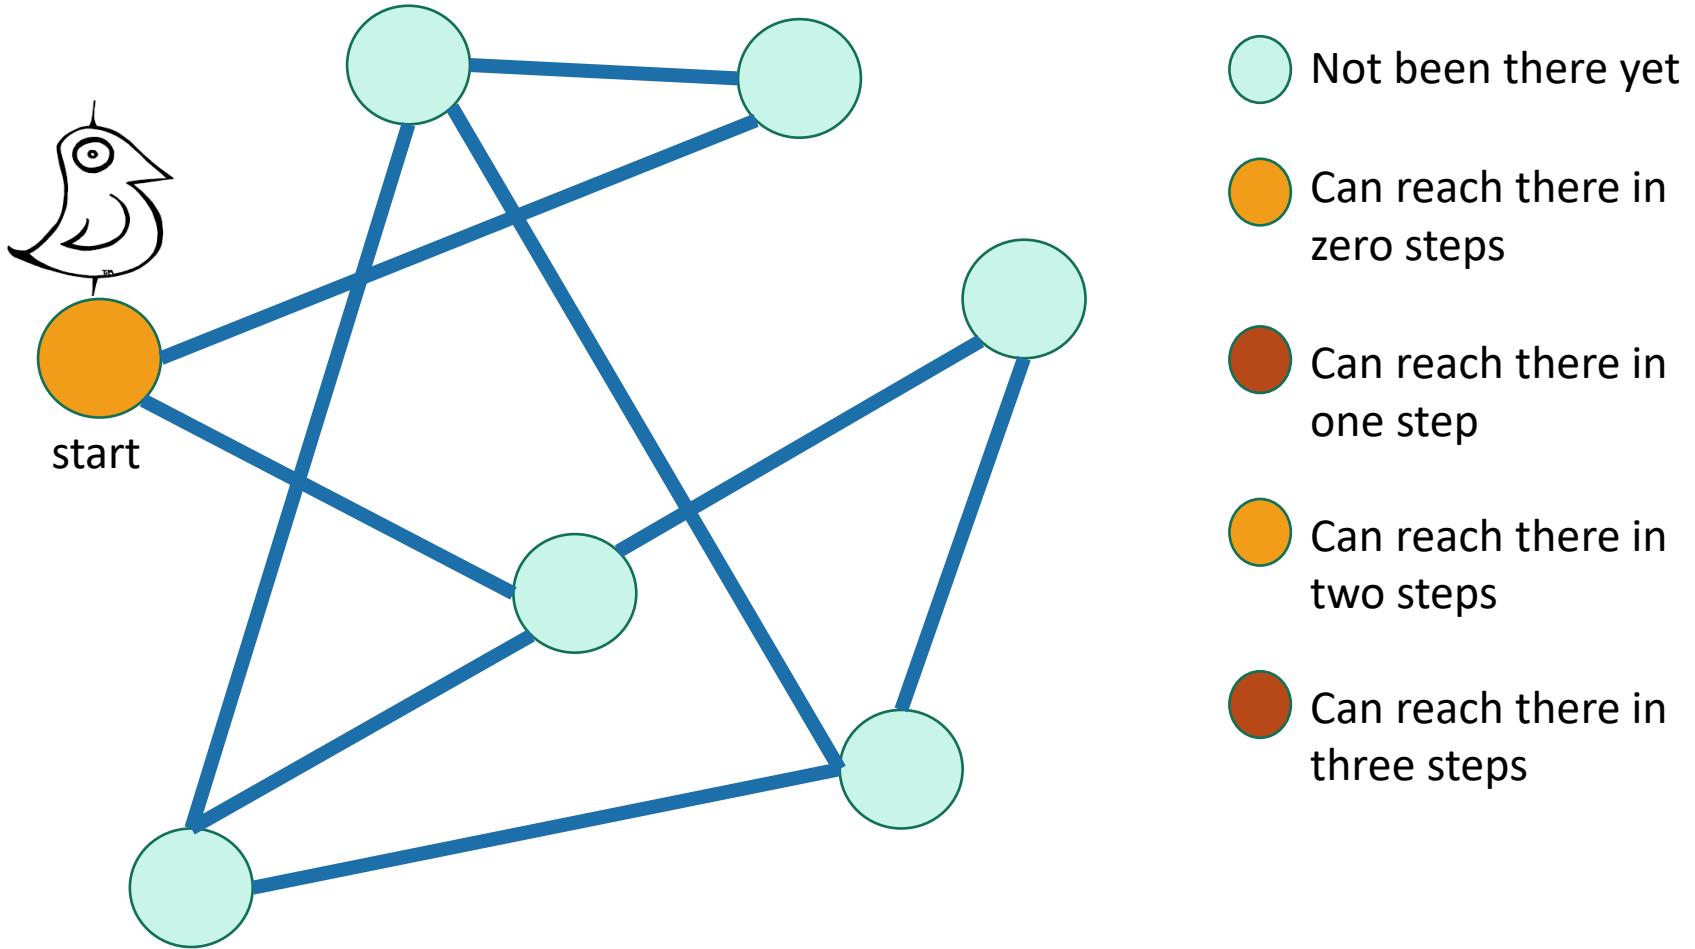
Breadth-First Search

For testing bipartite-ness



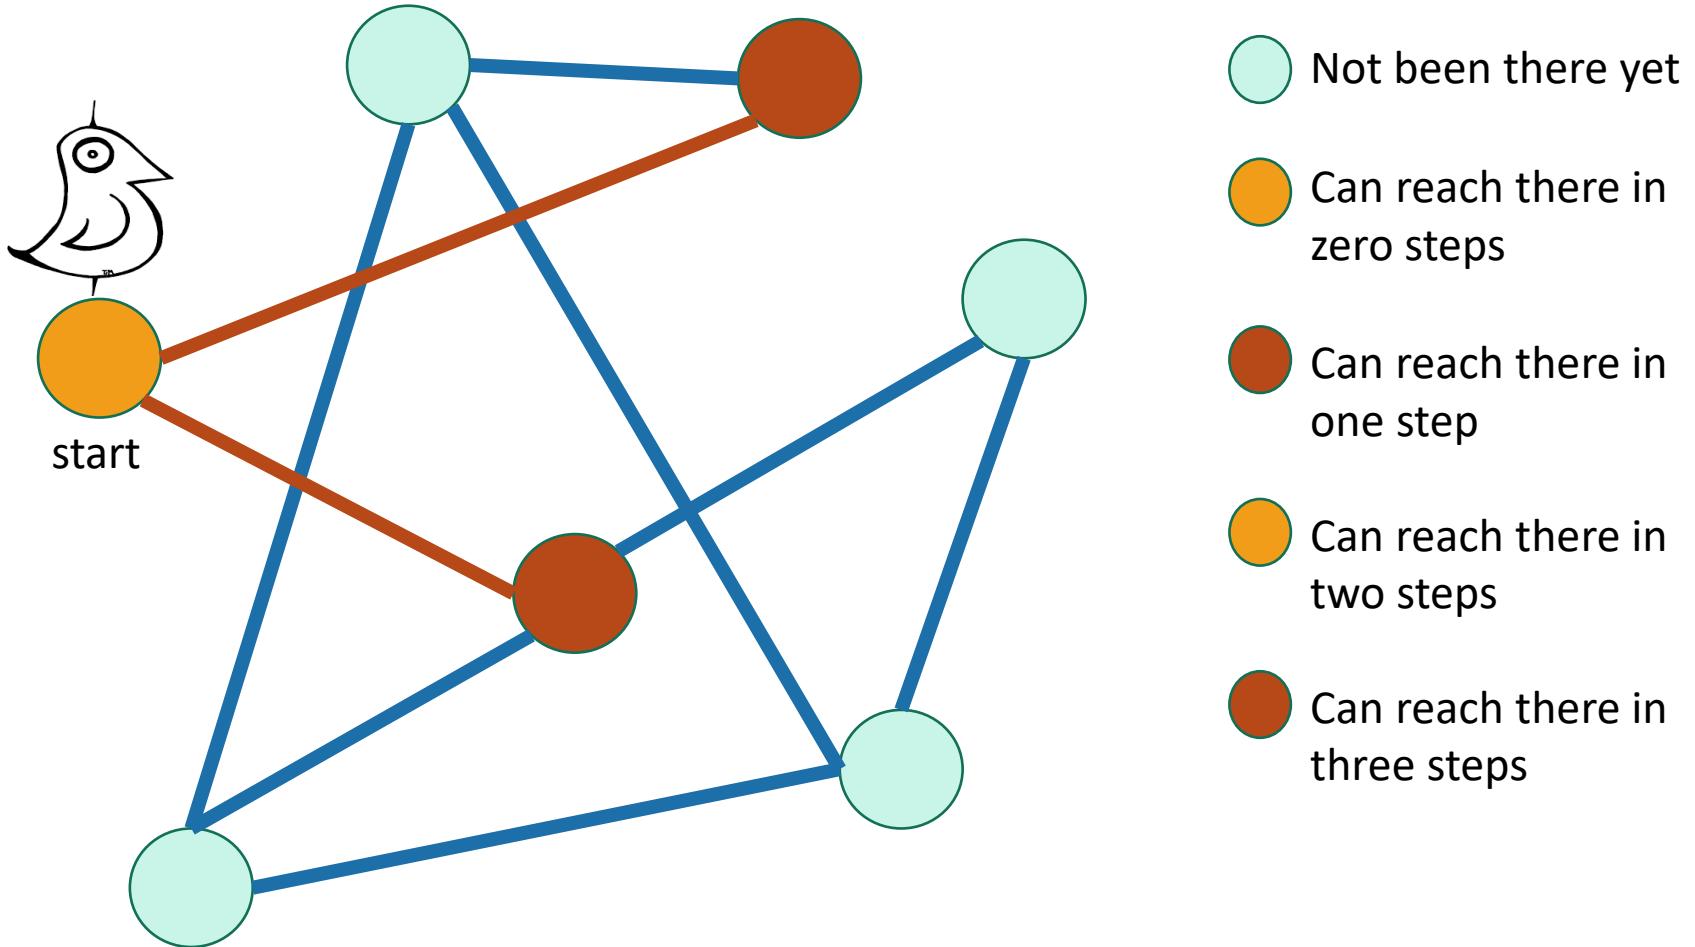
Breadth-First Search

For testing bipartite-ness



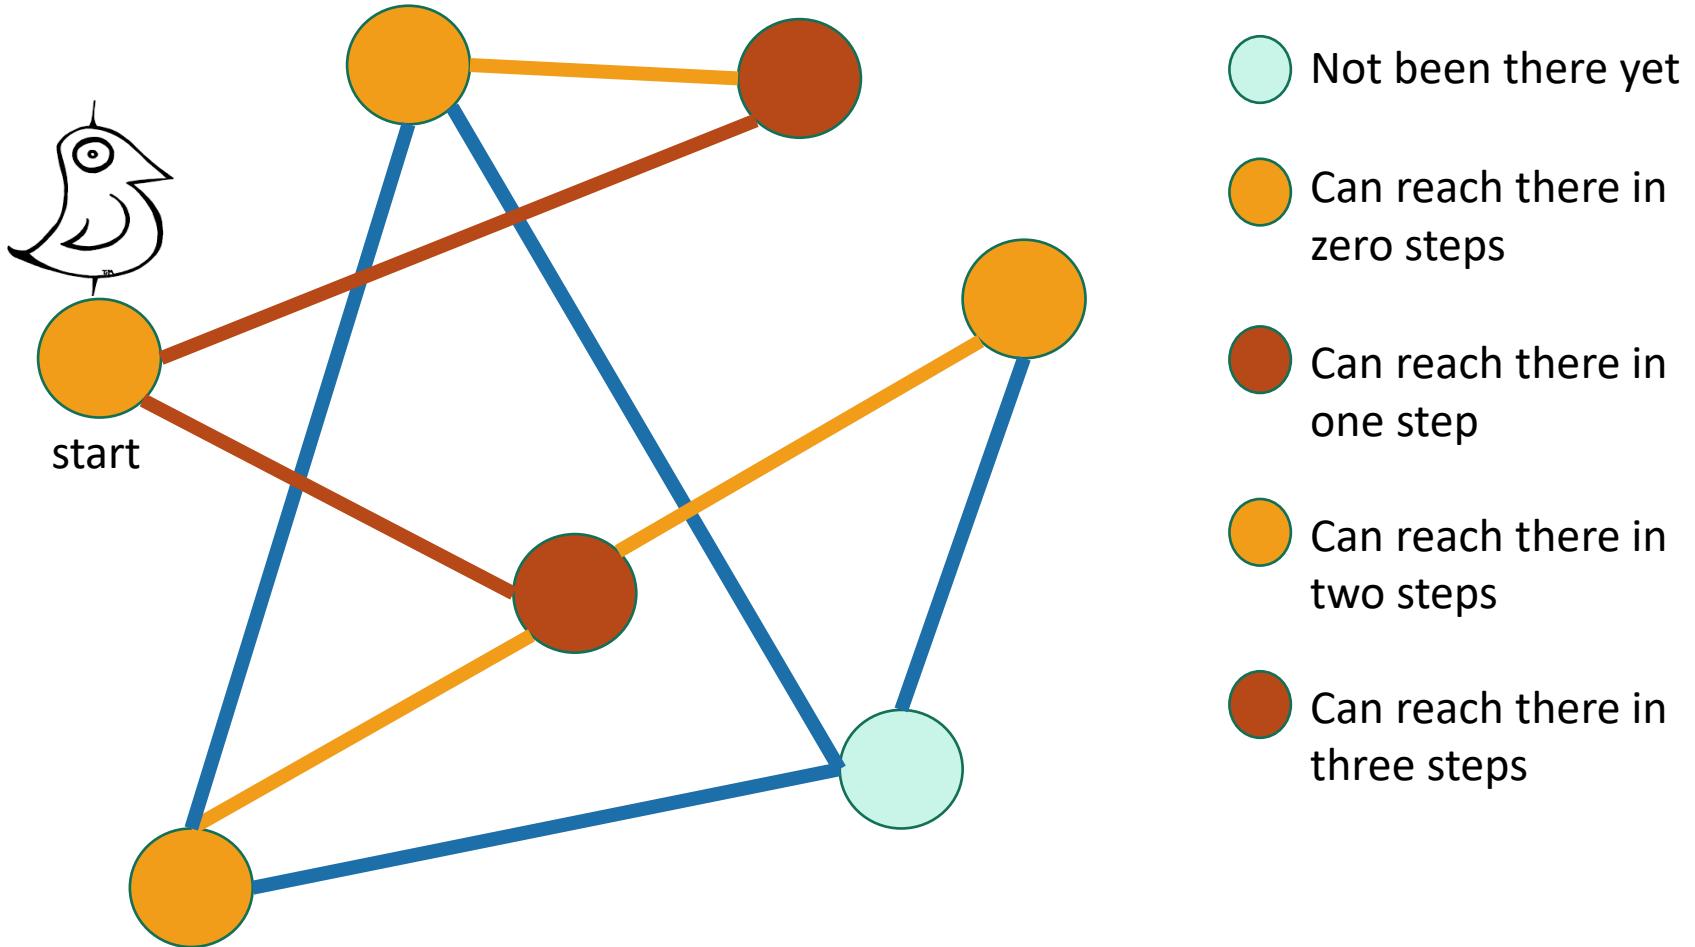
Breadth-First Search

For testing bipartite-ness



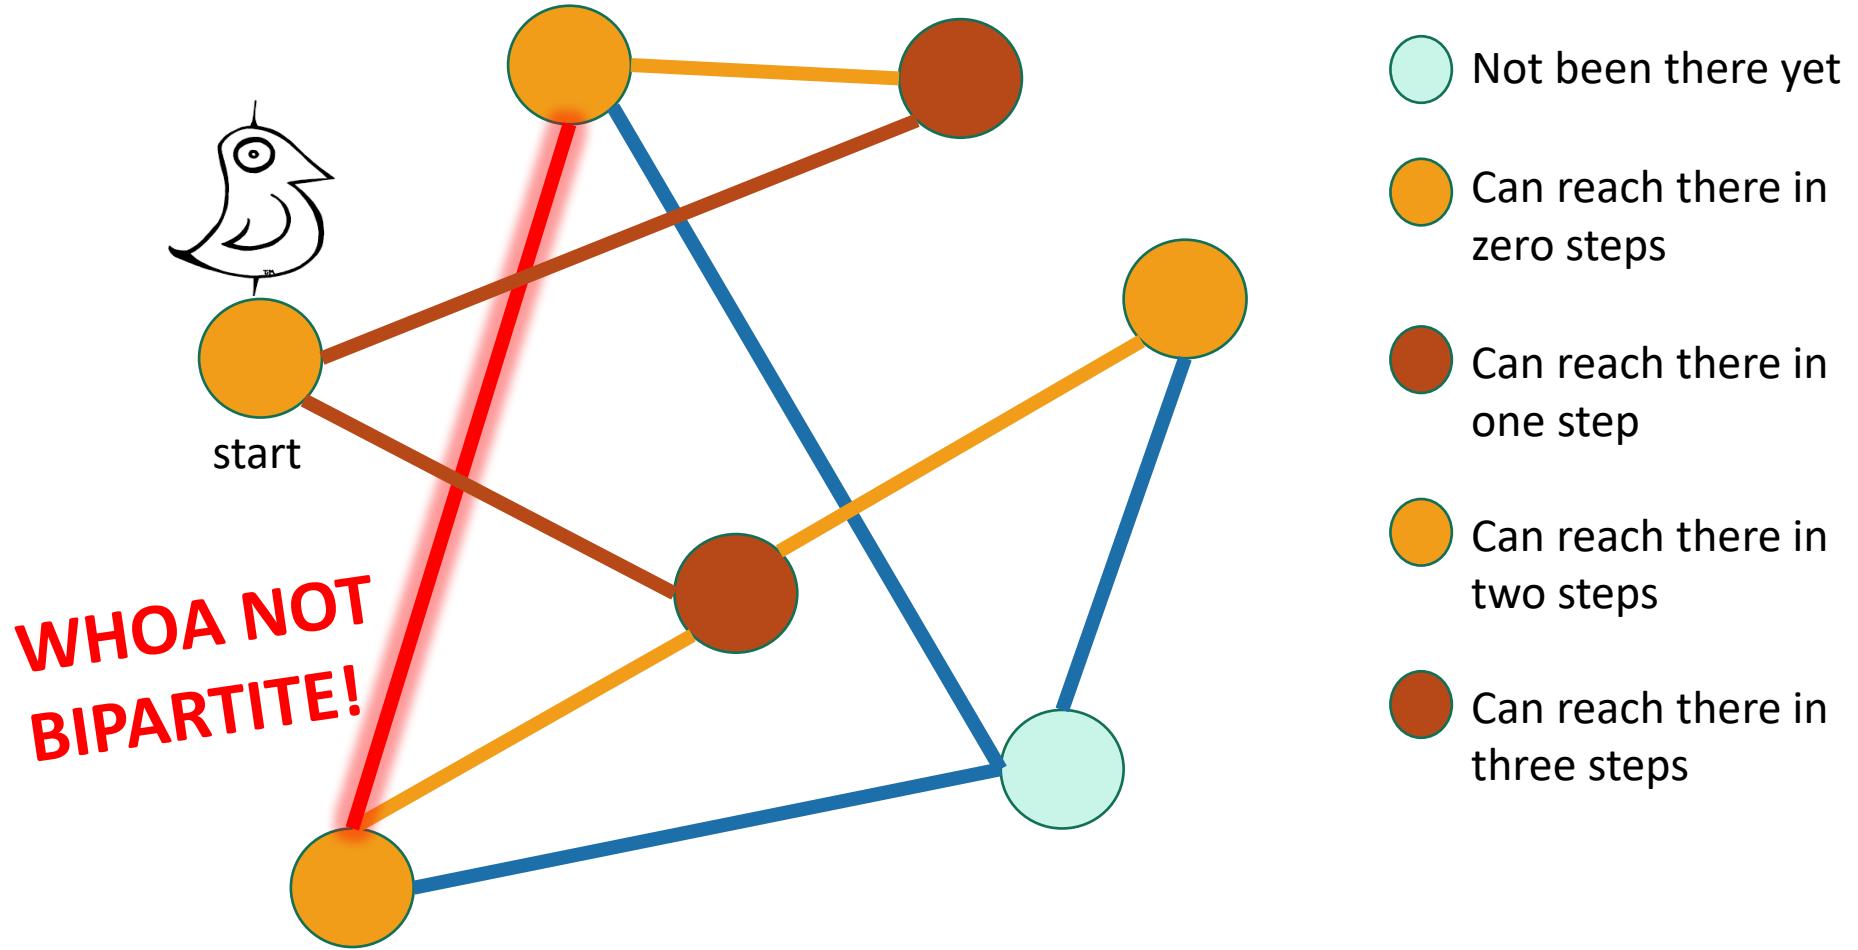
Breadth-First Search

For testing bipartite-ness



Breadth-First Search

For testing bipartite-ness



What did we just learn?

BFS can be used to detect bipartite-ness in time $O(n + m)$.

- Consider a hash table of size $m = 1000$ and a corresponding hash function
$$h(k) = \lfloor m (kA \bmod 1) \rfloor$$
 for . Compute the locations to which the keys 61, 62, 63, 64, and 65 are mapped. A is equal to 0.618033